

Aprendizado de Máquina e Reconhecimento de Padrões



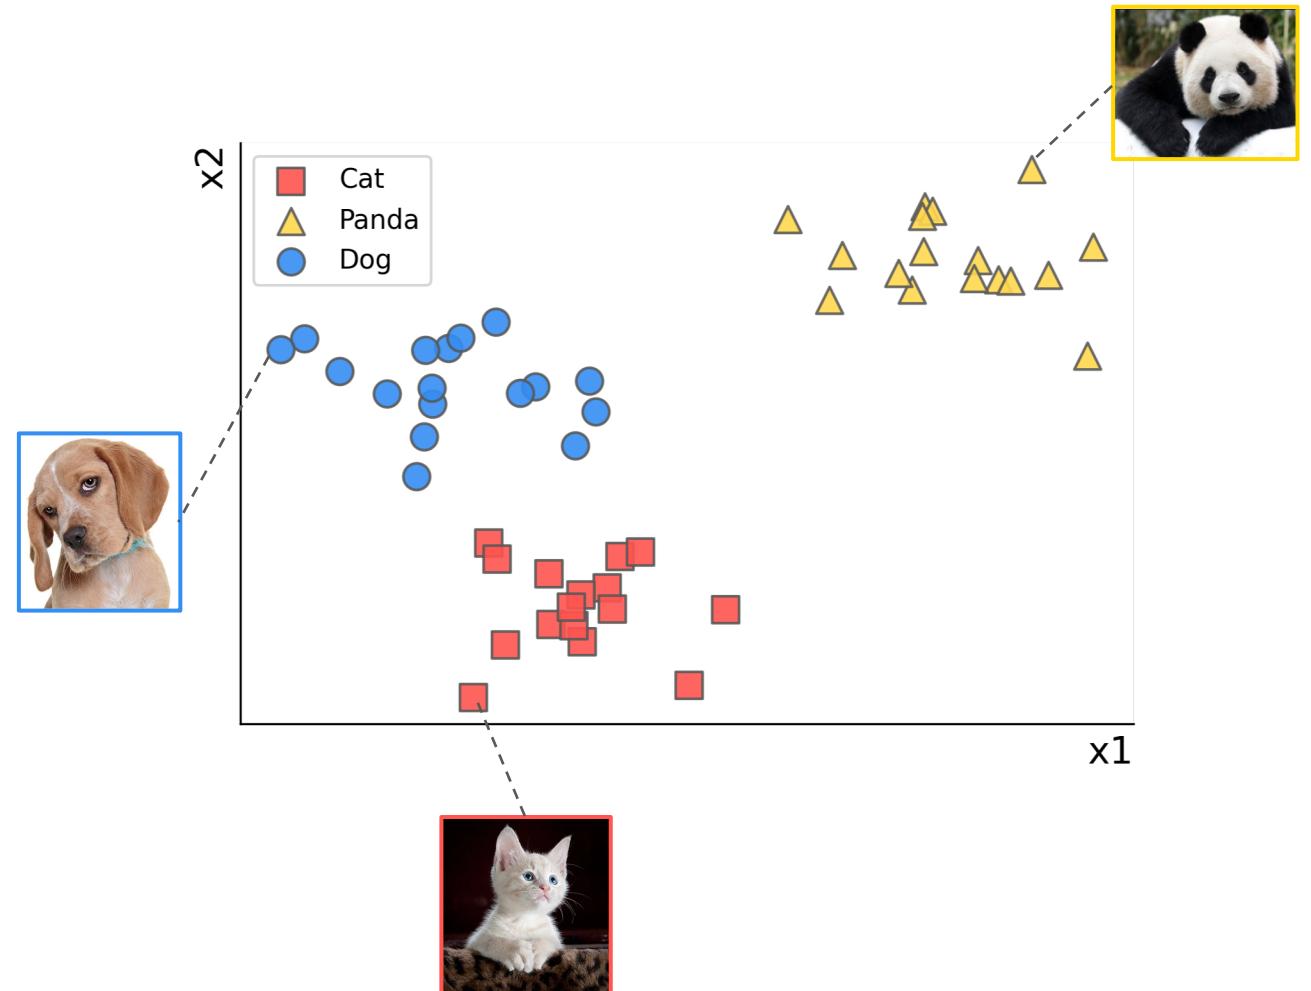
Kmeans

Prof. Dr. Samuel Martins (Samuka)
samuel.martins@ifsp.edu.br



Supervised Learning (e.g., classification)

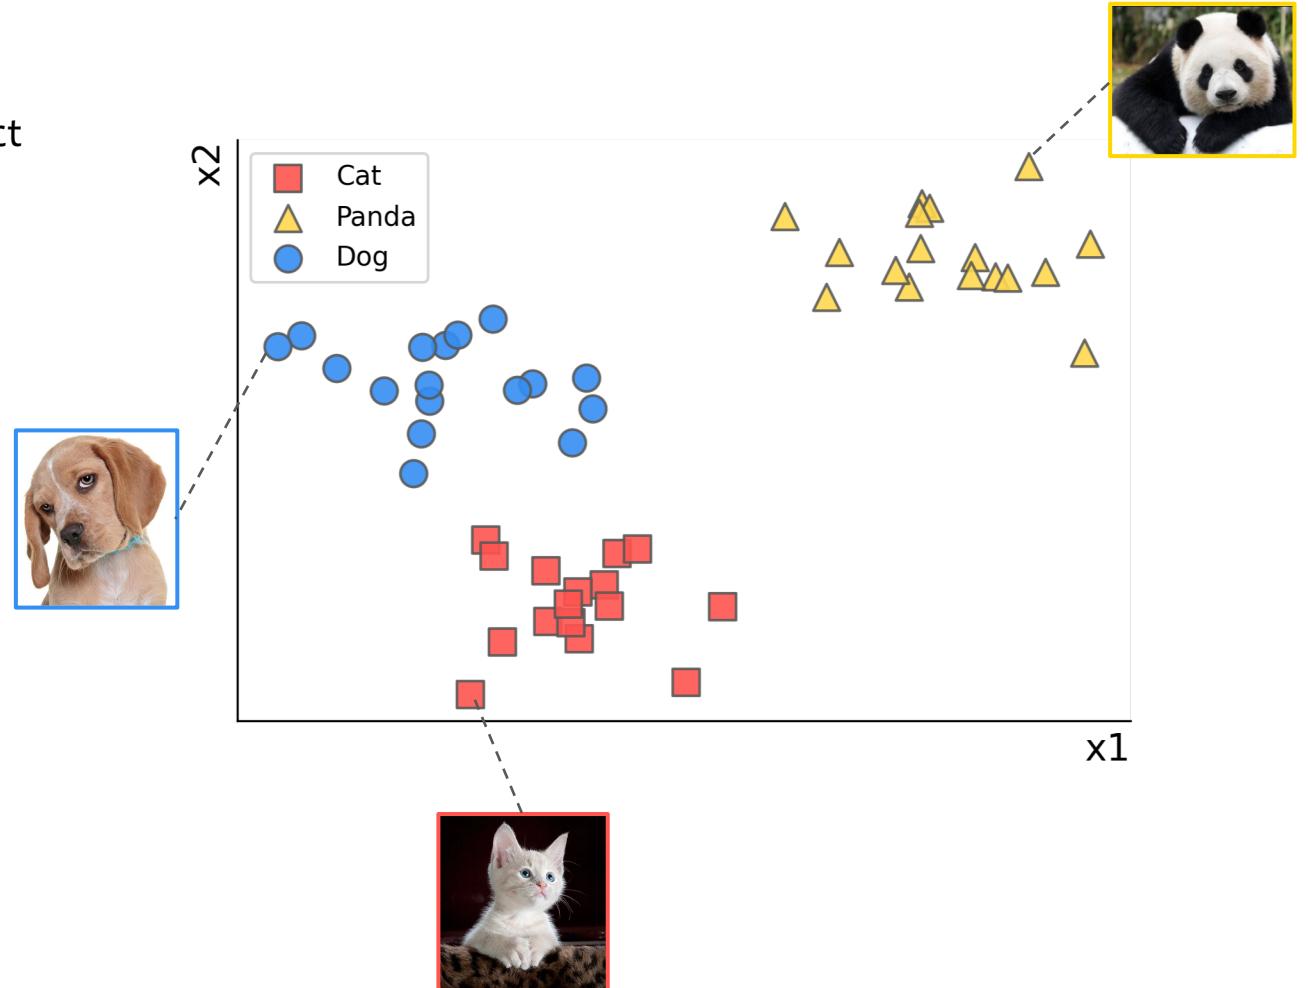
Previously **annotated/labeled training data**.



Supervised Learning (e.g., classification)

Previously **annotated/labeled training data**.

Mostly, **manual annotation**: human experts **manually inspect** and tag each sample (label assignment / data annotation).

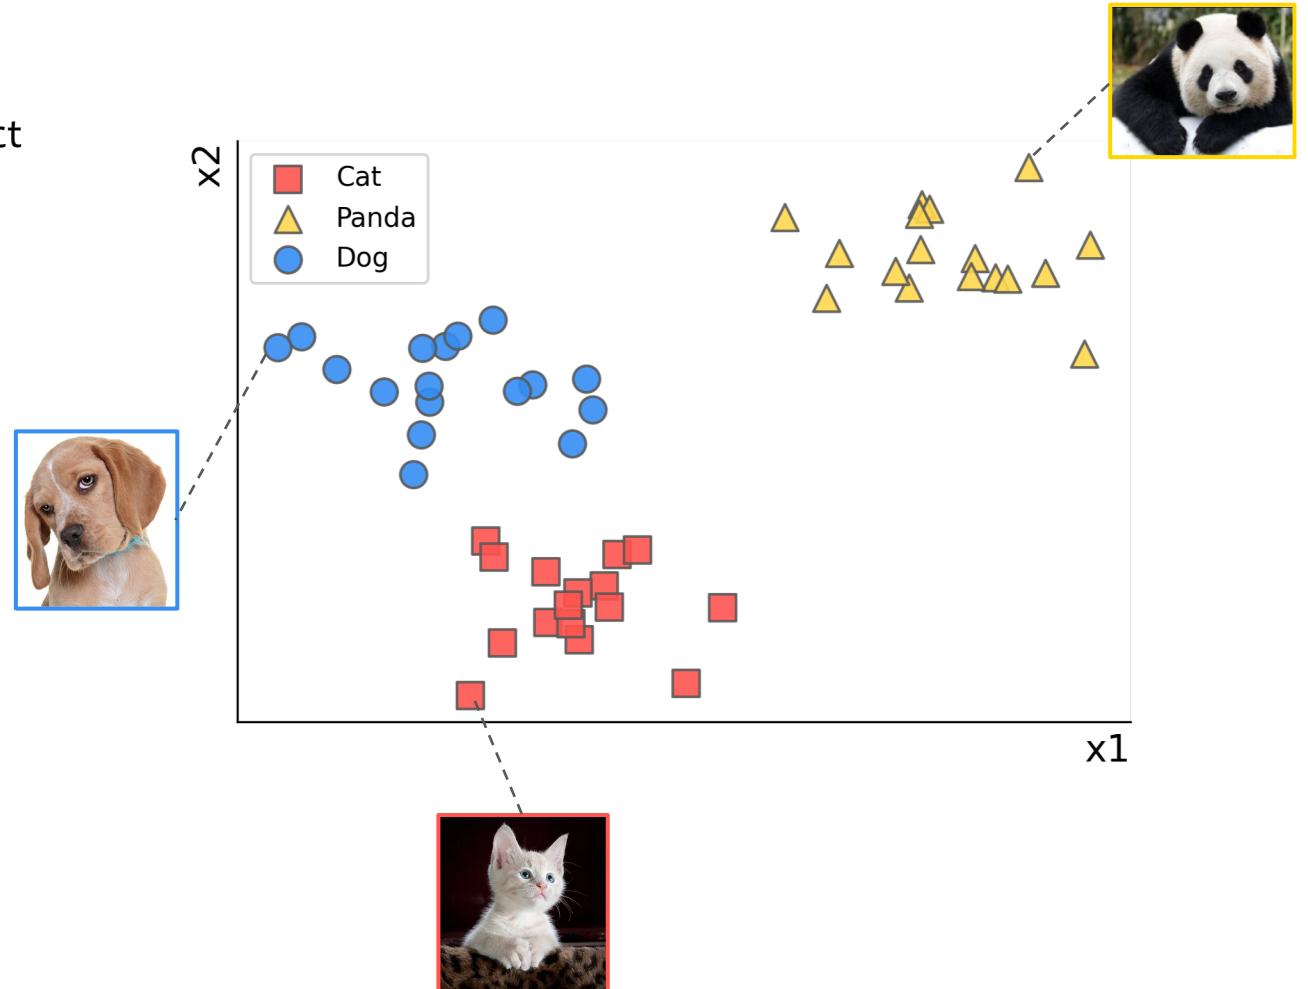


Supervised Learning (e.g., classification)

Previously **annotated/labeled training data**.

Mostly, **manual annotation**: human experts **manually inspect** and tag each sample (label assignment / data annotation).

Long, costly, error-prone, and tedious task, so it will usually only be done on a small subset of the available samples.



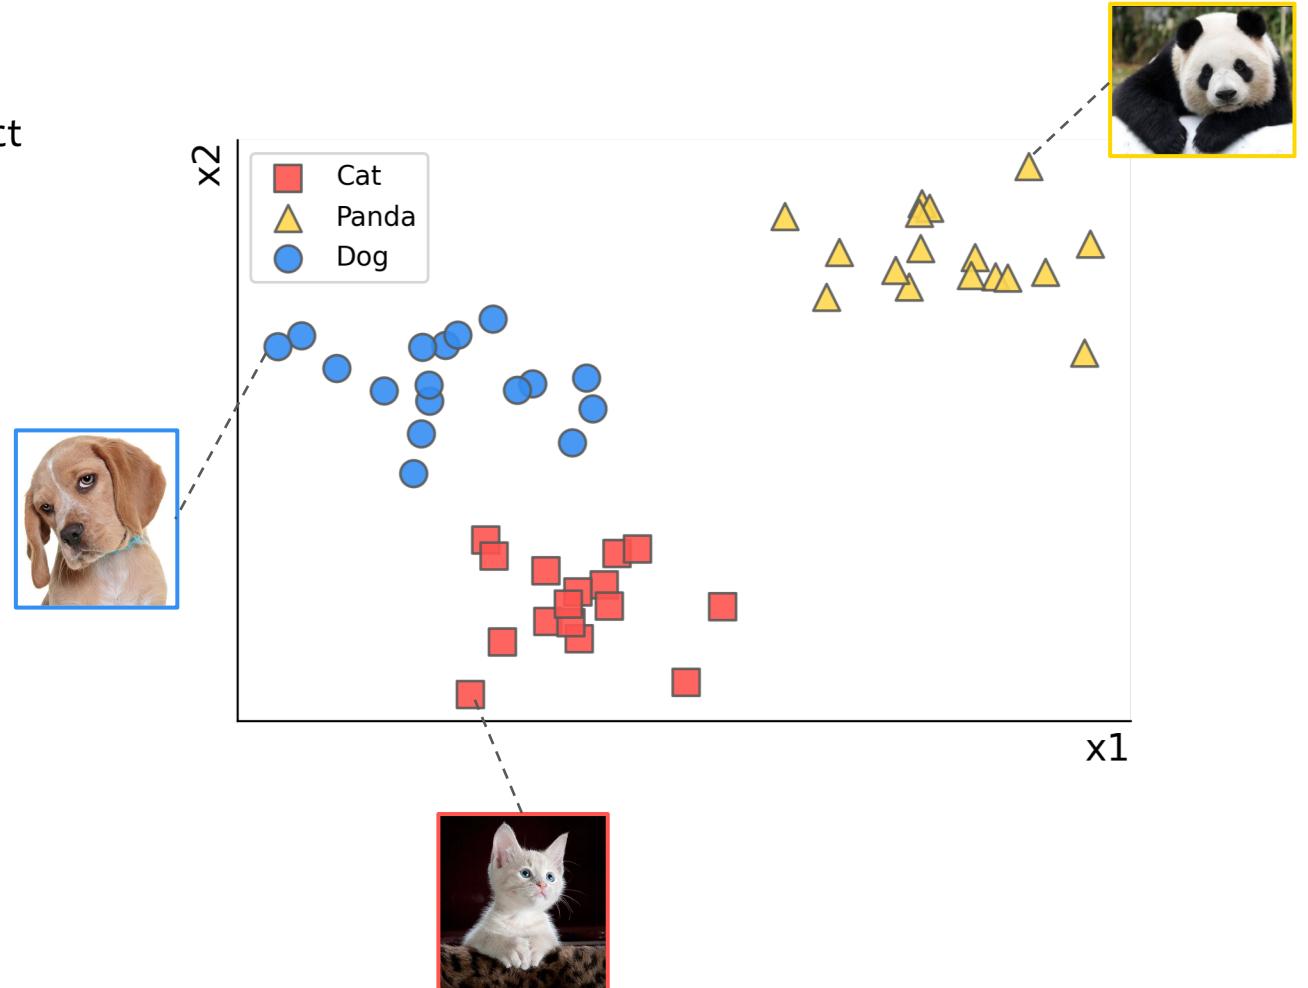
Supervised Learning (e.g., classification)

Previously **annotated/labeled training data**.

Mostly, **manual annotation**: human experts **manually inspect** and tag each sample (label assignment / data annotation).

Long, costly, error-prone, and tedious task, so it will usually only be done on a small subset of the available samples.

As a result, the **labeled dataset** will be **quite small**, and the classifier's performance will be disappointing.



Supervised Learning (e.g., classification)

Previously **annotated/labeled training data**.

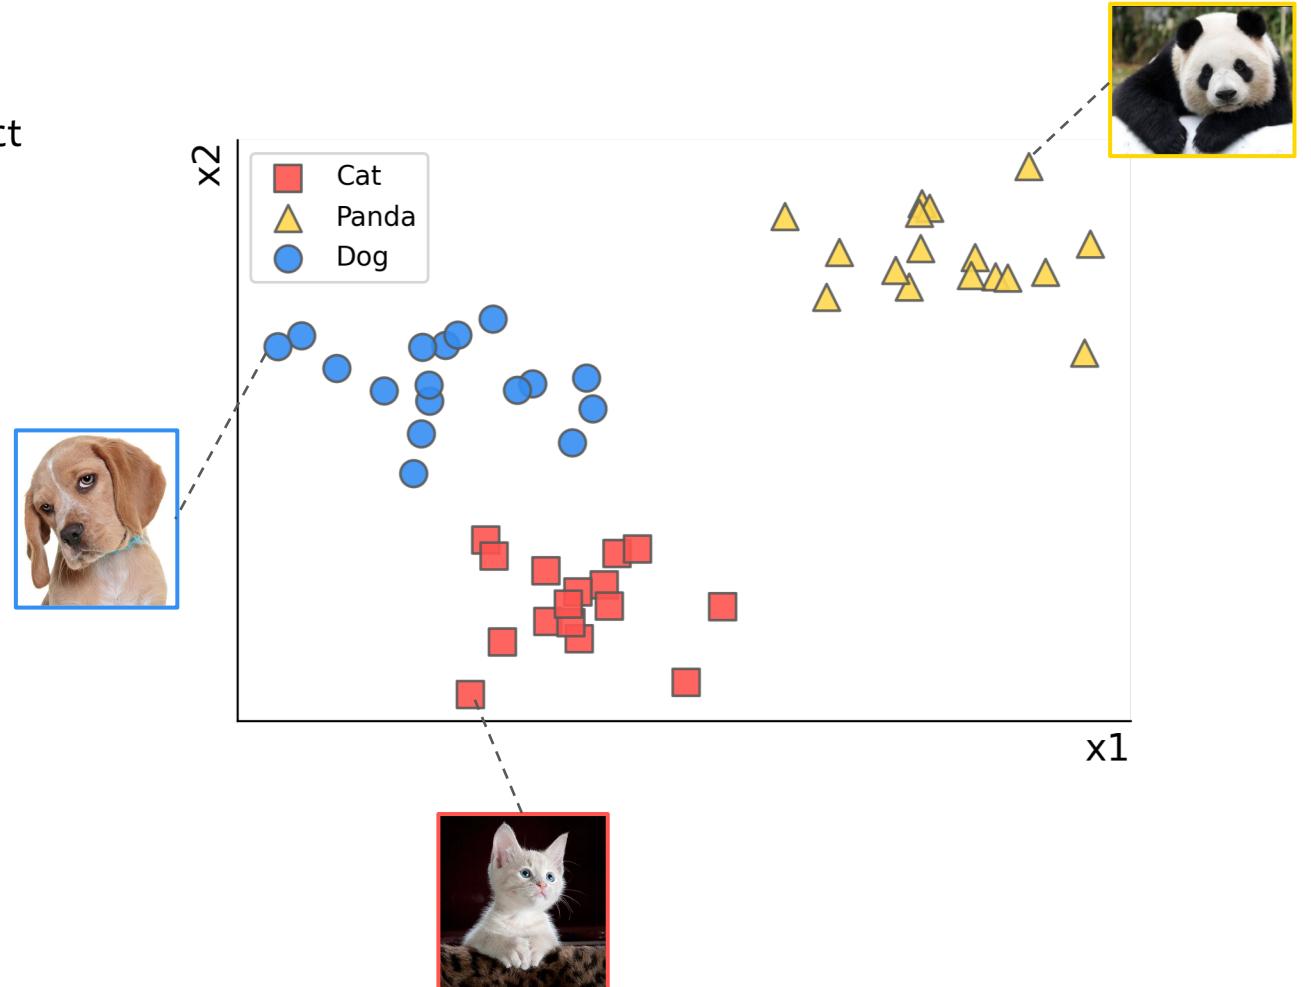
Mostly, **manual annotation**: human experts **manually inspect** and tag each sample (label assignment / data annotation).

Long, costly, error-prone, and tedious task, so it will usually only be done on a small subset of the available samples.

As a result, the **labeled dataset** will be **quite small**, and the classifier's performance will be disappointing.



These limitations are more accentuated in some applications, e.g., medical diagnosis, for demanding specific experts often unavailable for data annotation.



Supervised Learning (e.g., classification)

Previously **annotated/labeled training data**.

Mostly, **manual annotation**: human experts
and tag each sample by hand.

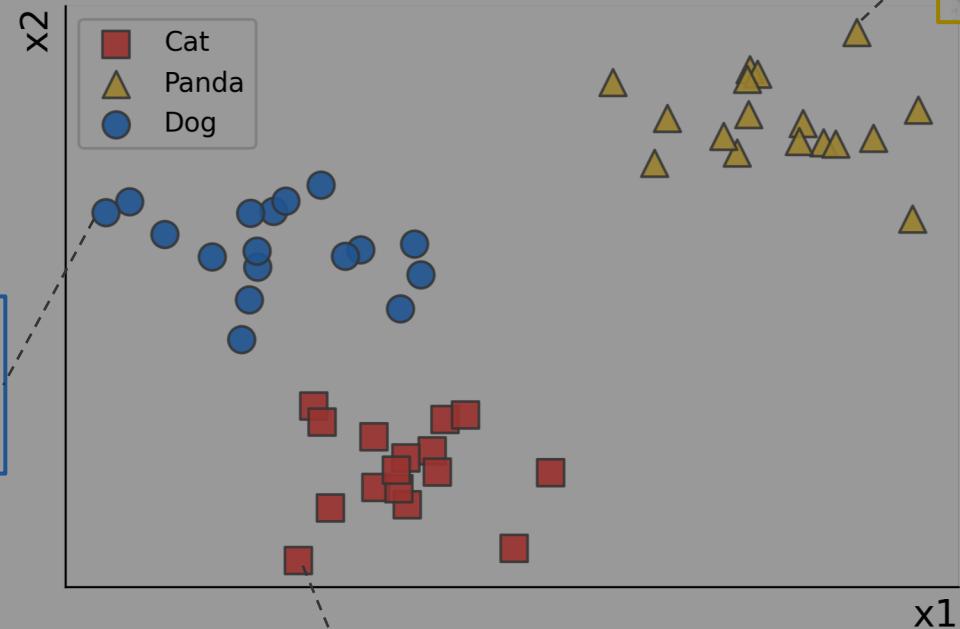
Wouldn't it be great if the algorithm could just
exploit the unlabeled data without needing
humans to label every sample?

Long, costly, and tedious task, so it will
usually only be done on a small subset of the available
samples.

As a result, the **labeled dataset** will be **quite small**, and
the classifier's performance will be disappointing.



These limitations are more accentuated in
some applications, e.g., medical diagnosis,
for demanding specific experts often
unavailable for data annotation.



Supervised Learning (e.g., classification)

Previously **annotated/labeled training data**.

Mostly, **manual annotation**: human experts
and tag each sample.

Wouldn't it be great if the algorithm could just
exploit the unlabeled data without needing
humans to label every sample?

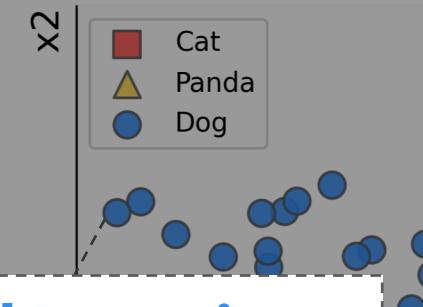
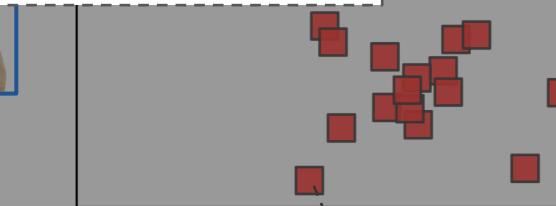
Long, costly, tedious task, so it will
usually only be done on a small subset of the available
samples.

Unsupervised Learning

As a result, the **labeled dataset** will be **quite small**, and
the classifier's performance will be disappointing.



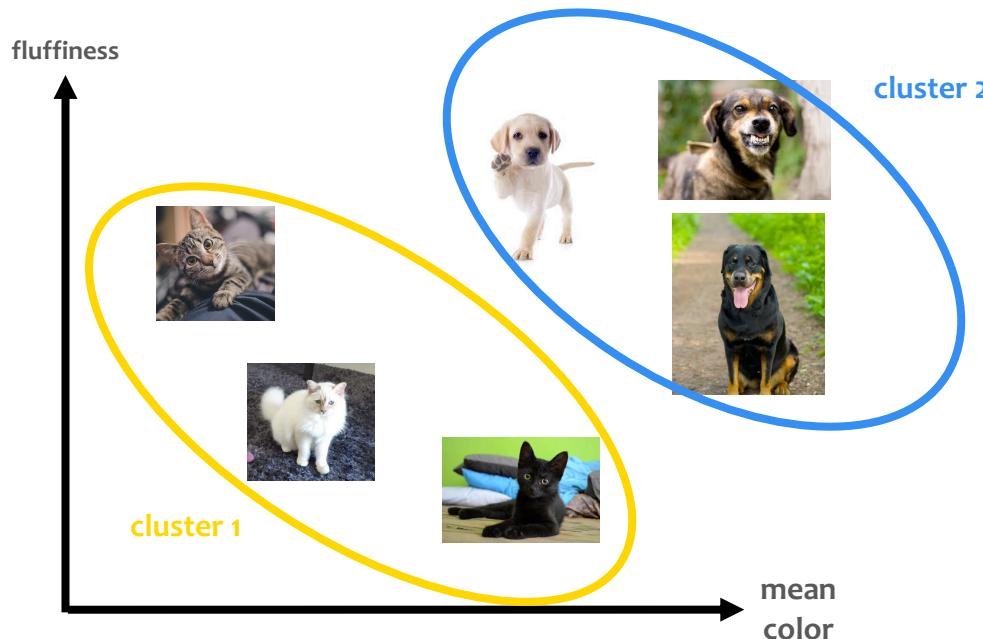
These limitations are more accentuated in
some applications, e.g., medical diagnosis,
for demanding specific experts often
unavailable for data annotation.



Unsupervised Learning

no supervision → no labels / annotations

Clustering

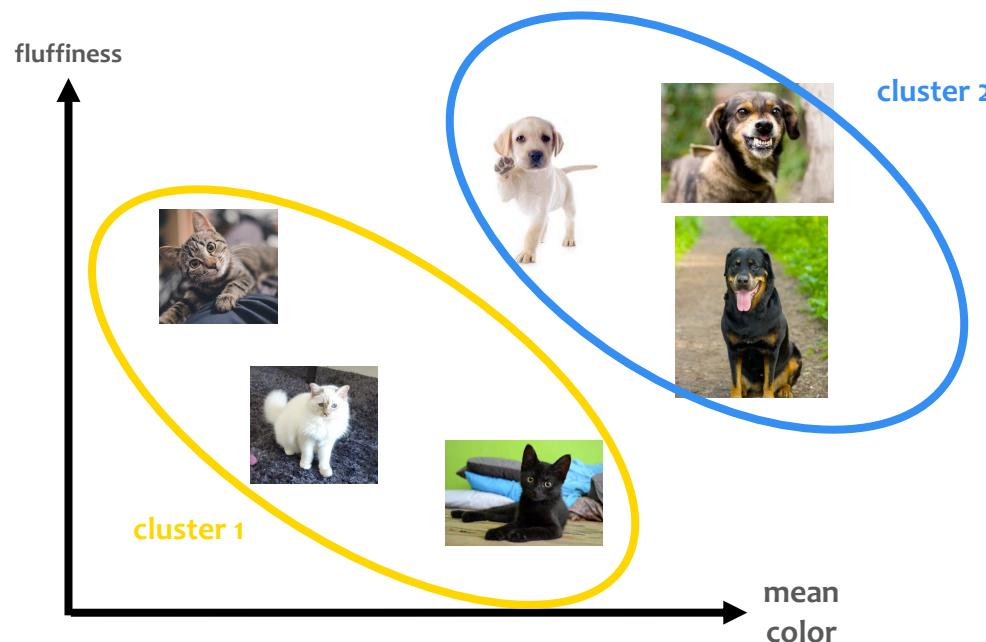


- k-Means
- DBSCAN
- Mean-shift
- OPF Clustering
- Hierarchical Clustering
- ...

Unsupervised Learning

no supervision → no labels / annotations

Clustering

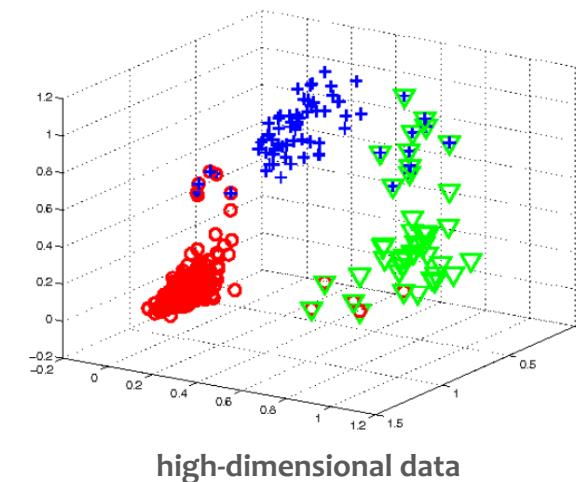


- k-Means
- DBSCAN
- Mean-shift
- OPF Clustering
- Hierarchical Clustering
- ...

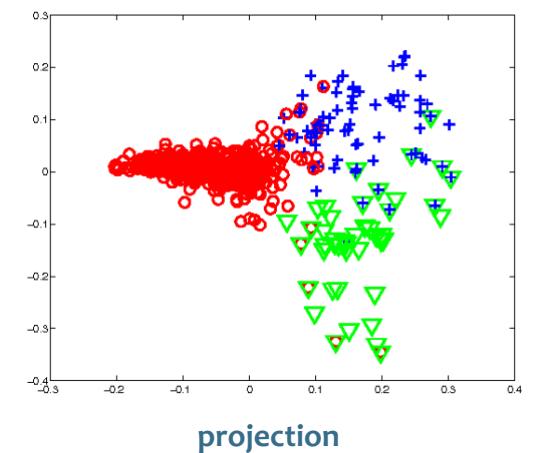
Dimensionality Reduction / Projection

Project a **high-dimensional data** on to a **lower-dimensional feature space** without losing much information.

Projection on to a **2D** or **3D** space is used to **visualize and understand** the data.



high-dimensional data



projection

- Principal Component Analysis (PCA)
- Locally Linear Embedding (LLE)
- t-SNE
- UMAP
- ...

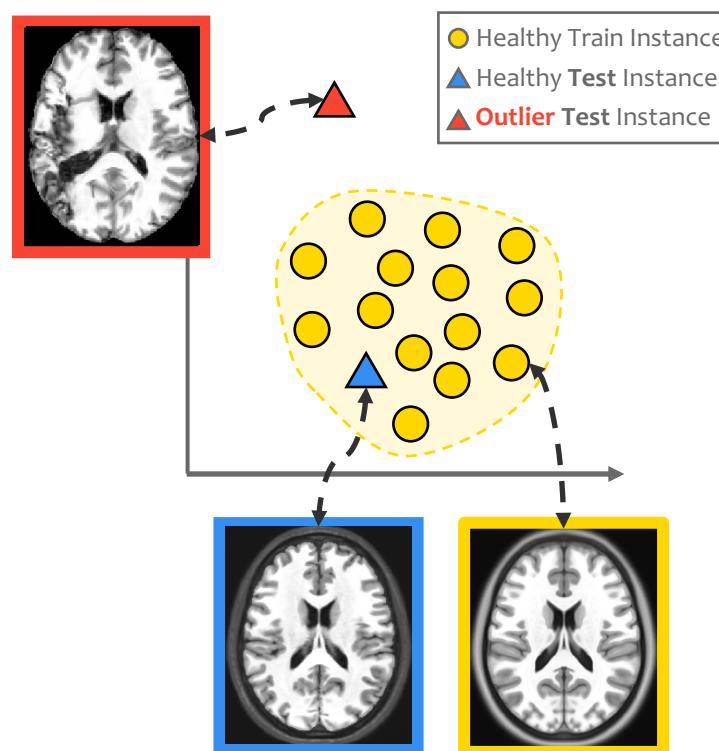
Unsupervised Learning

Anomaly/Novelty/Outlier Detection

The algorithm is fed with **training data** containing **only** a **single class** (normal, control, healthy).

New unseen instances (test) that **DOES NOT** look a **normal one** is an **outlier**, being considered as an **anomaly**.

The **training set** must be very **clean** and **representative**.



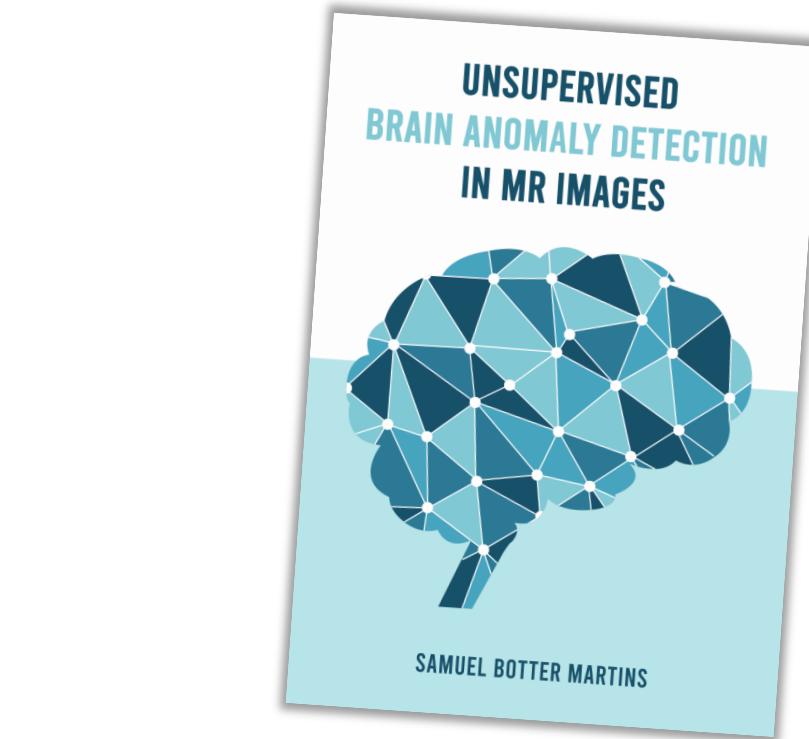
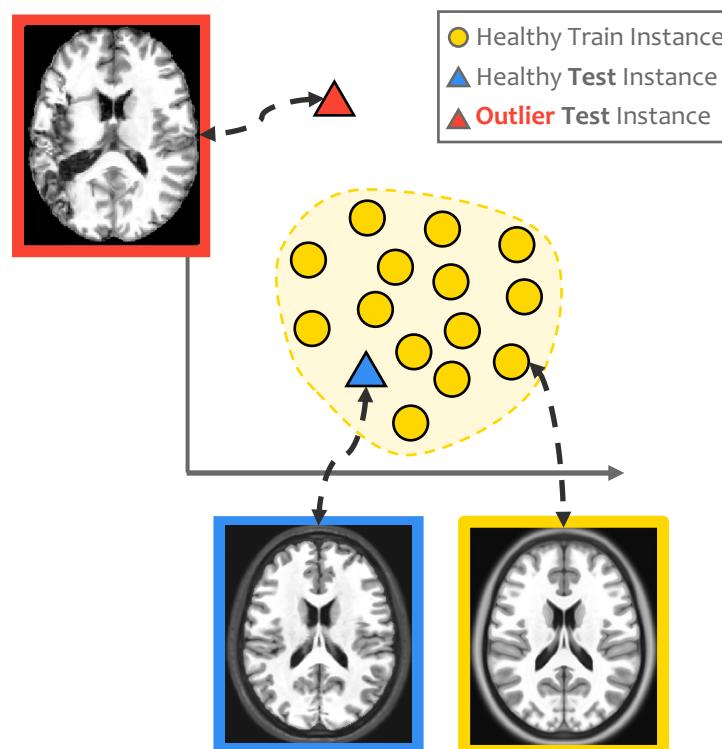
Unsupervised Learning

Anomaly/Novelty/Outlier Detection

The algorithm is fed with **training data** containing **only** a **single class** (normal, control, healthy).

New unseen instances (test) that **DOES NOT** look a **normal one** is an **outlier**, being considered as an **anomaly**.

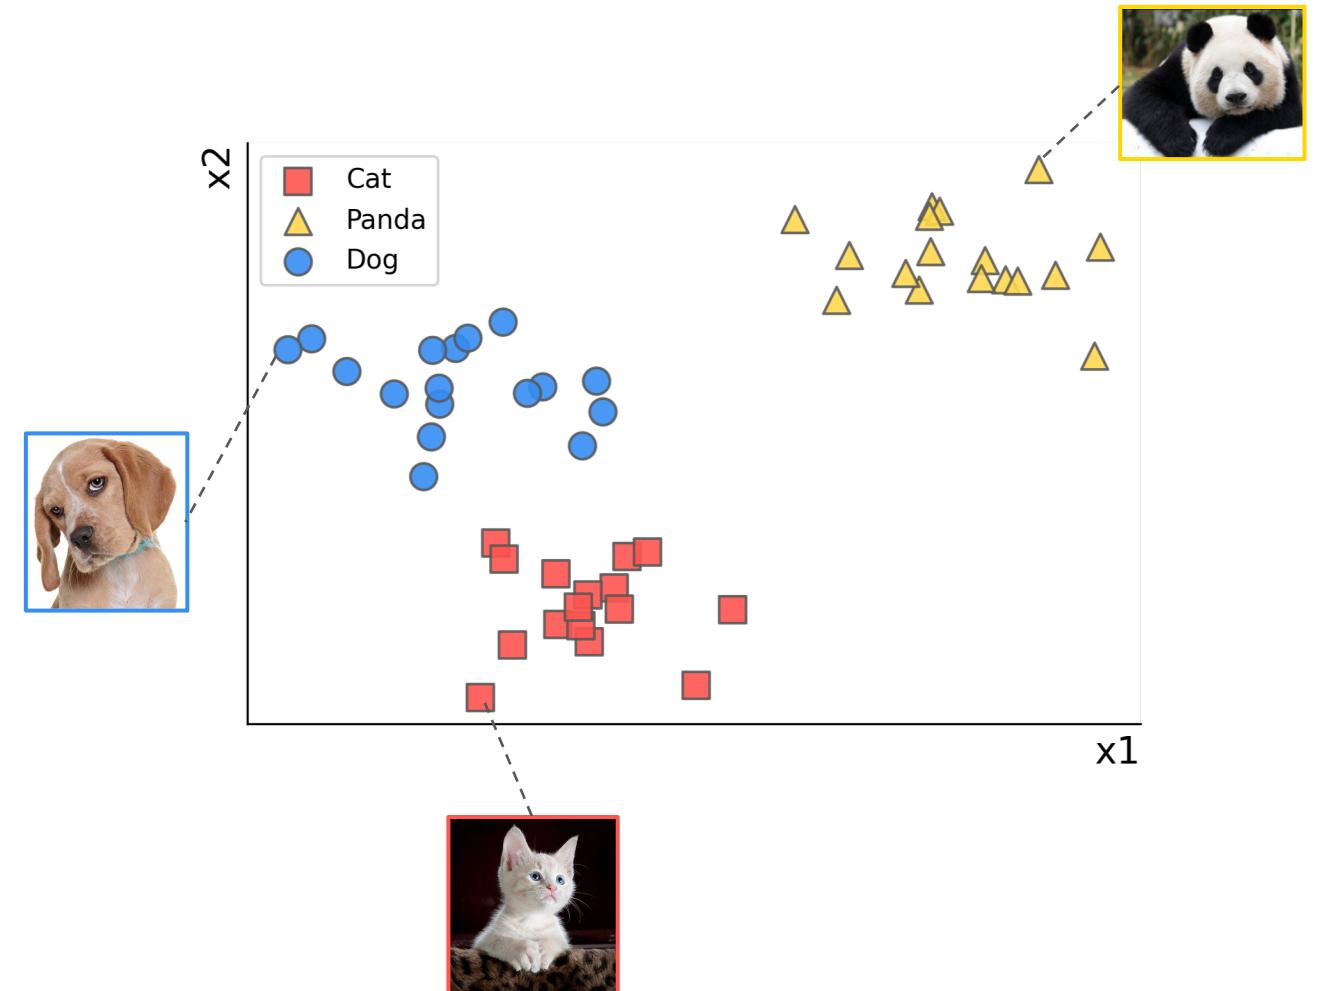
The **training set** must be very **clean** and **representative**.



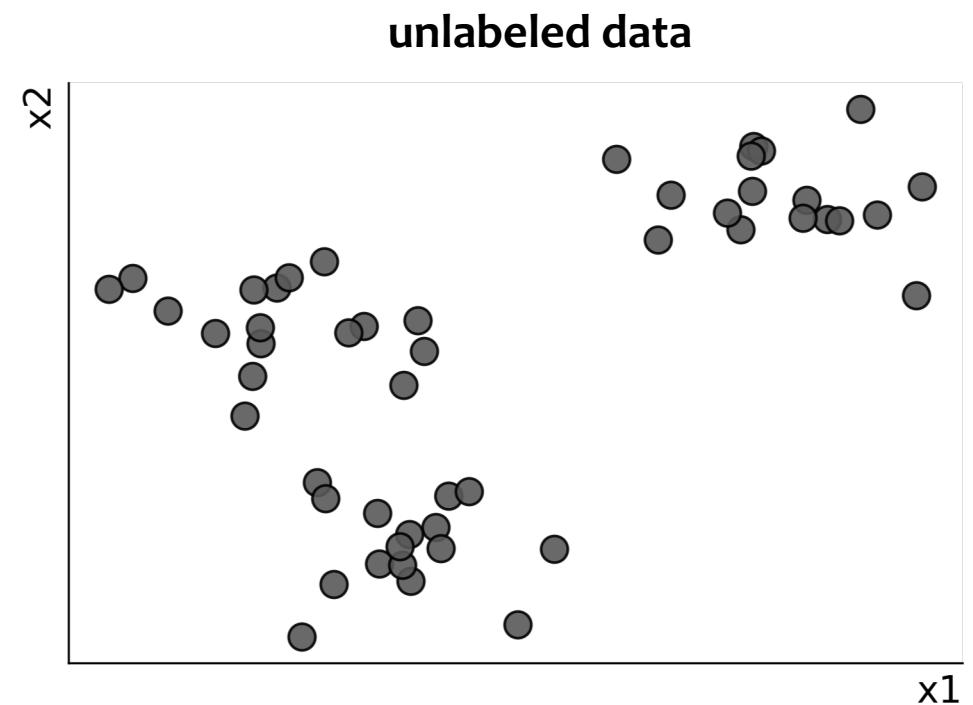
<https://research.rug.nl/en/publications/unsupervised-brain-anomaly-detection-in-mr-images>

Let's start with **Clustering!**

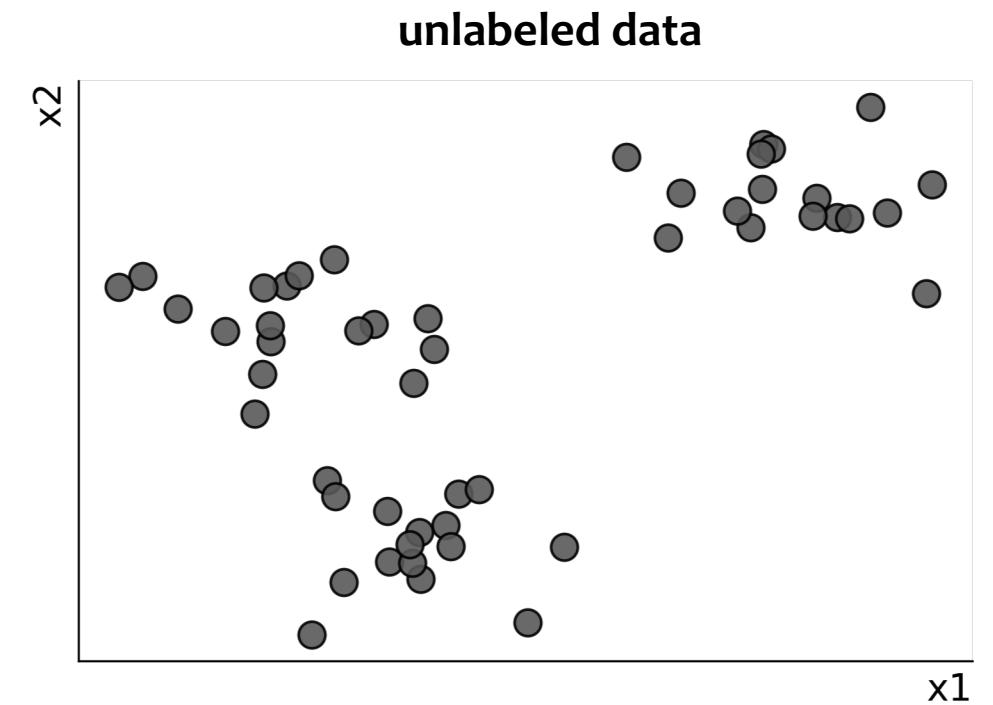
What if instead of having a set of previously labeled data...



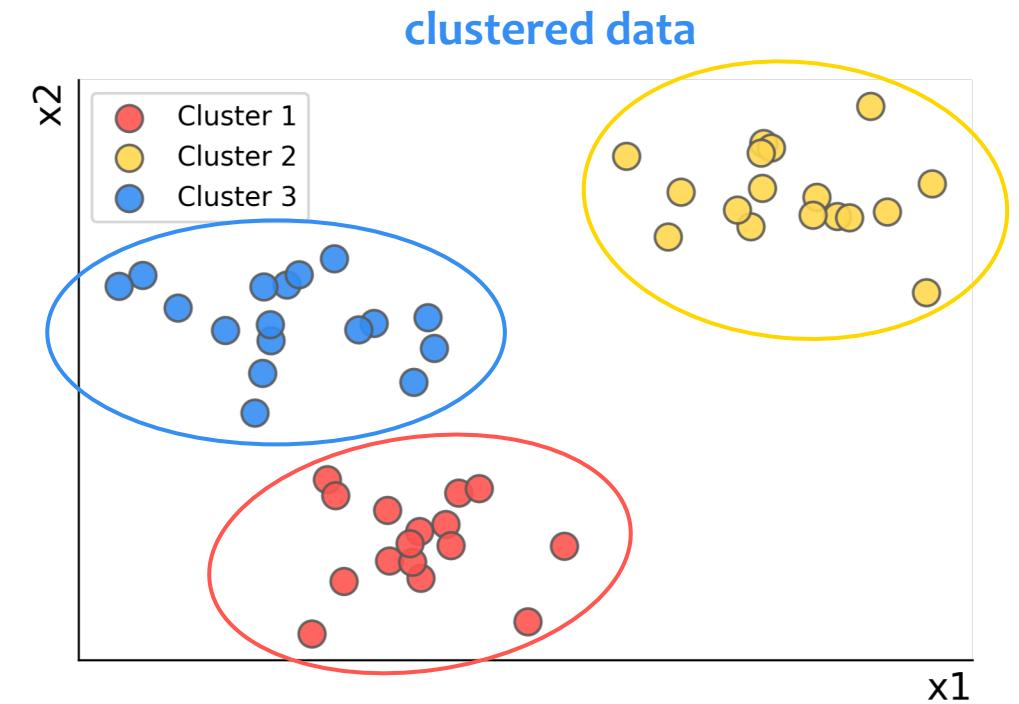
What if instead of having a set of previously labeled data, the **data were unlabeled**?



One could use some **clustering algorithm** to group similar instances together into *clusters/groups*.

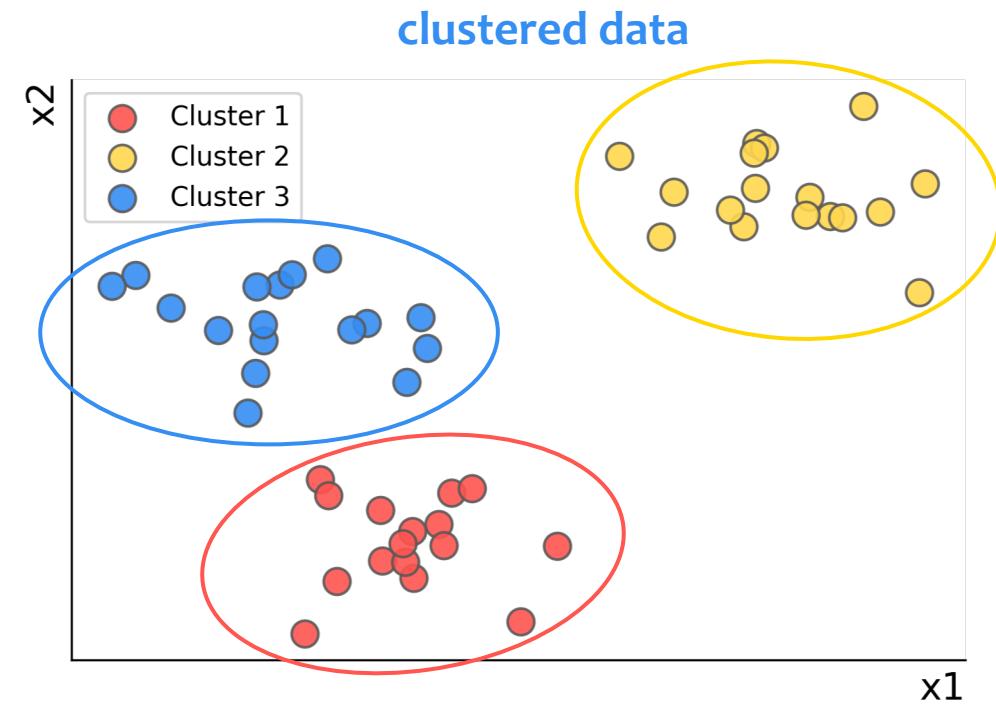


One could use some **clustering algorithm** to group similar instances together into *clusters/groups*.



One could use some **clustering algorithm** to group **similar instances** together into **clusters/groups**.

One may then **analyze each cluster** individually in order to find some specific pattern/behaviour/characteristic between their instances.

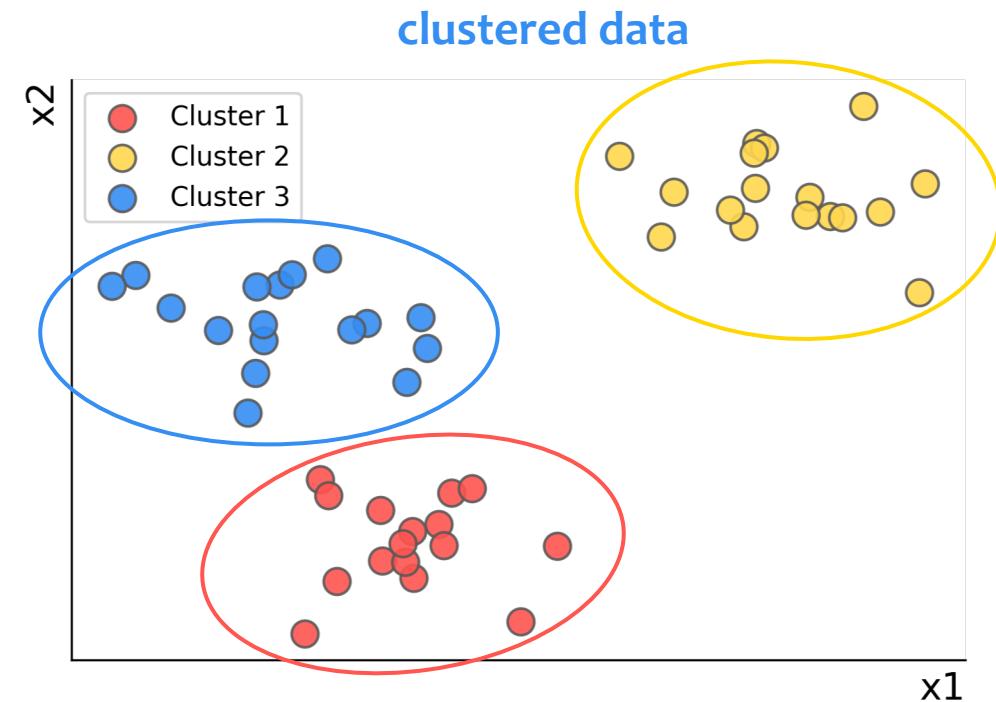


One could use some **clustering algorithm** to group **similar instances** together into **clusters/groups**.

One may then **analyze each cluster** individually in order to find some specific pattern/behaviour/characteristic between their instances.



This is often done during **exploratory data analysis**.



Clustering Approaches

Clustering Approaches



Partition methods

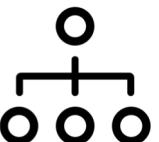
- Given a dataset with m samples, a partition method constructs **k partitions**, $k \leq m$, where each partition represents a **cluster**, satisfying:
 - Each **cluster** must contain at least one sample;
 - Each sample must belong to exactly one **cluster**;

Clustering Approaches



Partition methods

- Given a dataset with m samples, a partition method constructs **k partitions**, $k \leq m$, where each partition represents a **cluster**, satisfying:
 - Each **cluster** must contain at least one sample;
 - Each sample must belong to exactly one **cluster**;

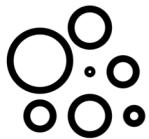


Hierarchical methods

- Creates a **hierarchical decomposition** of the given dataset (agglomerative or divisive).
- Hierarchical methods suffer from the fact that once a step (merge or split) is done, it can never be undone.

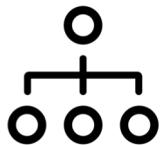
Clustering Approaches

Partition methods



- Given a dataset with m samples, a partition method constructs **k partitions**, $k \leq m$, where each partition represents a **cluster**, satisfying:
 - Each **cluster** must contain at least one sample;
 - Each sample must belong to exactly one **cluster**;

Hierarchical methods



- Creates a **hierarchical decomposition** of the given dataset (agglomerative or divisive).
- Hierarchical methods suffer from the fact that once a step (merge or split) is done, it can never be undone.

Density-based methods



- Clustering methods based on the notion of **density**;
- Their general idea is to continue growing the given **cluster** as long as the **density** (number of data points) in the neighborhood exceeds some threshold;
- They can be used to filter out noise (outliers) and discover **clusters** of arbitrary shape;

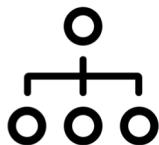
Clustering Approaches

Partition methods



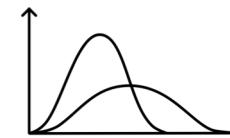
- Given a dataset with m samples, a partition method constructs **k partitions**, $k \leq m$, where each partition represents a **cluster**, satisfying:
 - Each **cluster** must contain at least one sample;
 - Each sample must belong to exactly one **cluster**;

Hierarchical methods



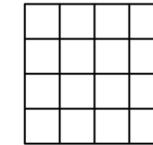
- Creates a **hierarchical decomposition** of the given dataset (agglomerative or divisive).
- Hierarchical methods suffer from the fact that once a step (merge or split) is done, it can never be undone.

Density-based methods



- Clustering methods based on the notion of **density**;
- Their general idea is to continue growing the given **cluster** as long as the **density** (number of data points) in the neighborhood exceeds some threshold;
- They can be used to filter out noise (outliers) and discover **clusters** of arbitrary shape;

Grid-based methods

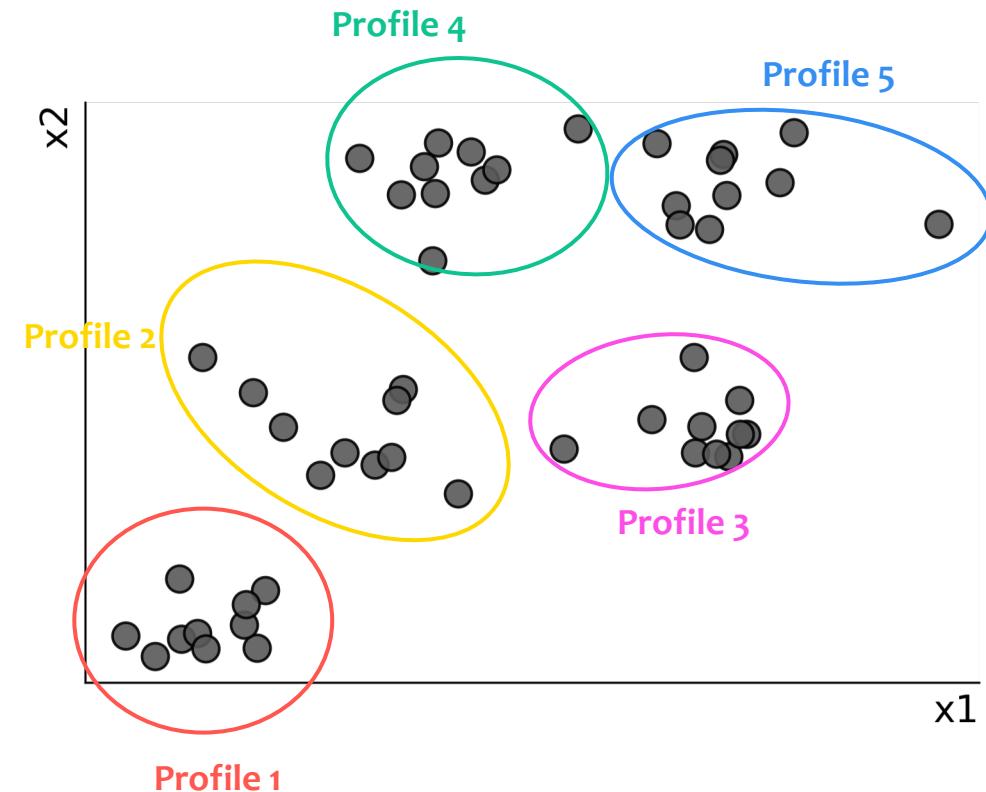


- First divide the space into regular grids, then perform **clustering** on the grids;

Clustering is used in several applications, including:

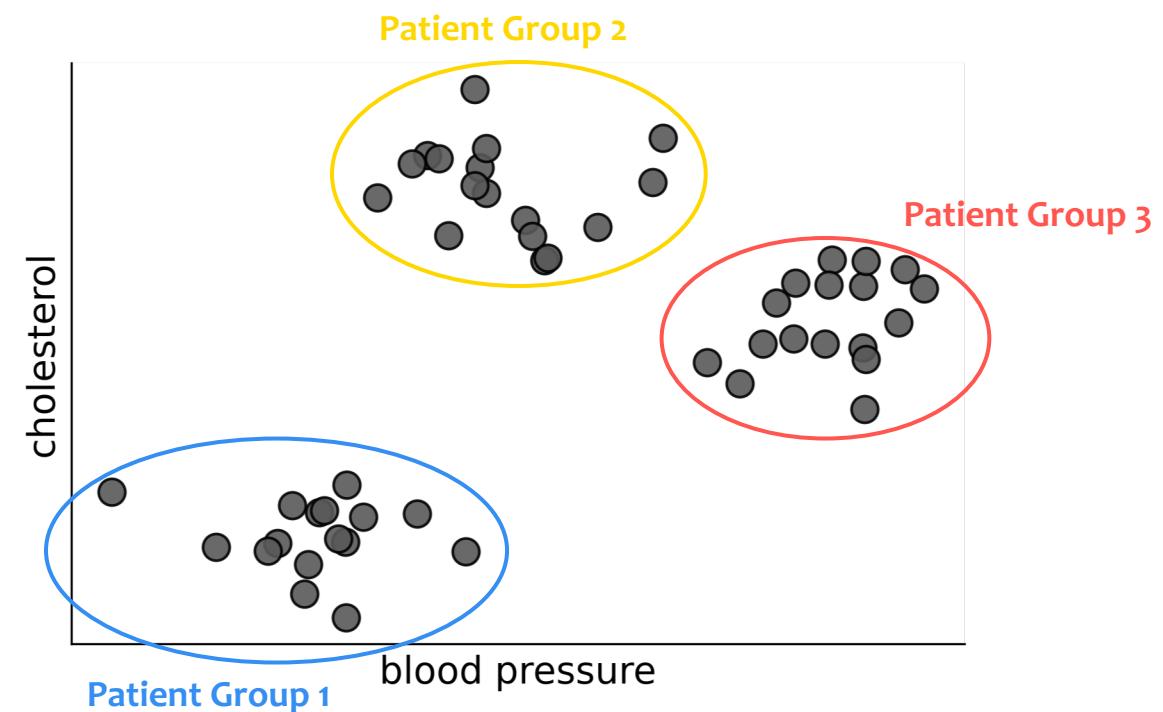
Clustering is used in several applications, including:

- Customer segmentation / recommender systems;



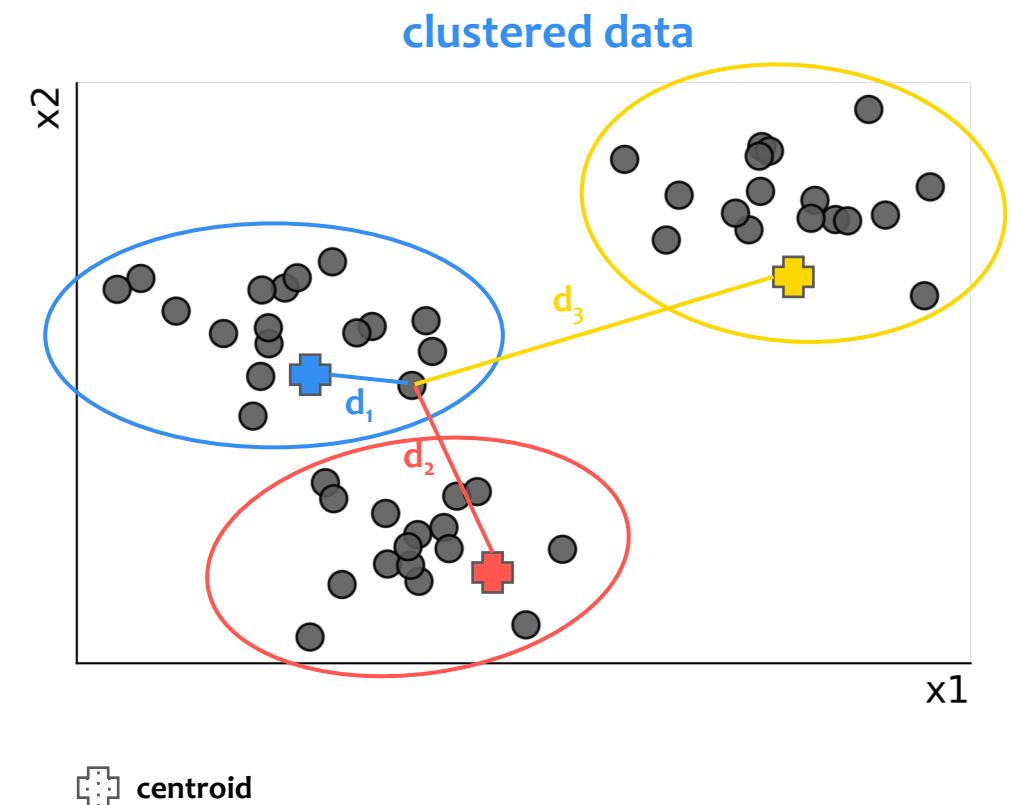
Clustering is used in several applications, including:

- Customer segmentation / recommender systems;
- Data Analysis;



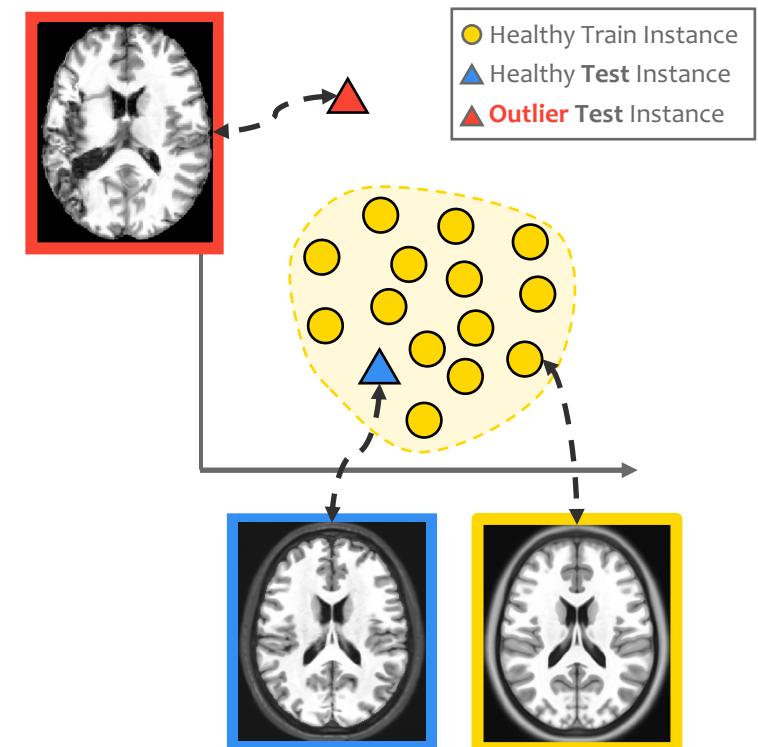
Clustering is used in several applications, including:

- Customer segmentation / recommender systems;
- Data Analysis;
- Dimensionality reduction technique;



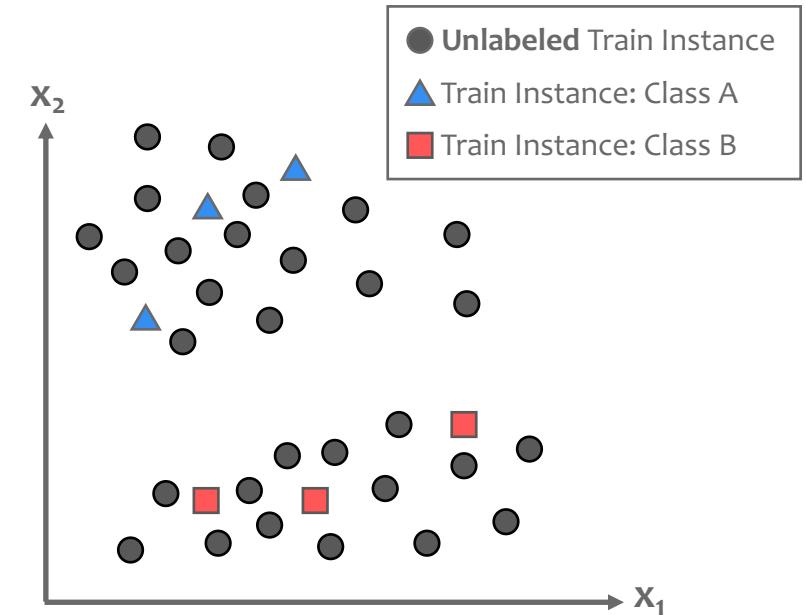
Clustering is used in several applications, including:

- Customer segmentation / recommender systems;
- Data Analysis;
- Dimensionality reduction technique;
- Anomaly detection;



Clustering is used in several applications, including:

- Customer segmentation / recommender systems;
- Data Analysis;
- Dimensionality reduction technique;
- Anomaly detection;
- Semi-supervised learning;



Clustering is used in several applications, including:

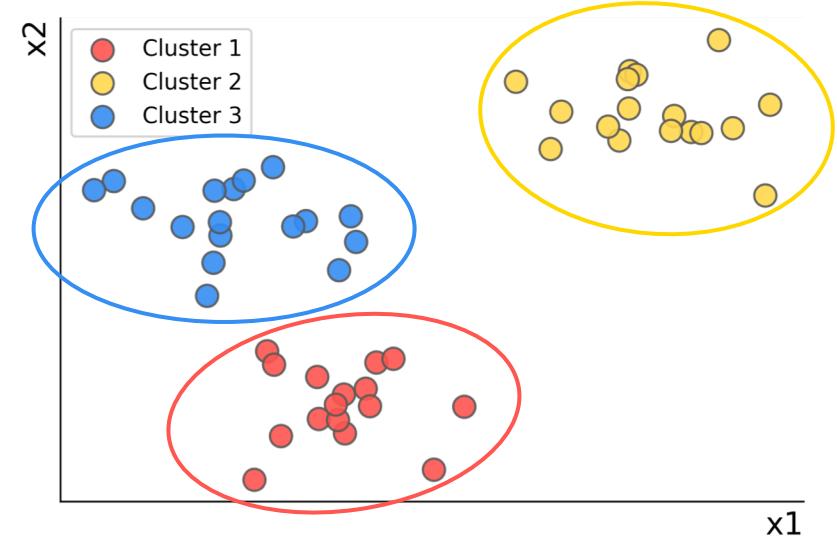
- Customer segmentation / recommender systems;
- Data Analysis;
- Dimensionality reduction technique;
- Anomaly detection;
- Semi-supervised learning;
- others...

Clustering is used in several applications, including:

- Customer segmentation / recommender systems;
- Data Analysis;
- Dimensionality reduction technique;
- Anomaly detection;
- Semi-supervised learning;
- others...

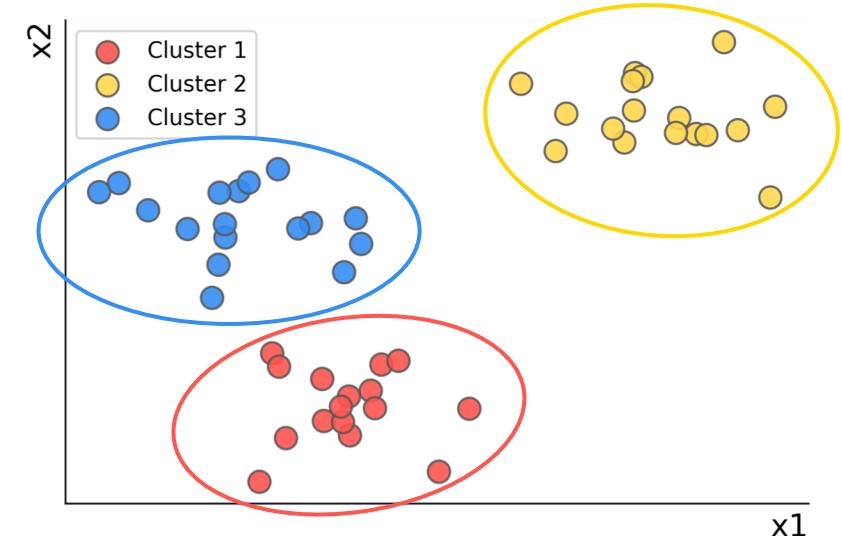
Let's see **KMeans**: the most
classical clustering algorithm!

KMeans



KMeans

- Proposed by Stuart Lloyd at Bell Labs in 1957 but only published in 1982;
- Sometimes referred to as Lloyd–Forgy.

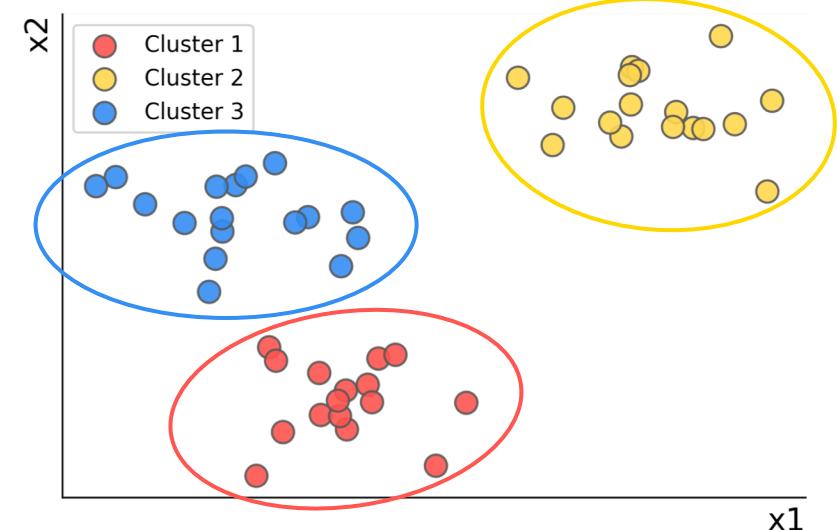


KMeans

- Proposed by Stuart Lloyd at Bell Labs in 1957 but only published in 1982;
- Sometimes referred to as Lloyd–Forgy.

Goal

- **To partition** the dataset into k **clusters**:
 - Each **cluster** must contain **at least one** sample;
 - Each sample must belong to **exactly one** **cluster**;
 - The **mass center (mean)** of a **cluster** corresponds to its **centroid**;



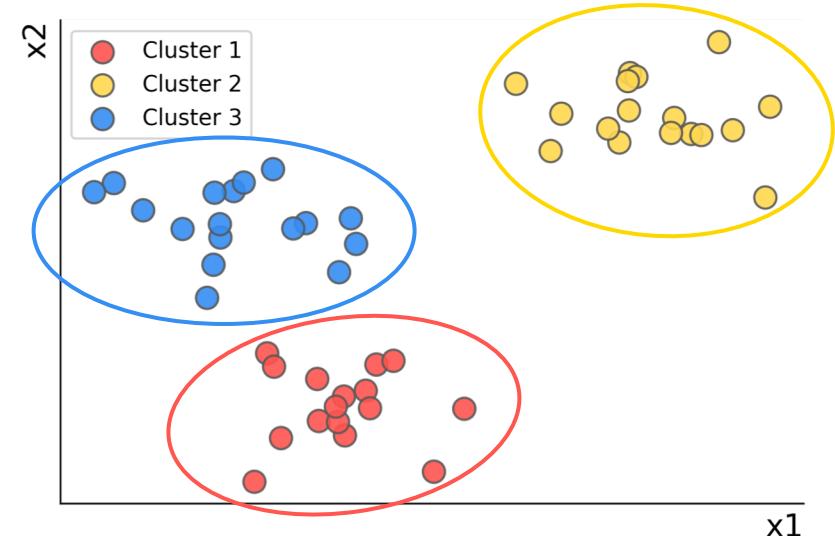
KMeans

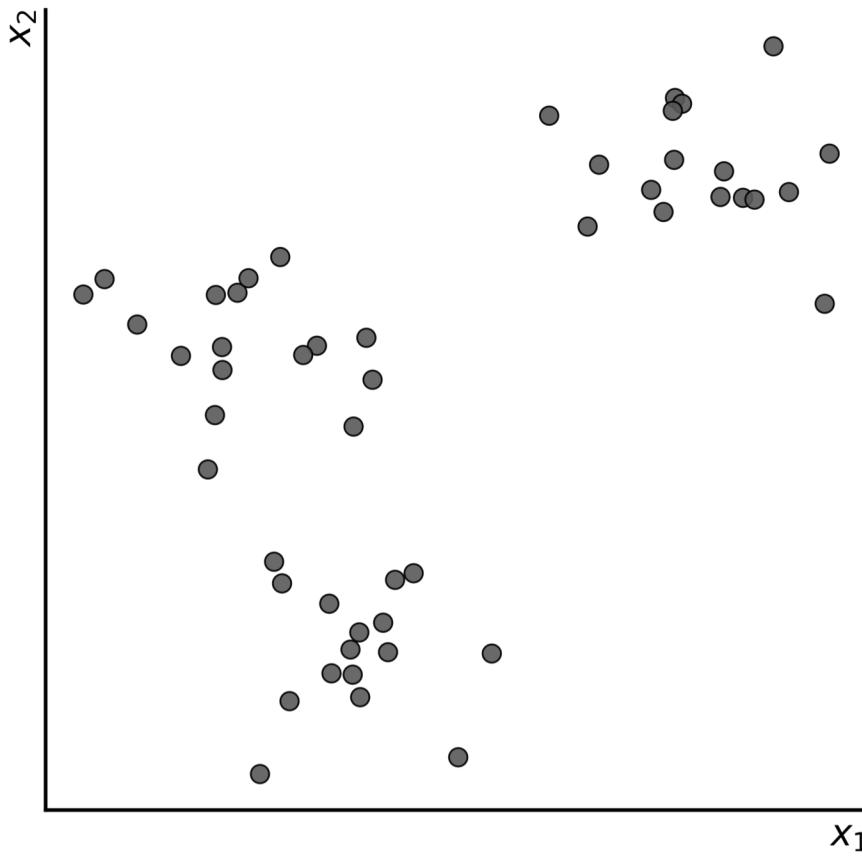
Initialization:

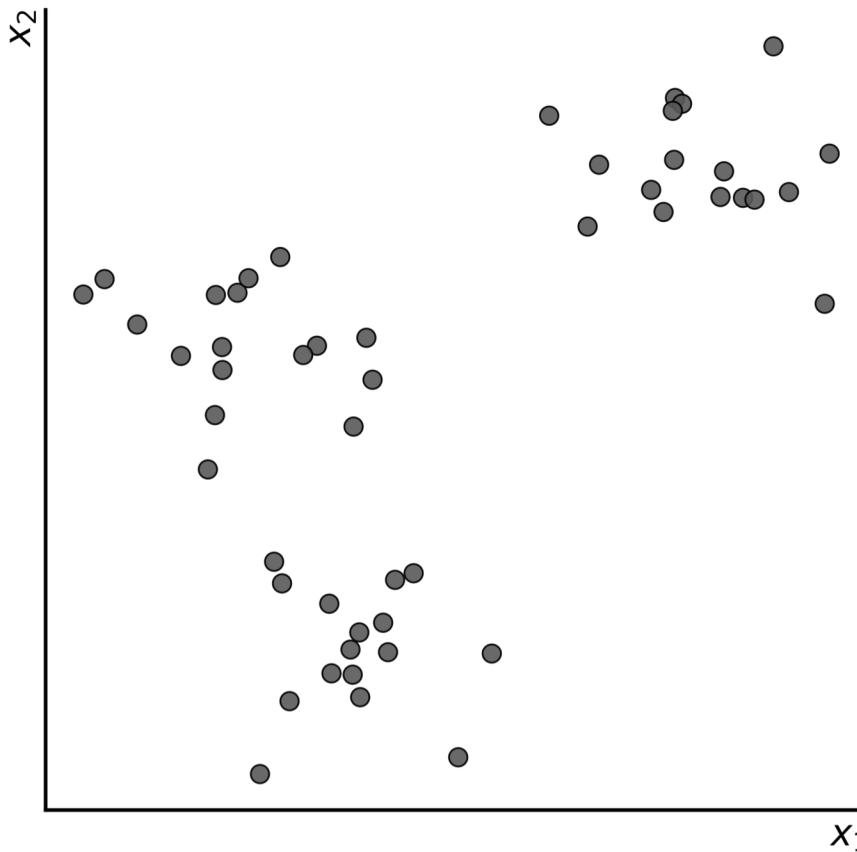
- 1) Choose the **number k** of **clusters**;
- 2) Choose the **initial k centroids** (cluster centers);

Learning:

- 3) Repeat until **convergence**:
 - 4) Assign each instance to **the nearest** centroid;
 - 5) Move the centroids to the centers of their **clusters**;
- 6) Assign each instance to **the final nearest** centroid;

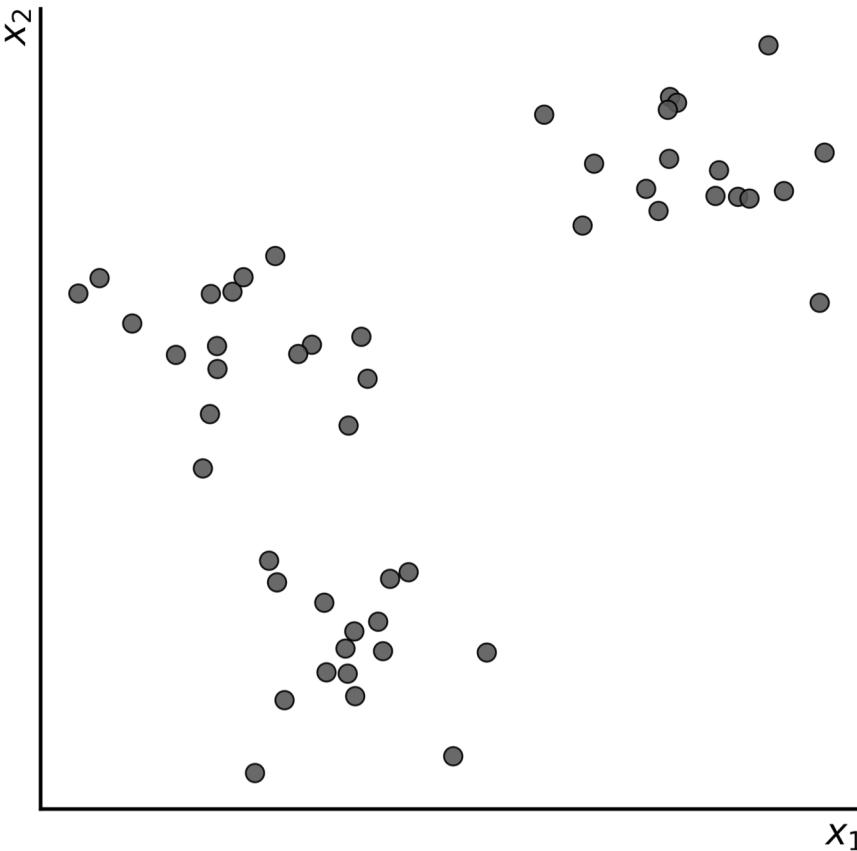






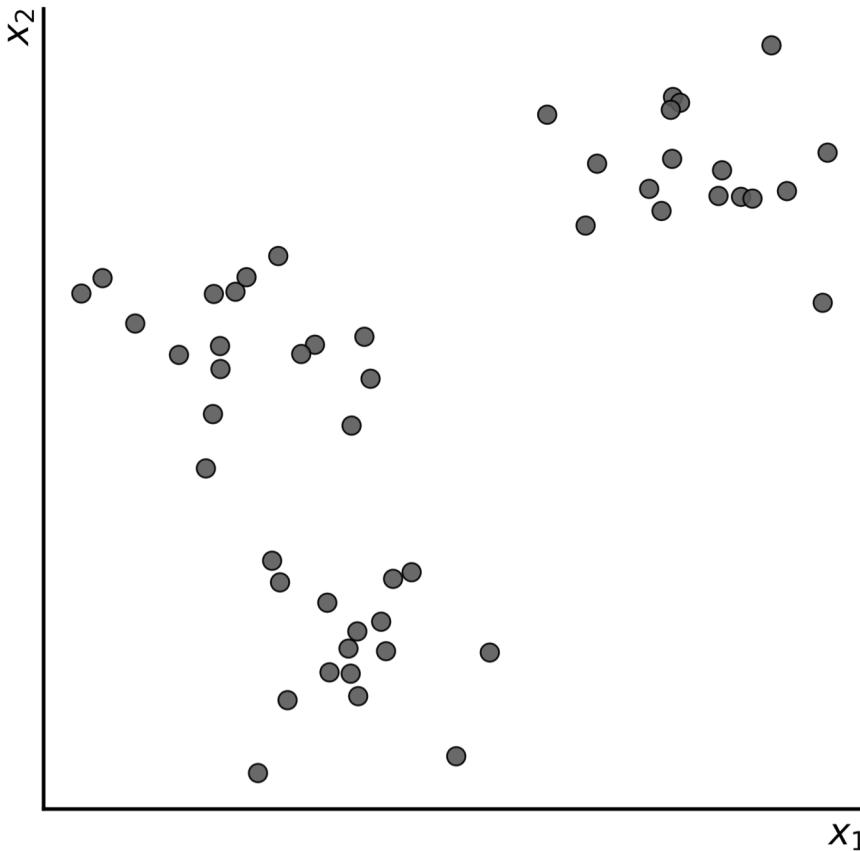
Consider that x_1 and x_2 have the **same scale**.

1) Choose the **number k** of clusters;



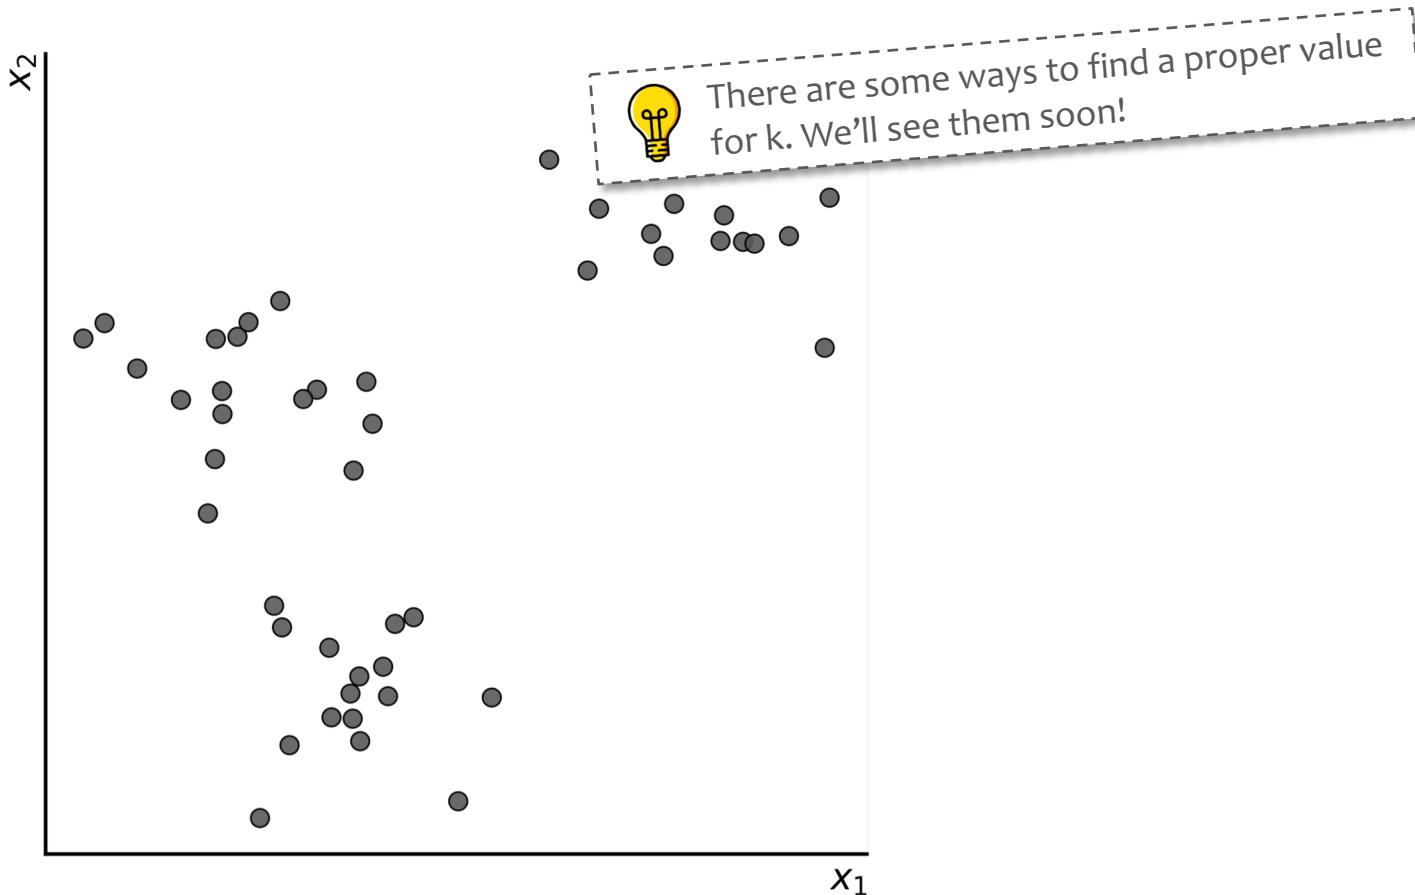
1) Choose the **number k** of clusters;

Let's consider: **k=3**

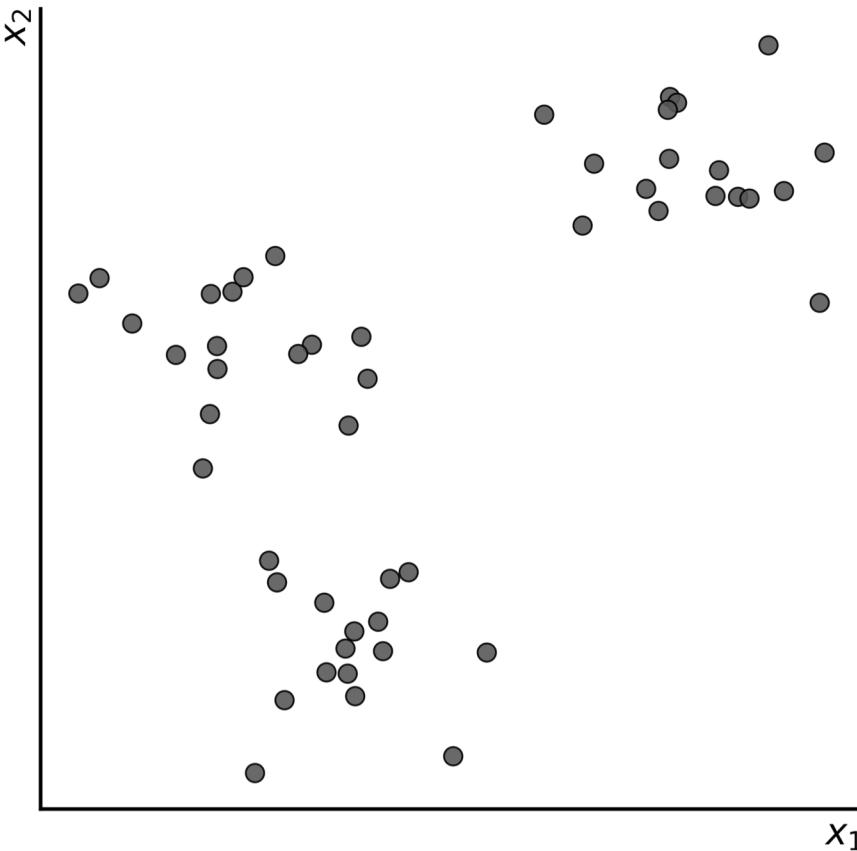


1) Choose the number k of clusters;

Let's consider: $k=3$

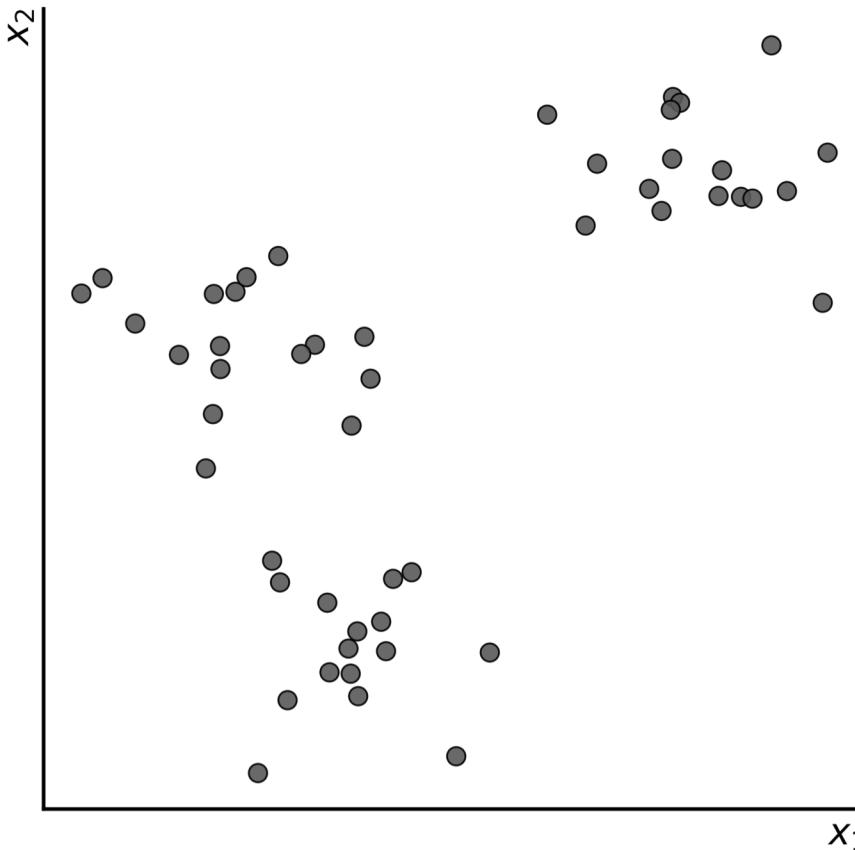


2) Choose the **initial k centroids** (cluster centers)



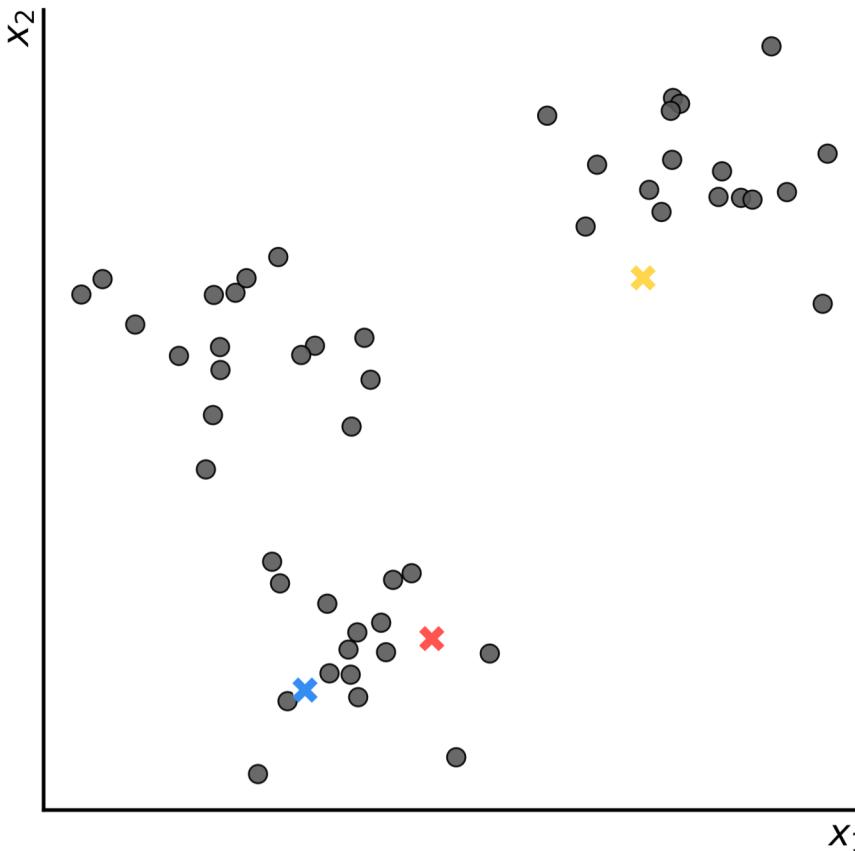
2) Choose the **initial k centroids** (cluster centers)

- random instances from the dataset (Forgy initialization); or
- random positions in the input feature space; or
- pre-defined k positions; or
- positions resulting from some optimization technique;

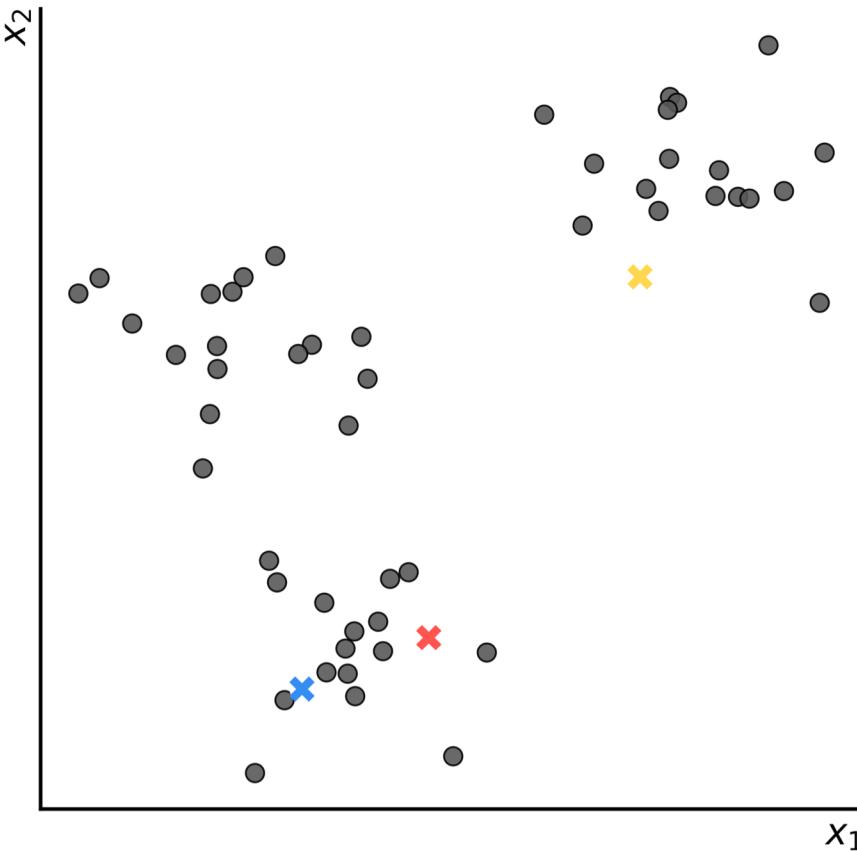


2) Choose the **initial k centroids** (cluster centers)

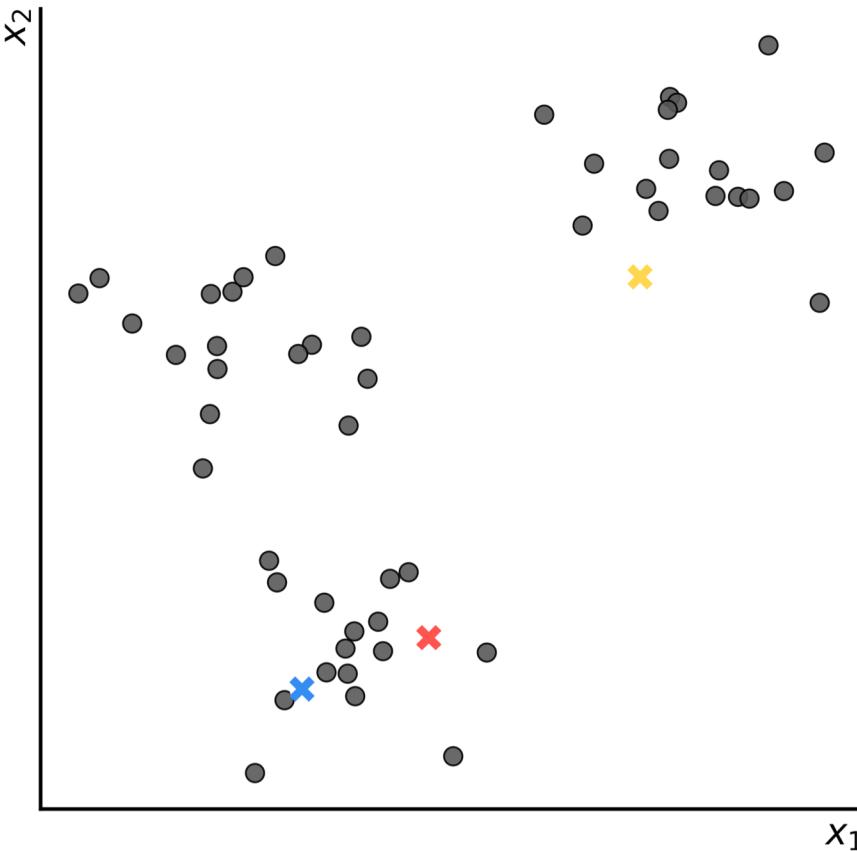
- random instances from the dataset (Forgy initialization); or
- **random positions in the input feature space; or**
- pre-defined k positions; or
- positions resulting from some optimization technique;



2) Choose the **initial k centroids** (cluster centers)

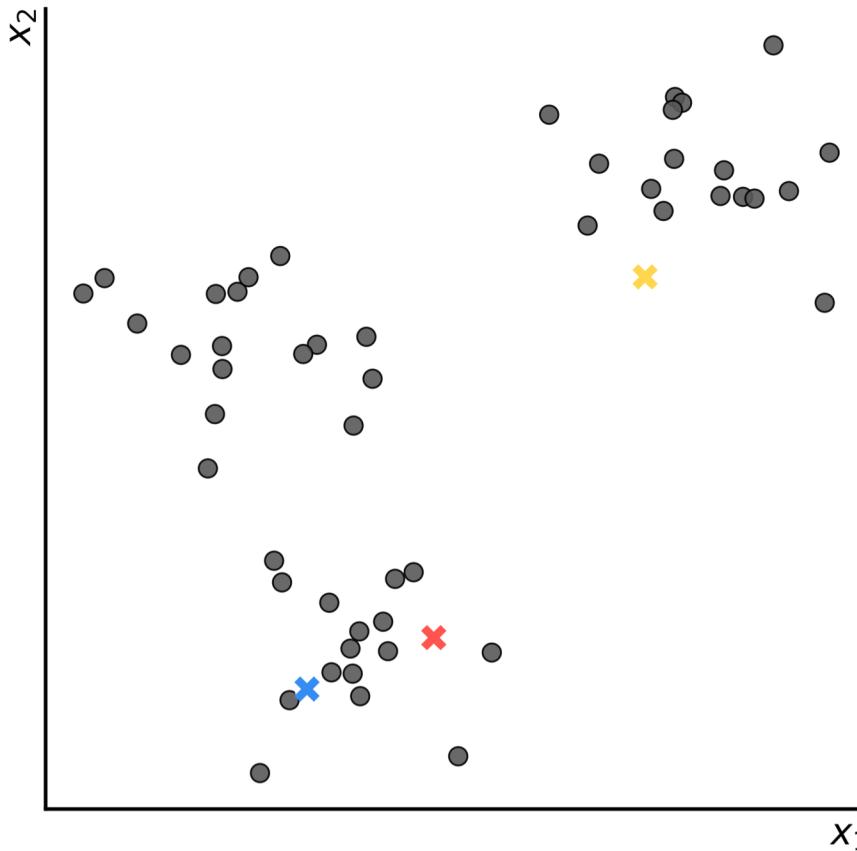


3) Repeat until **convergence**:



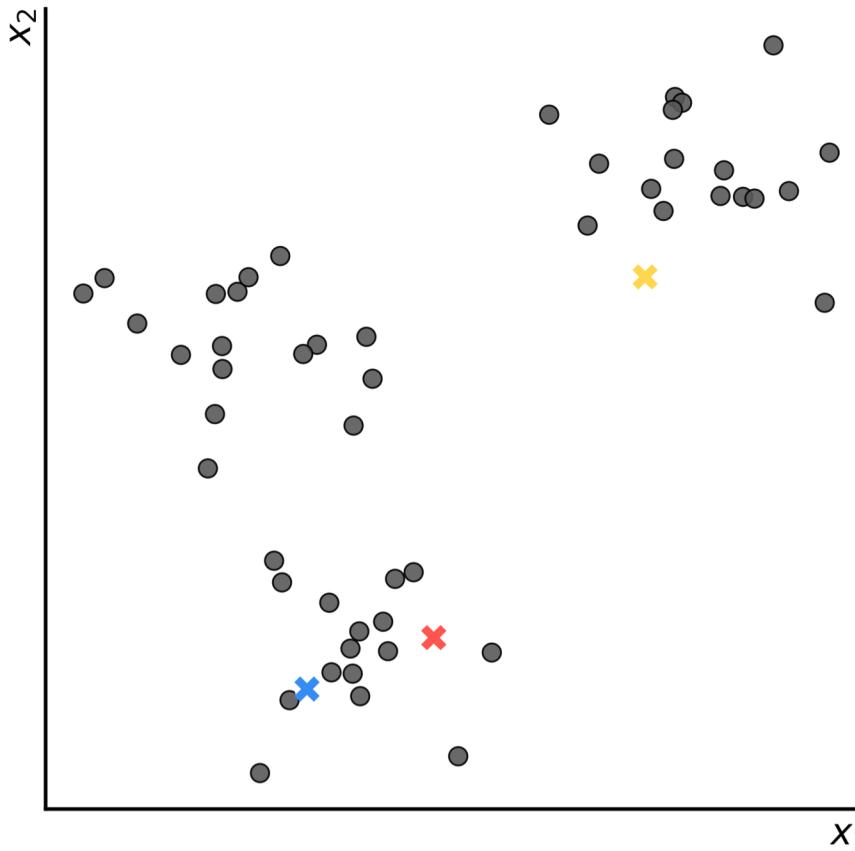
3) Repeat until **convergence**:

- No reassignment occurs; and/or
- Centroids do not move significantly; and/or
- Maximum number of iterations reached;

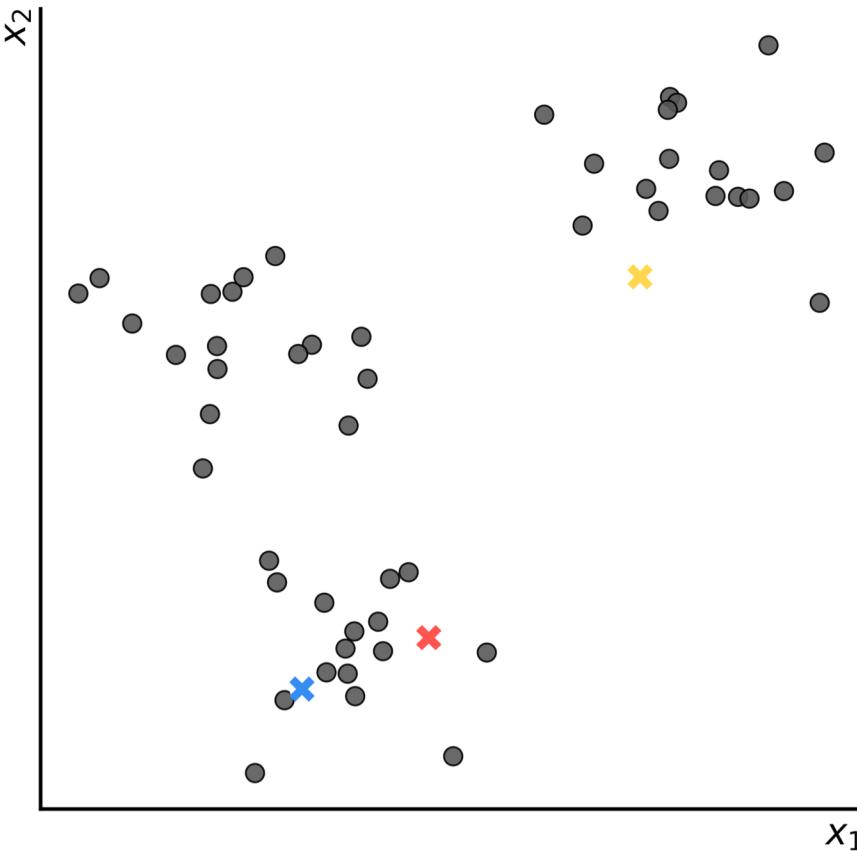


3) Repeat until **convergence**:

- No reassignment occurs; and/or
- **Centroids do not move significantly**; and/or
- Maximum number of iterations reached;



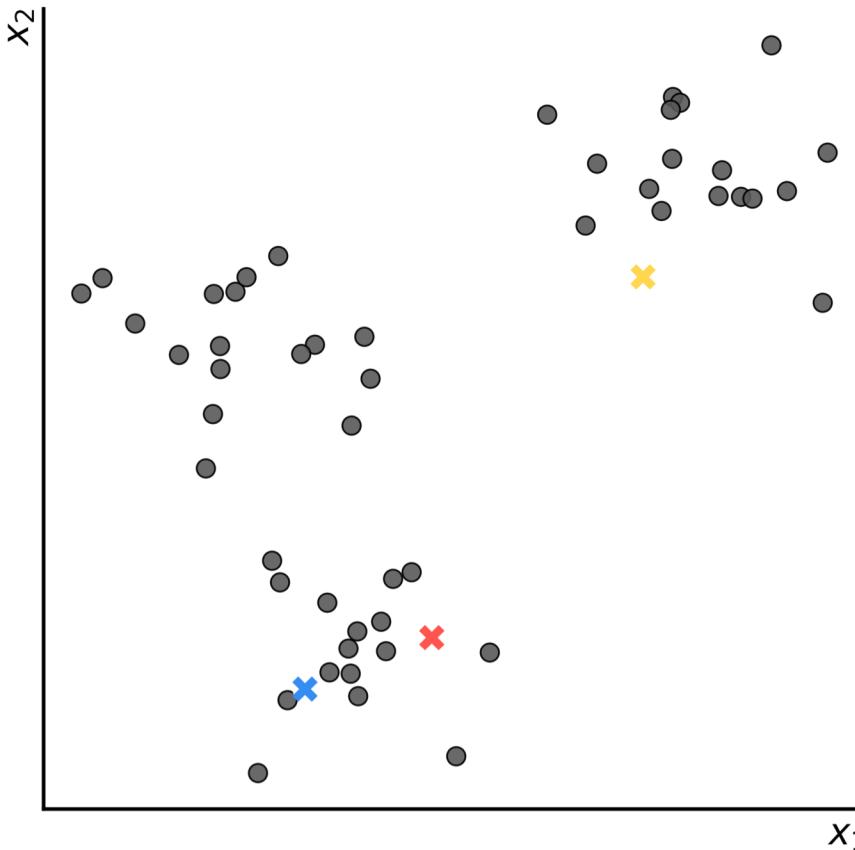
3) Repeat until **centroids do not move significantly**:



3) Repeat until **centroids do not move significantly**:

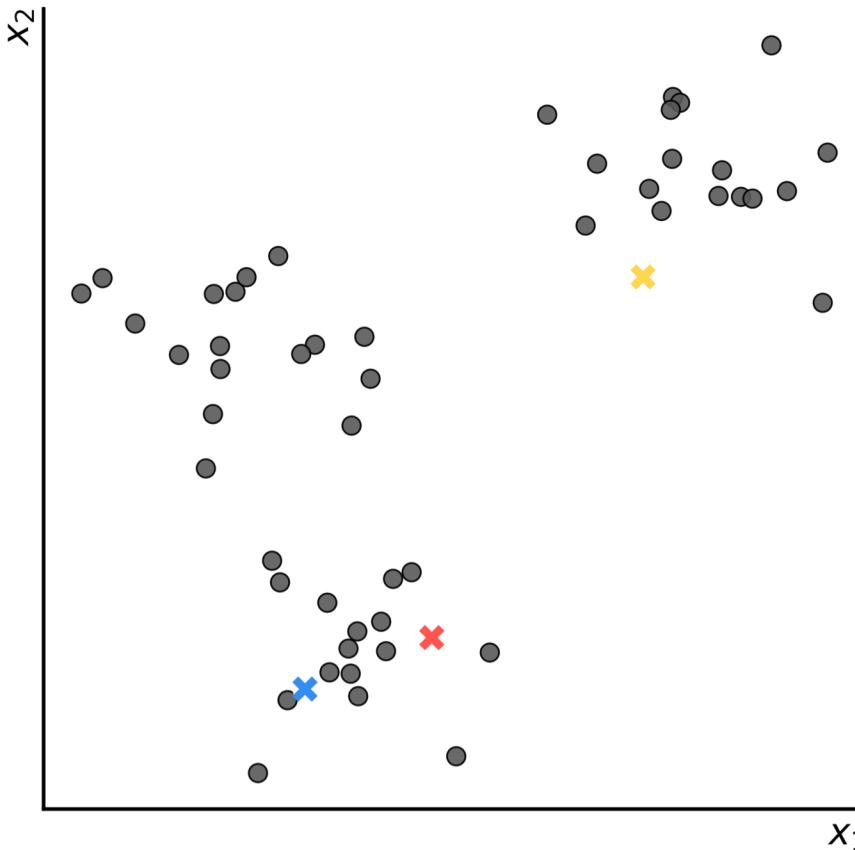
iteration 1

- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



4) Assign each instance to **the nearest centroid**:

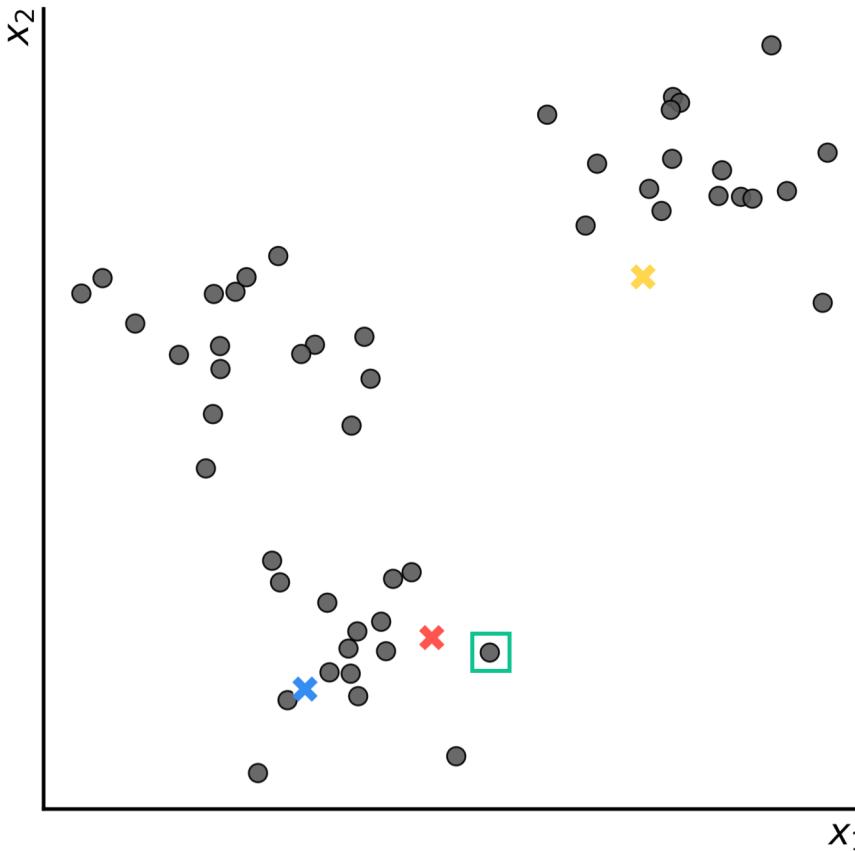
iteration 1



4) Assign each instance to **the nearest** centroid:

iteration 1

- For each sample $x^{(i)}$:

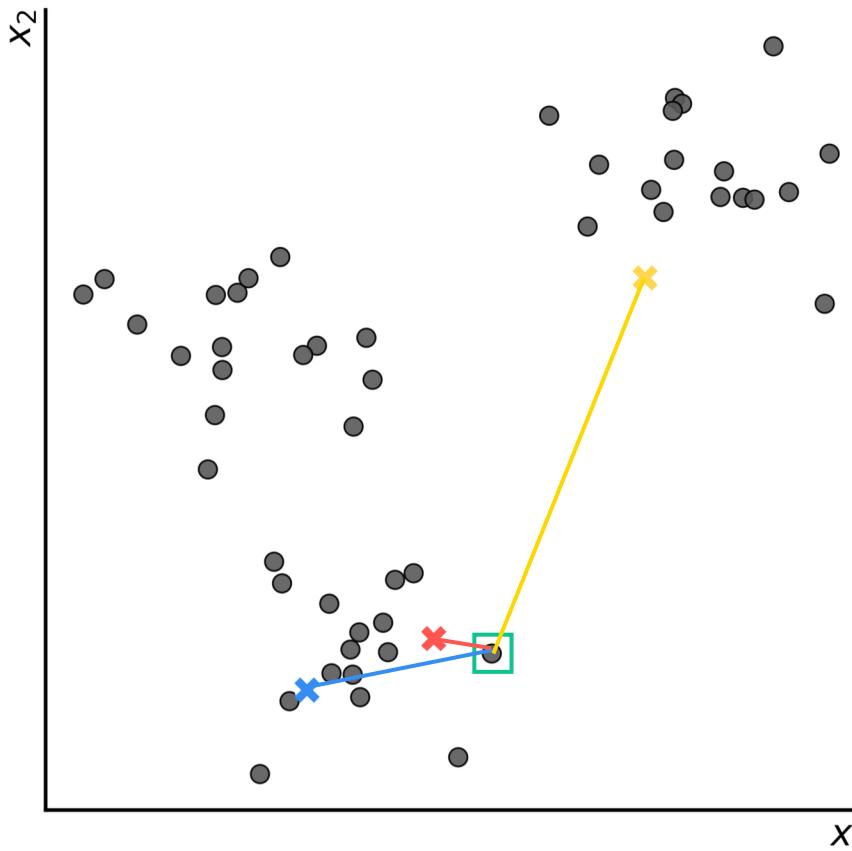


4) Assign each instance to **the nearest centroid**:

- For each sample $x^{(i)}$:
 - Compute the distance from $x^{(i)}$ to each centroid $c^{(j)}$;

 **Euclidean distance** is the common metric (default) used in KMeans.

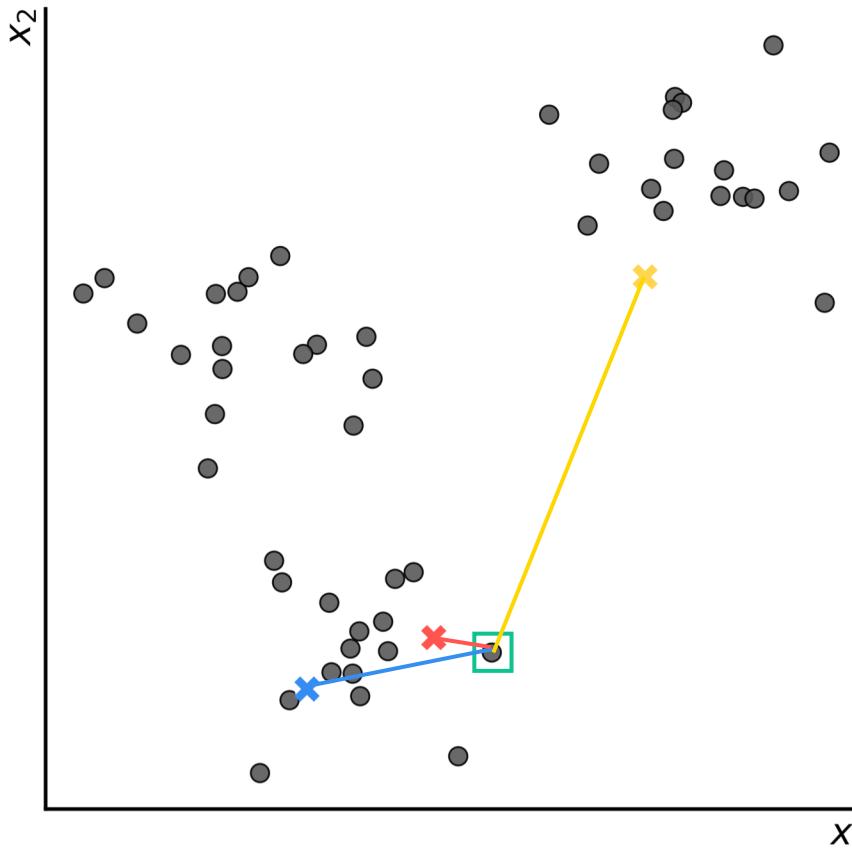
iteration 1



4) Assign each instance to **the nearest** centroid:

iteration 1

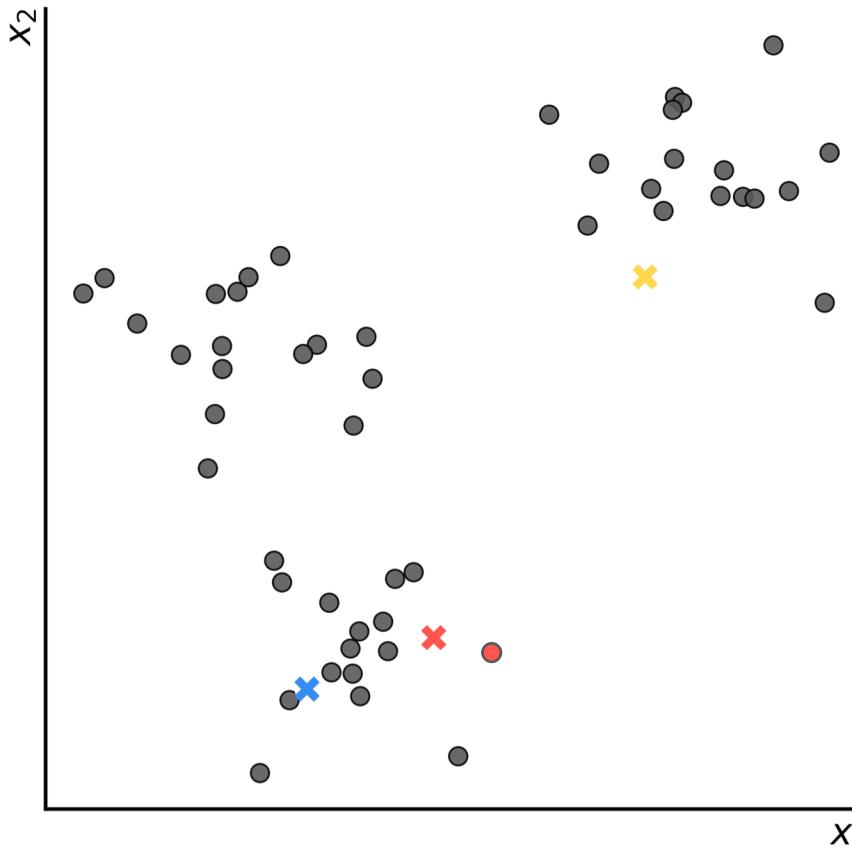
- For each sample $x^{(i)}$:
 - Compute the distance from $x^{(i)}$ to each centroid $c^{(j)}$;
 - Assign the sample to the nearest cluster center with distance $d^{(i)} = \min_j d(x^{(i)}, c^{(j)})$



4) Assign each instance to **the nearest** centroid:

iteration 1

- For each sample $x^{(i)}$:
 - Compute the distance from $x^{(i)}$ to each centroid $c^{(j)}$;
 - Assign the sample to the nearest cluster center with distance $d^{(i)} = \min_j d(x^{(i)}, c^{(j)})$

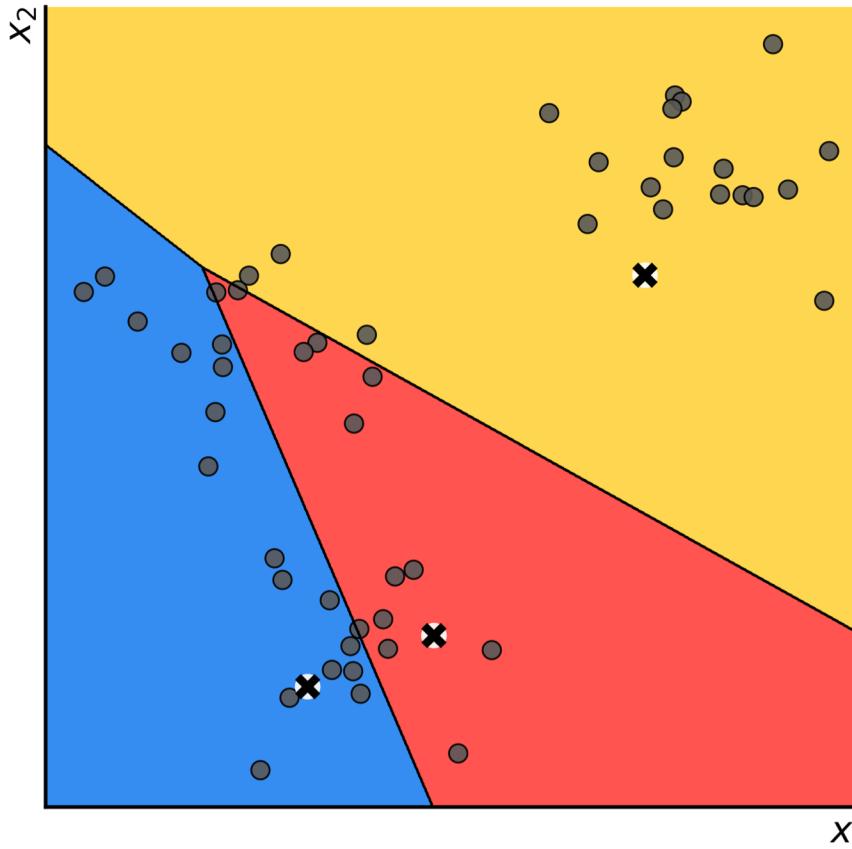


4) Assign each instance to **the nearest** centroid:

iteration 1

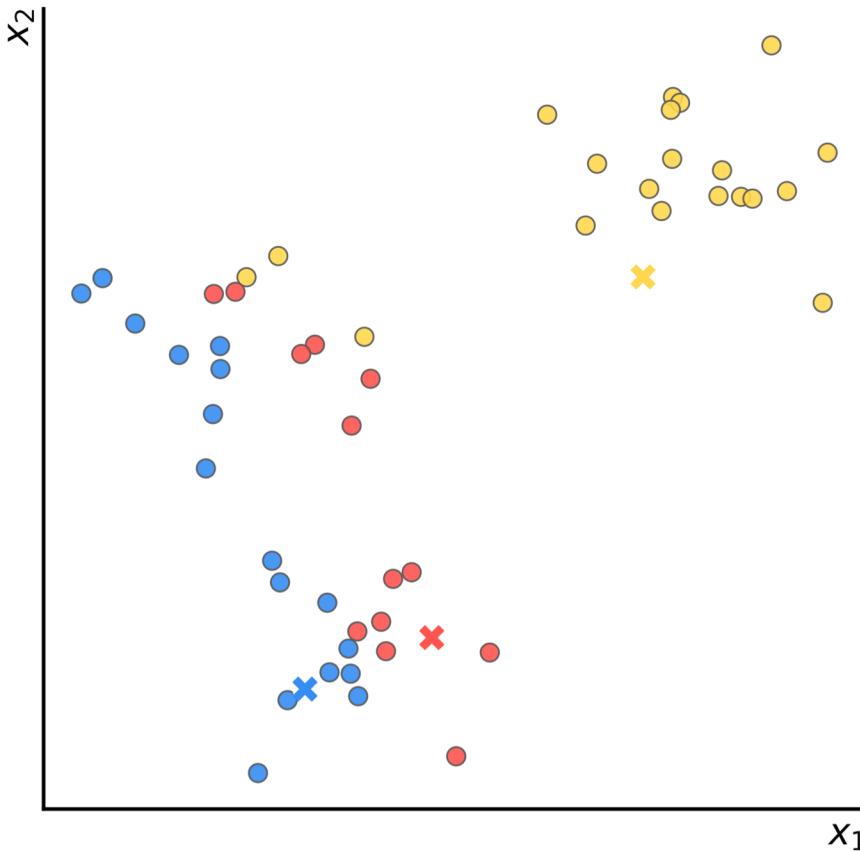
- For each sample $x^{(i)}$:
 - Compute the distance from $x^{(i)}$ to each centroid $c^{(j)}$;
 - Assign the sample to the nearest cluster center with distance $d^{(i)} = \min_j d(x^{(i)}, c^{(j)})$

 K-Means **decision boundaries**
(Voronoi diagram/tessellation)



5) Move the centroids to the centers of their clusters

iteration 1



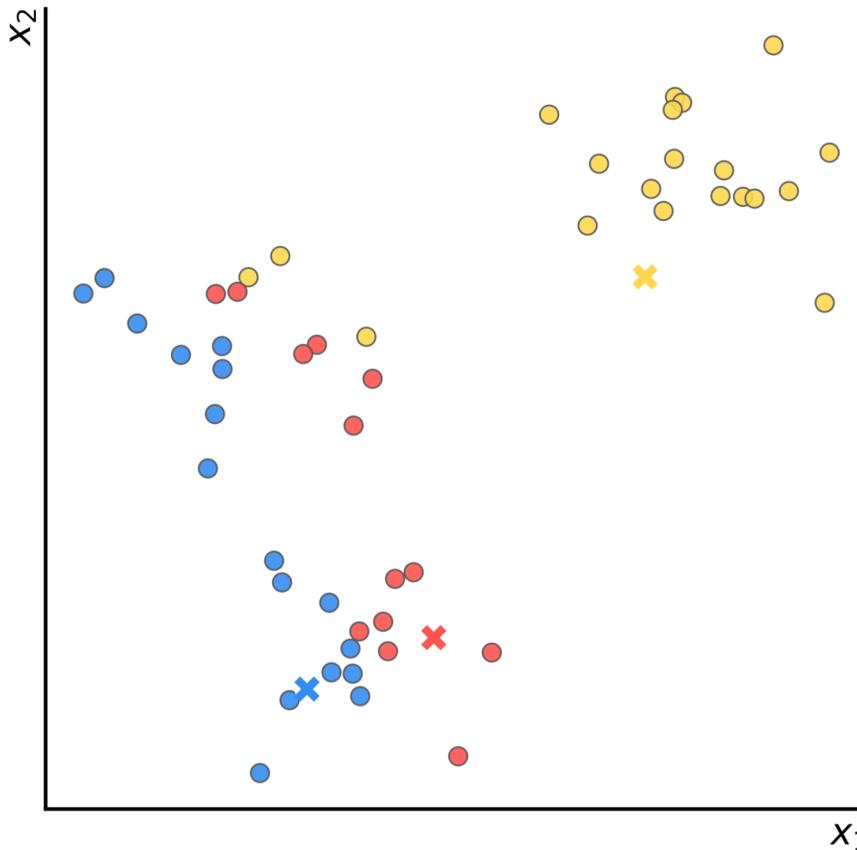
5) Move the centroids to the centers of their clusters

iteration 1

- For each cluster $c^{(j)} = [c_1^{(j)}, c_2^{(j)}, \dots]$:

- $c^{(j)} = \frac{1}{m^{(j)}} \sum_{i=1}^{m^{(j)}} x^{(j)}$

number of samples
in **cluster j**



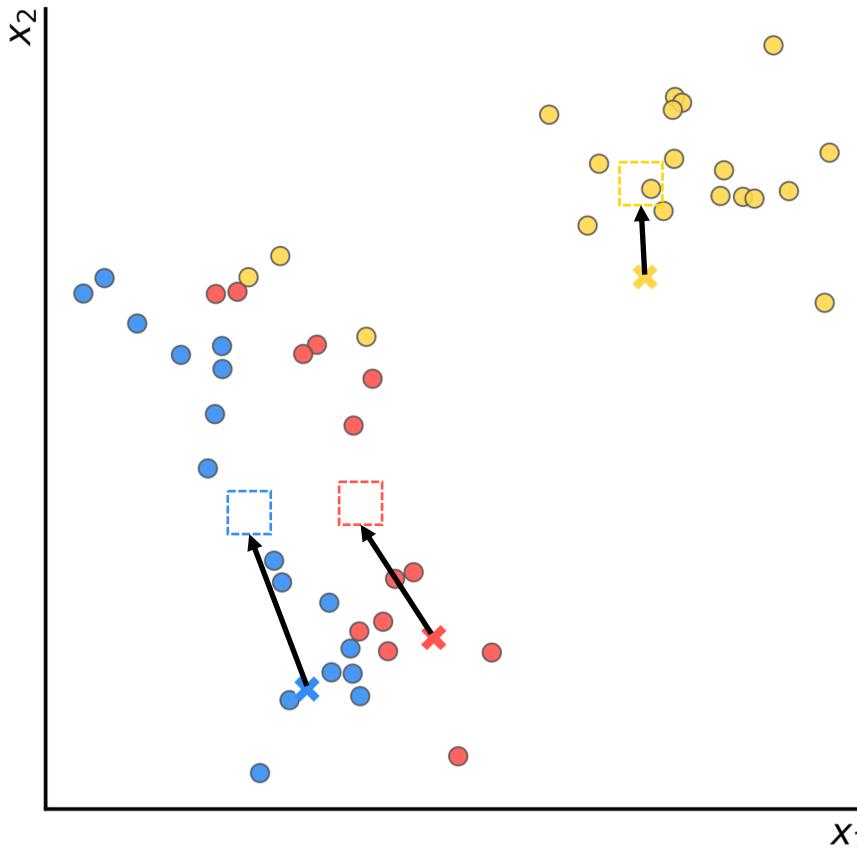
5) Move the centroids to the centers of their clusters

iteration 1

- For each cluster $c^{(j)} = [c_1^{(j)}, c_2^{(j)}, \dots]$:

- $c^{(j)} = \frac{1}{m^{(j)}} \sum_{i=1}^{m^{(j)}} x^{(j)}$

number of samples
in **cluster j**



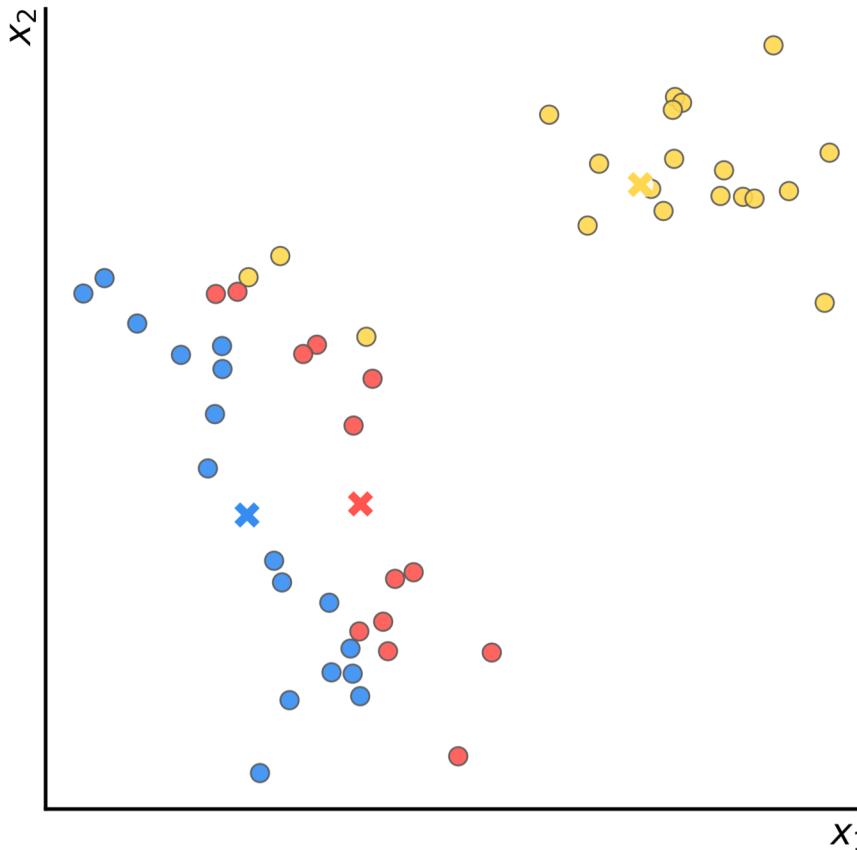
5) Move the centroids to the centers of their clusters

iteration 1

- For each cluster $c^{(j)} = [c_1^{(j)}, c_2^{(j)}, \dots]$:

- $c^{(j)} = \frac{1}{m^{(j)}} \sum_{i=1}^{m^{(j)}} x^{(j)}$

number of samples
in **cluster j**



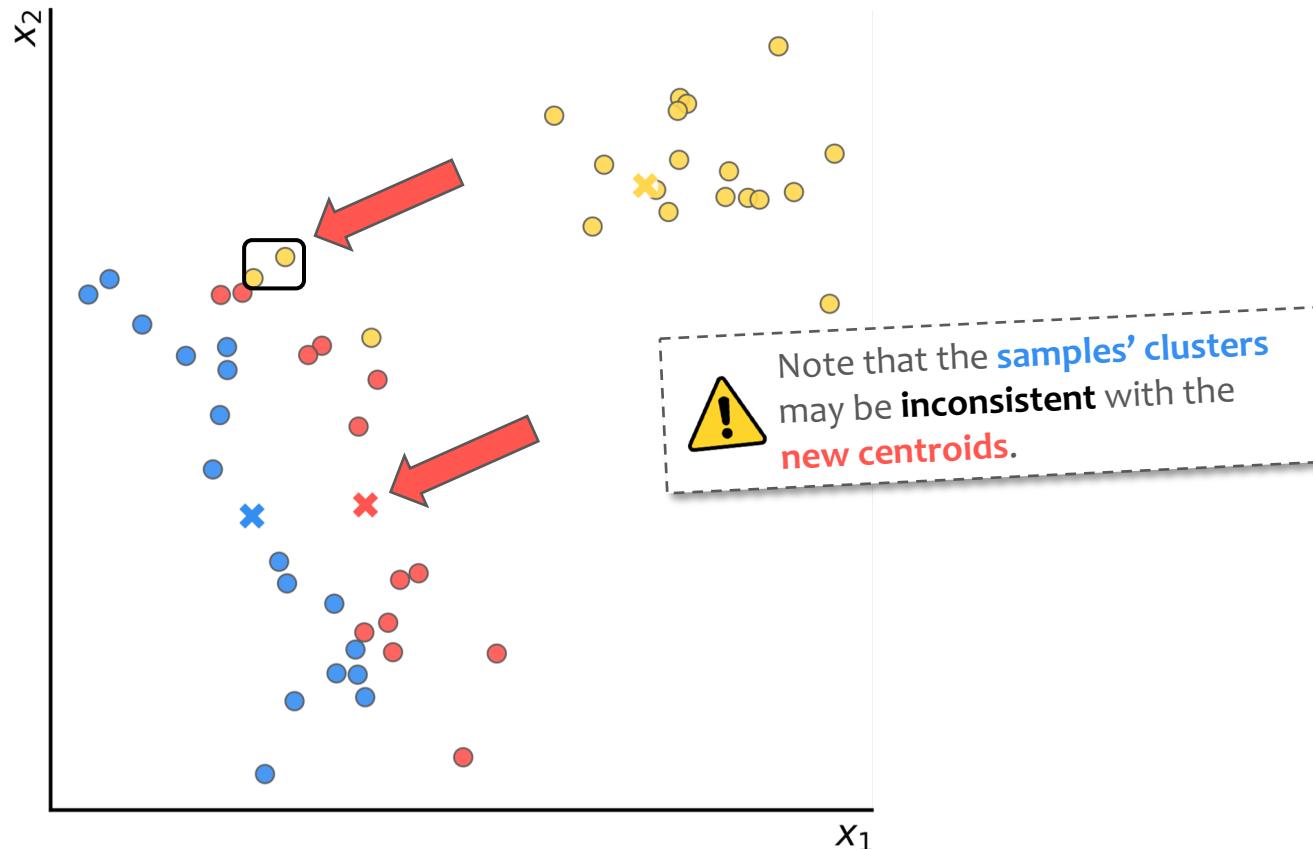
5) Move the centroids to the centers of their clusters

iteration 1

- For each cluster $c^{(j)} = [c_1^{(j)}, c_2^{(j)}, \dots]$:

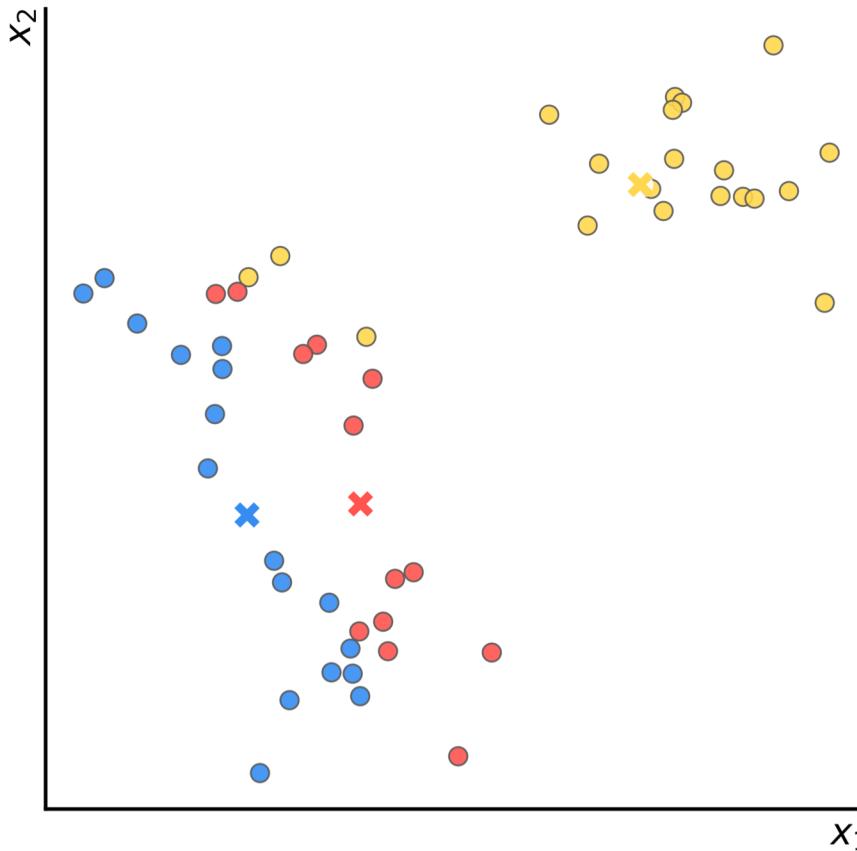
- $c^{(j)} = \frac{1}{m^{(j)}} \sum_{i=1}^{m^{(j)}} x^{(j)}$

number of samples
in **cluster j**



3) Repeat until **centroids do not move significantly**:

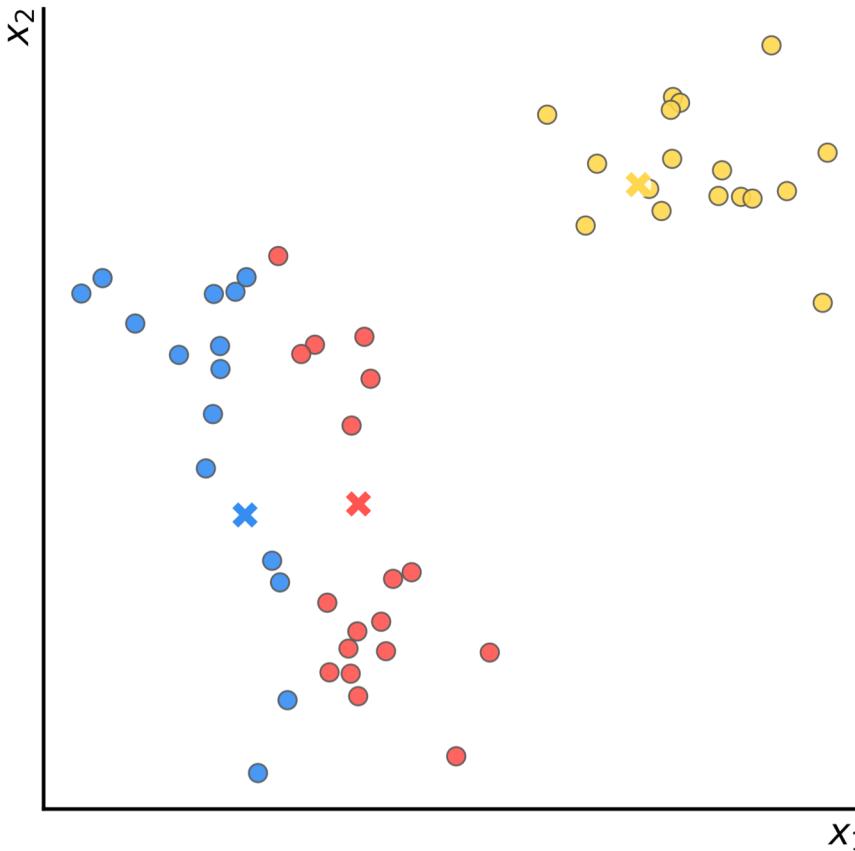
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 2

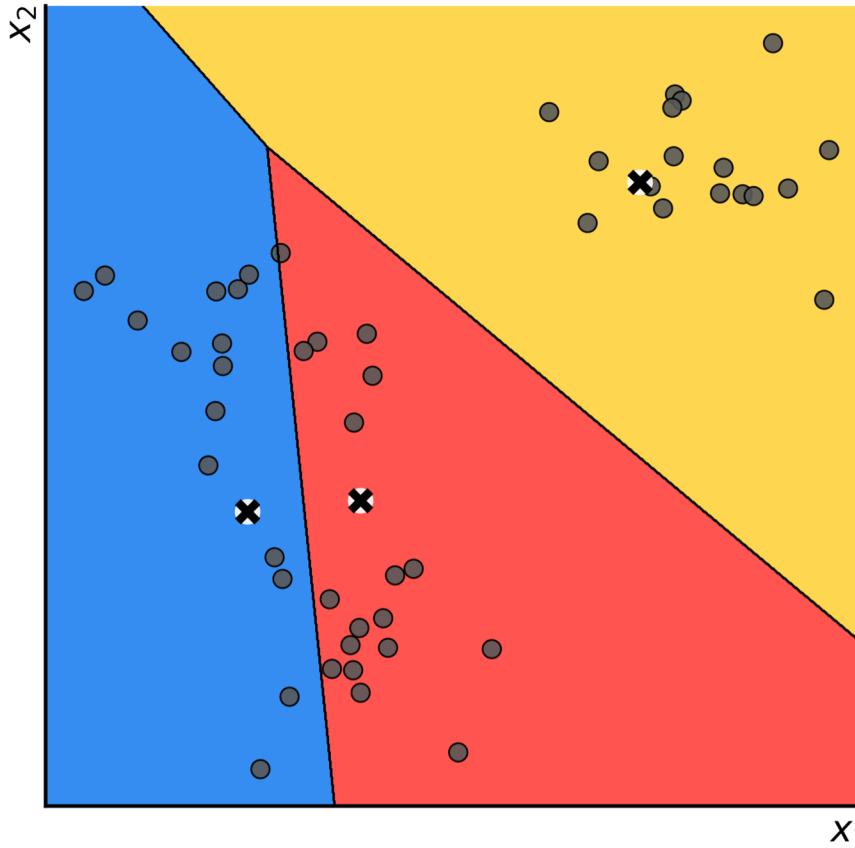
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 2

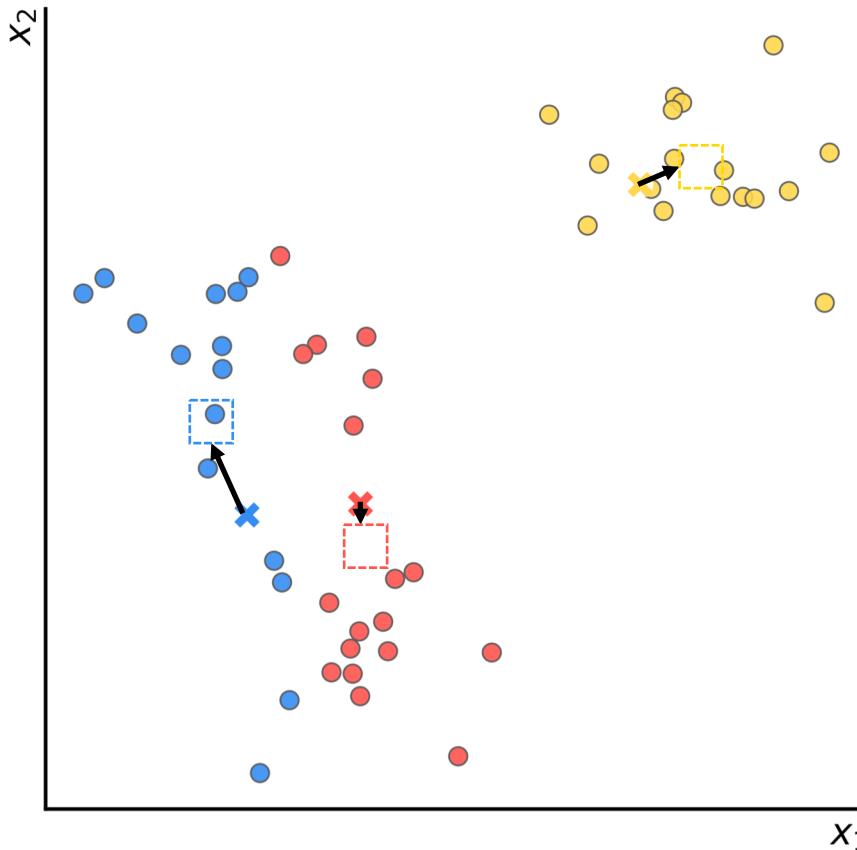
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 2

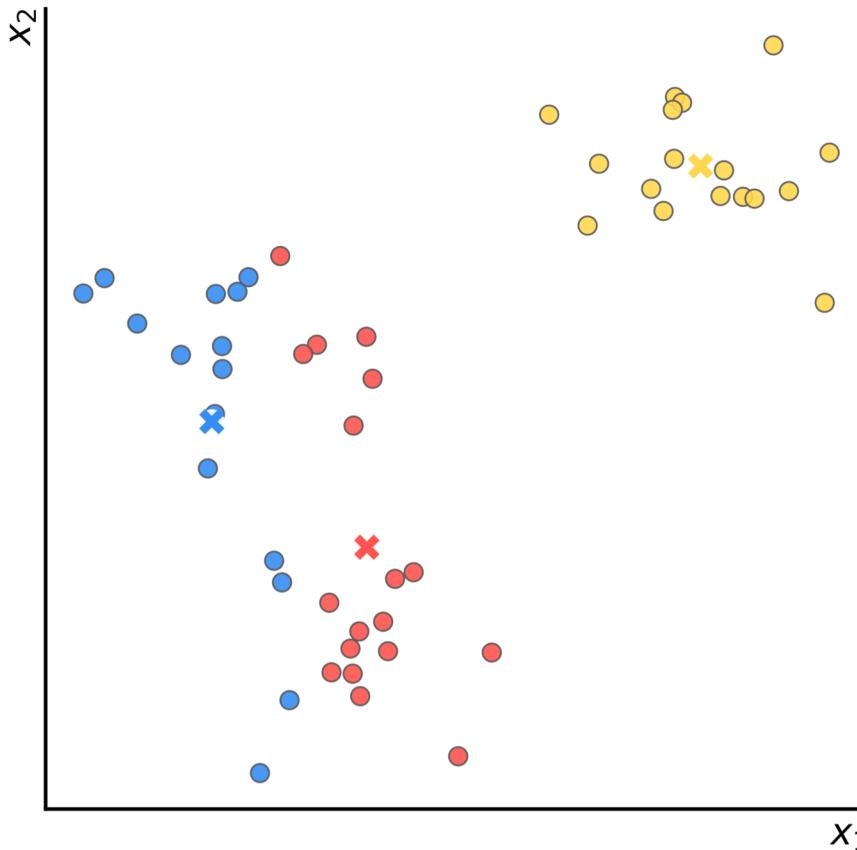
- 4) Assign each instance to **the nearest centroid**;
- 5) Move the centroids to the (new) centers of their **clusters** 



3) Repeat until **centroids do not move significantly**:

iteration 2

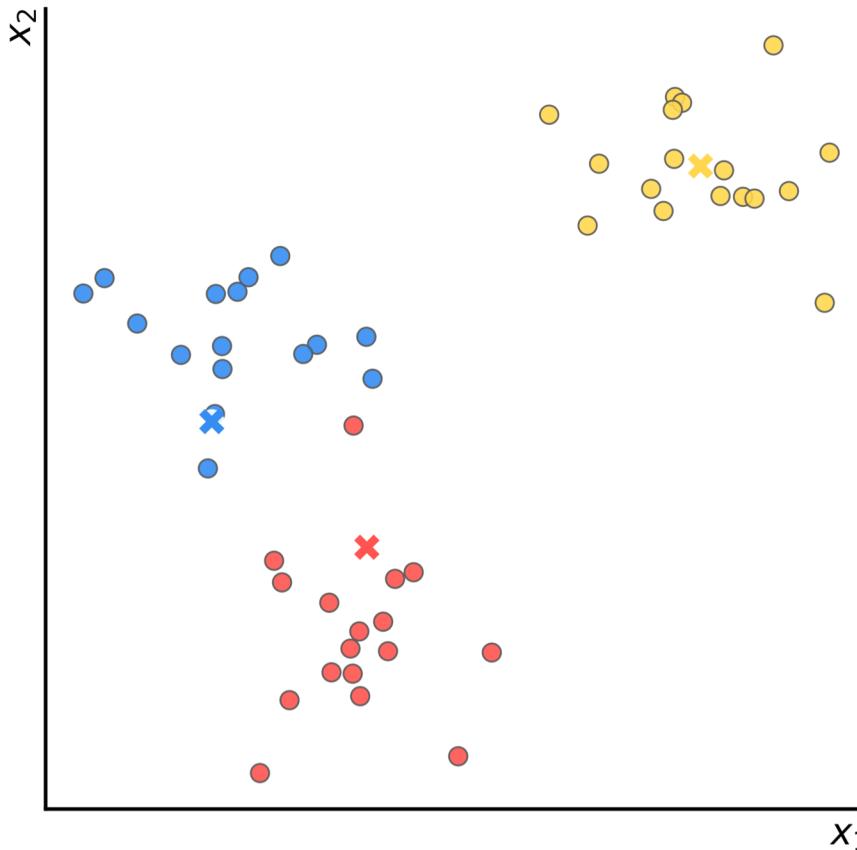
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters** 



 Note that the **samples' clusters** may be **inconsistent** with the **new centroids**.

3) Repeat until **centroids do not move significantly**:

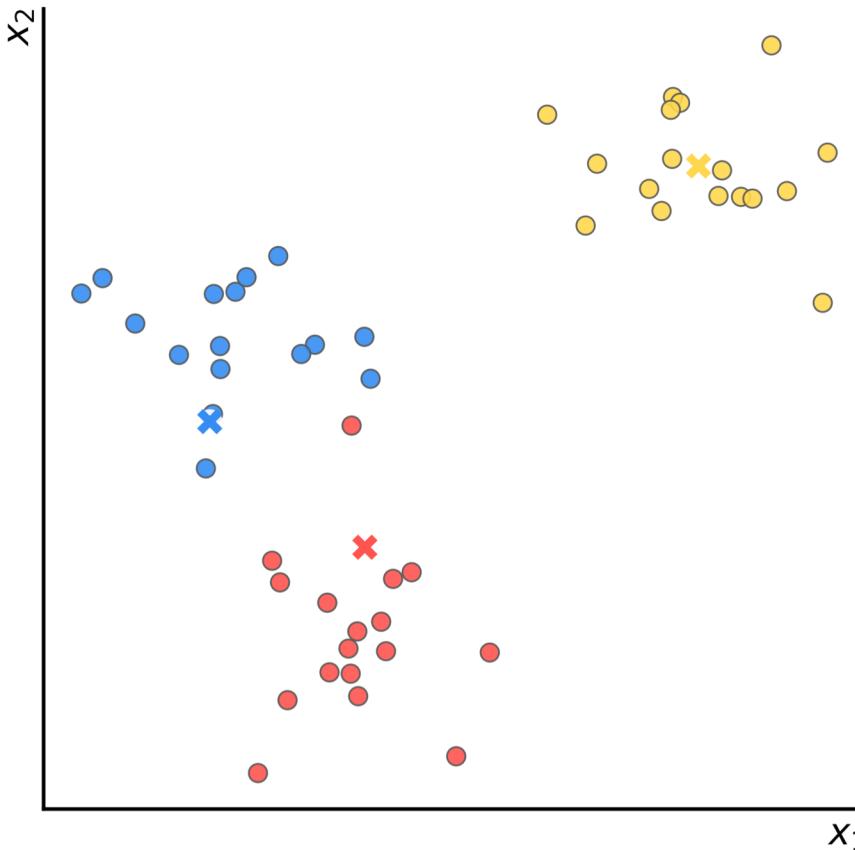
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 3

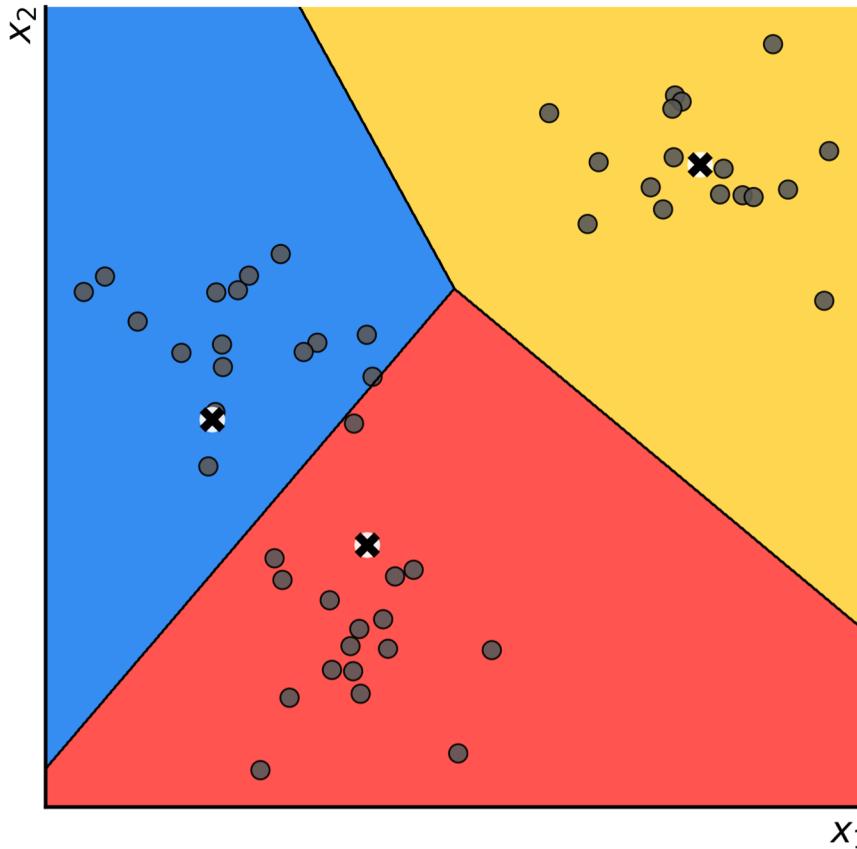
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 3

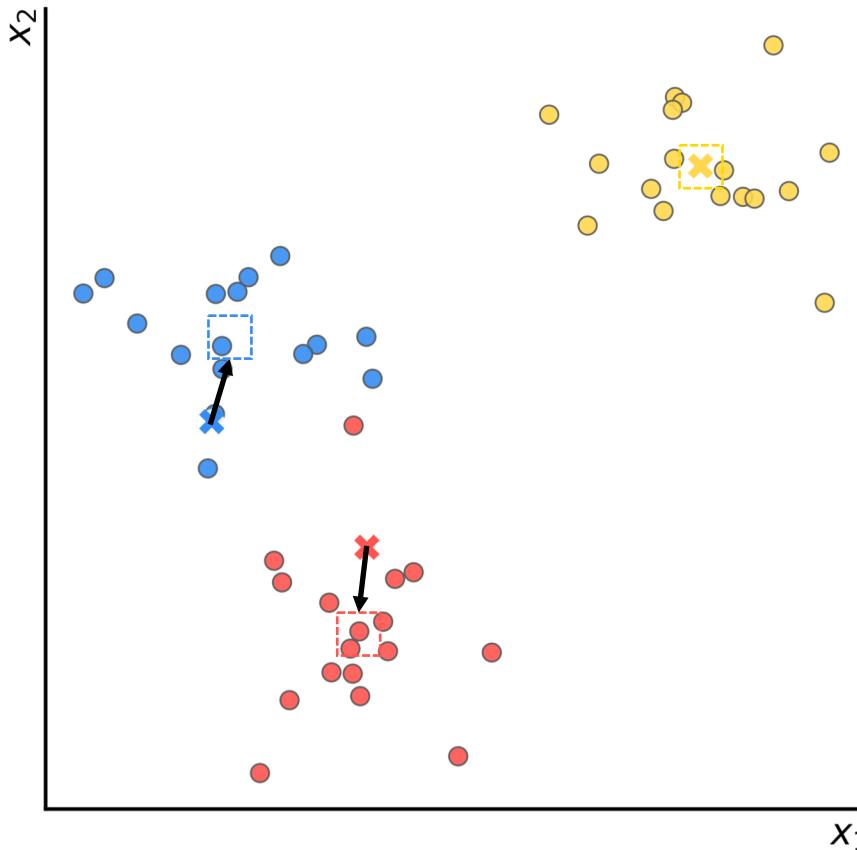
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 3

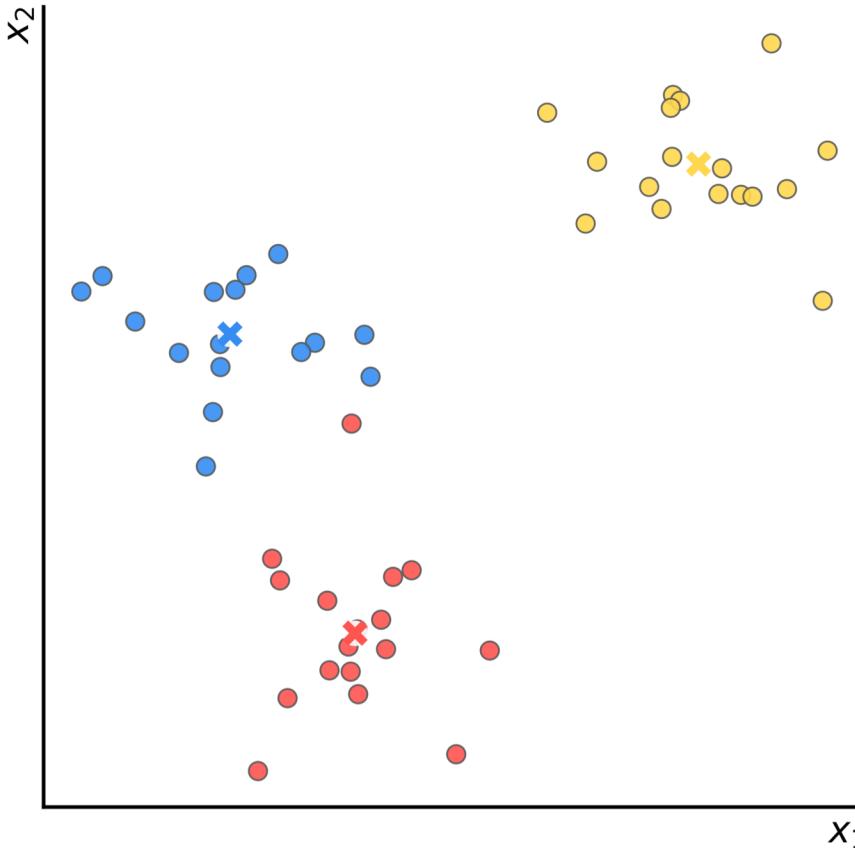
- 4) Assign each instance to **the nearest centroid**;
- 5) Move the centroids to the (new) centers of their **clusters** 



3) Repeat until **centroids do not move significantly**:

iteration 3

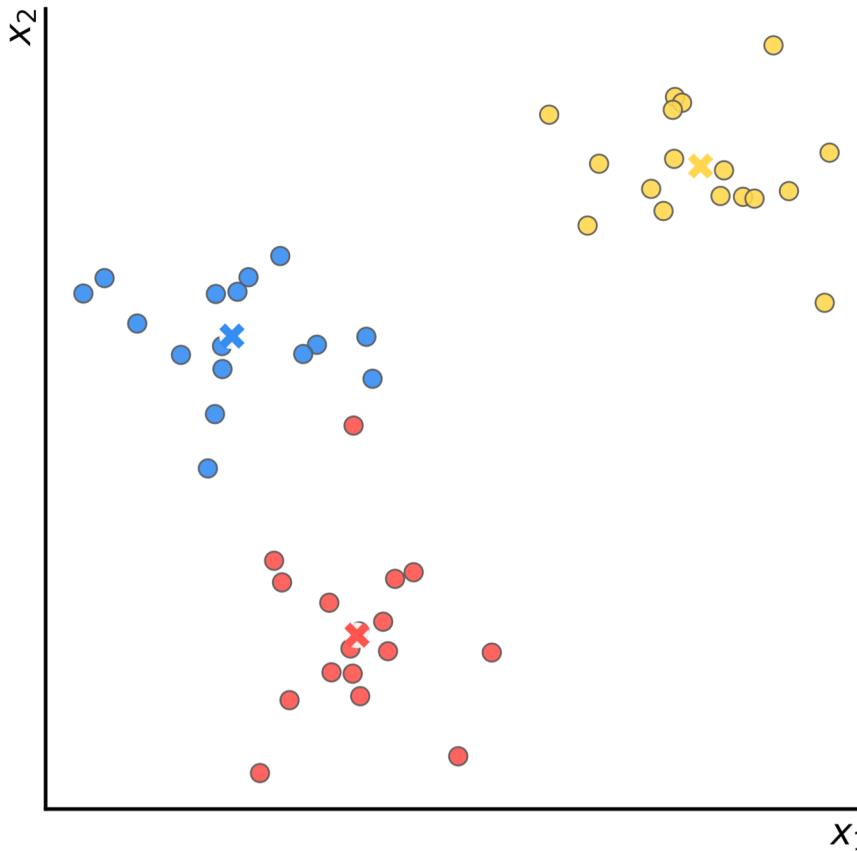
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters** 



 Note that the **samples' clusters** may be **inconsistent** with the **new centroids**.

3) Repeat until **centroids do not move significantly**:

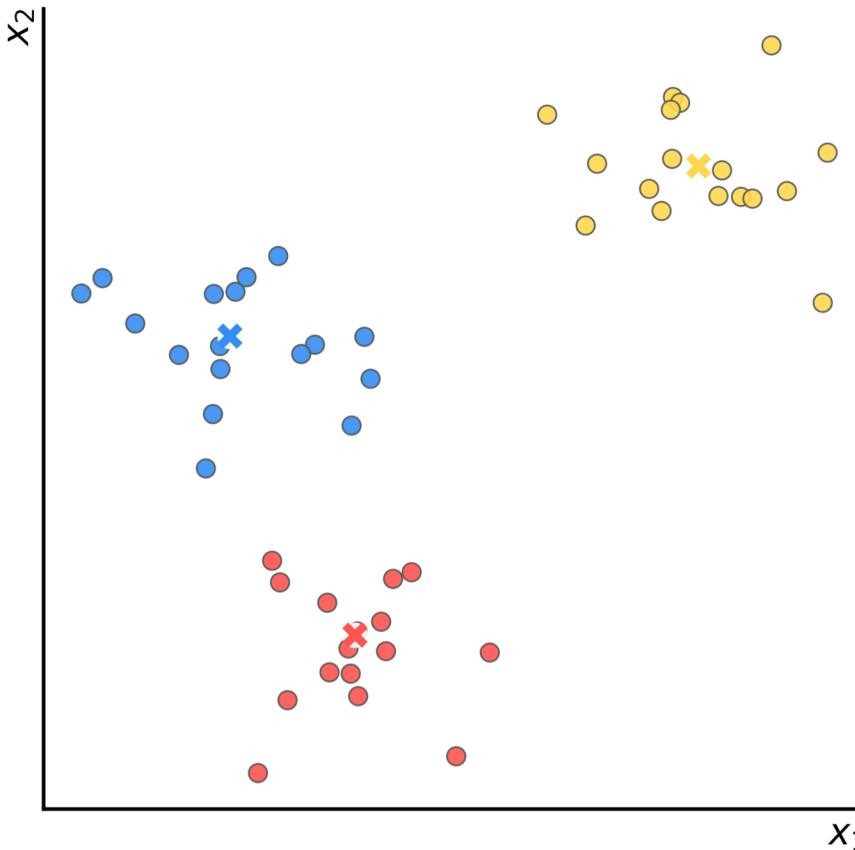
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 4

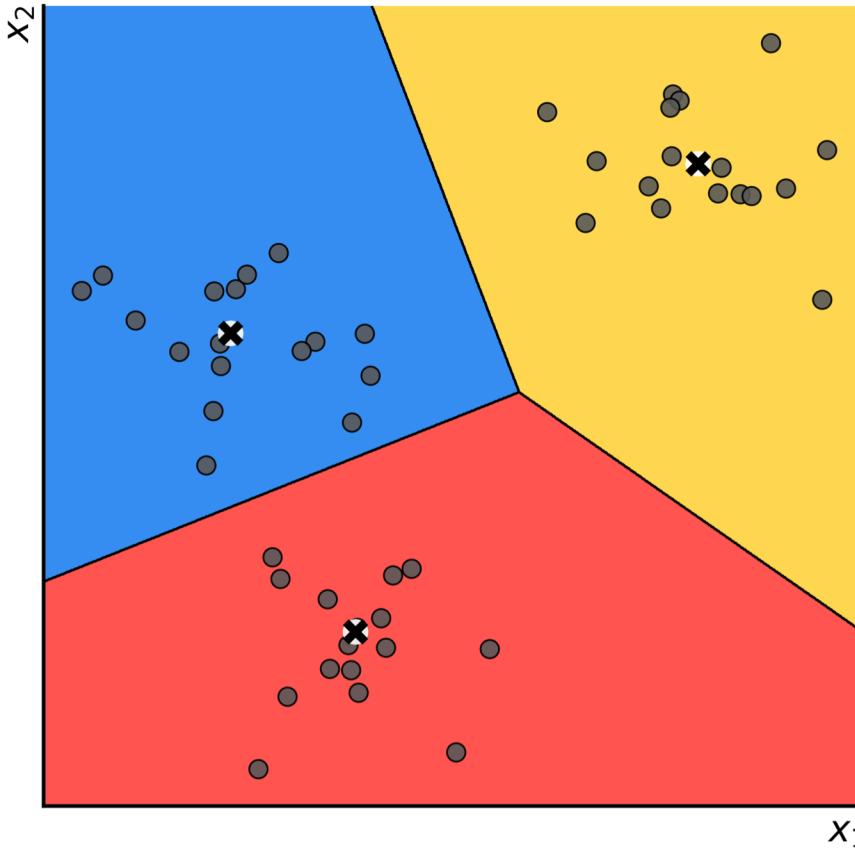
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 4

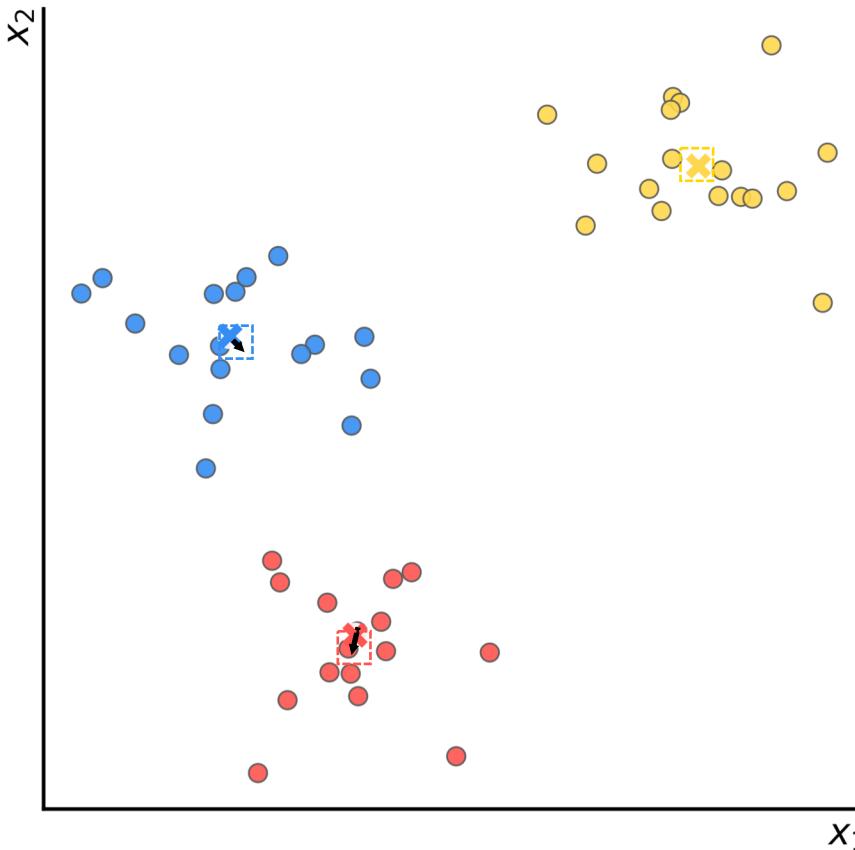
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 4

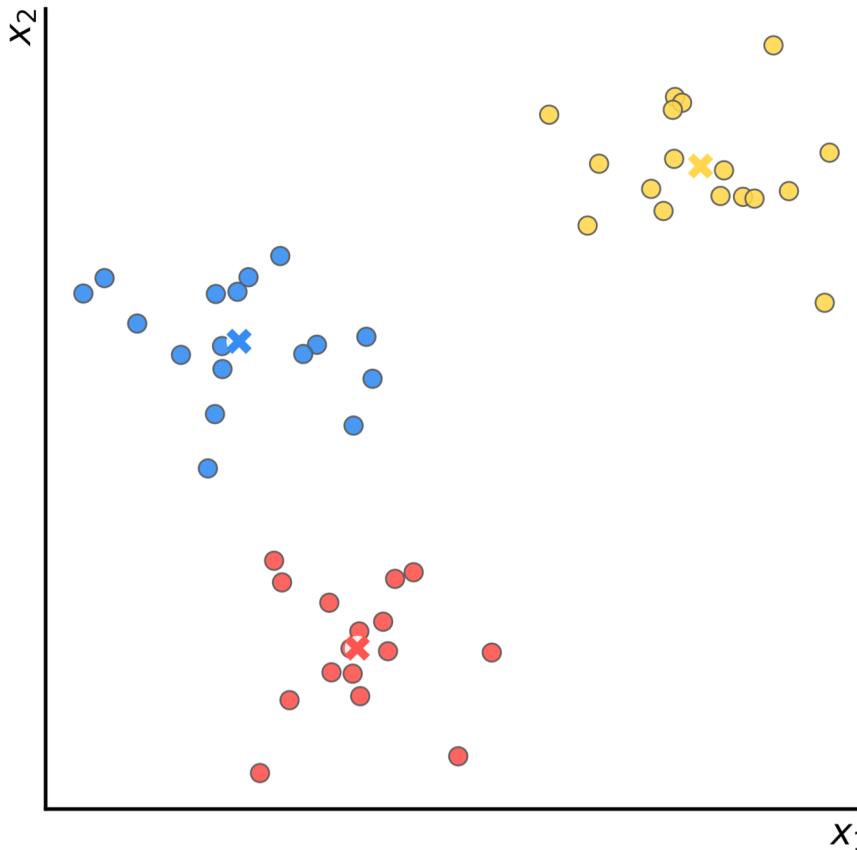
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters** 



3) Repeat until **centroids do not move significantly**:

iteration 4

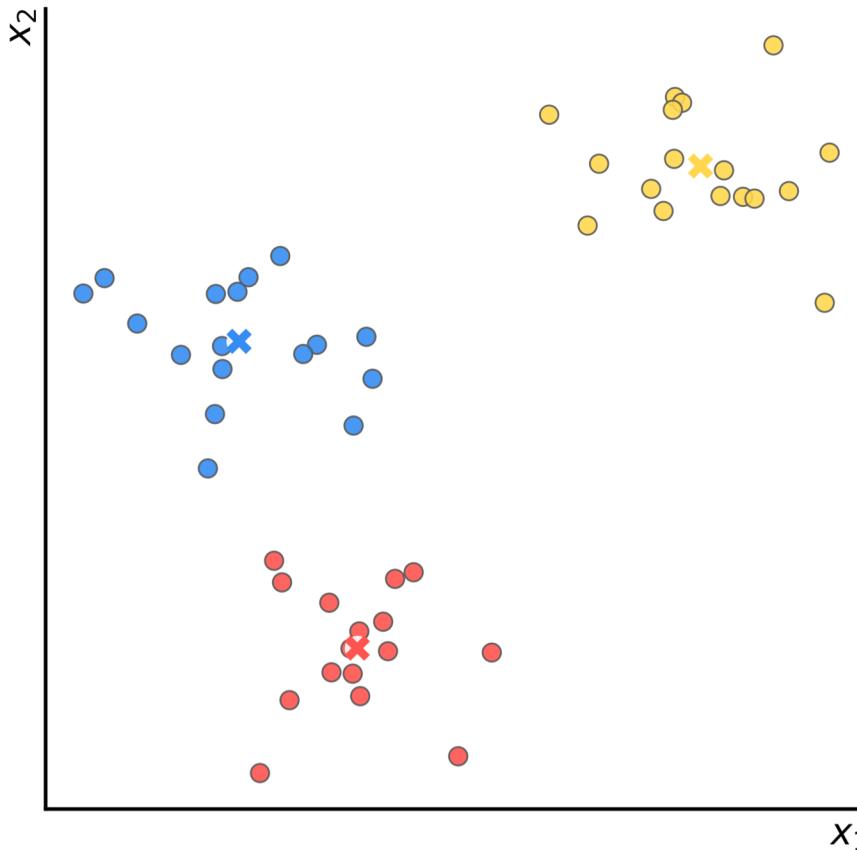
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters** 



 Note that the **samples' clusters** may be **inconsistent** with the **new centroids**.

3) Repeat until **centroids do not move significantly**:

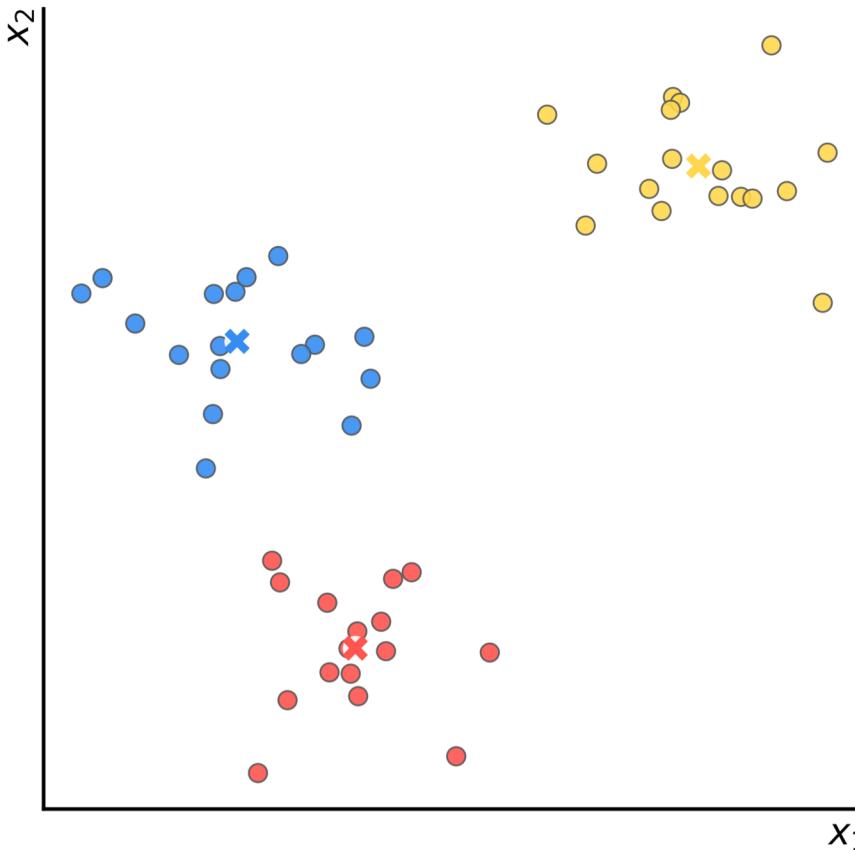
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 5

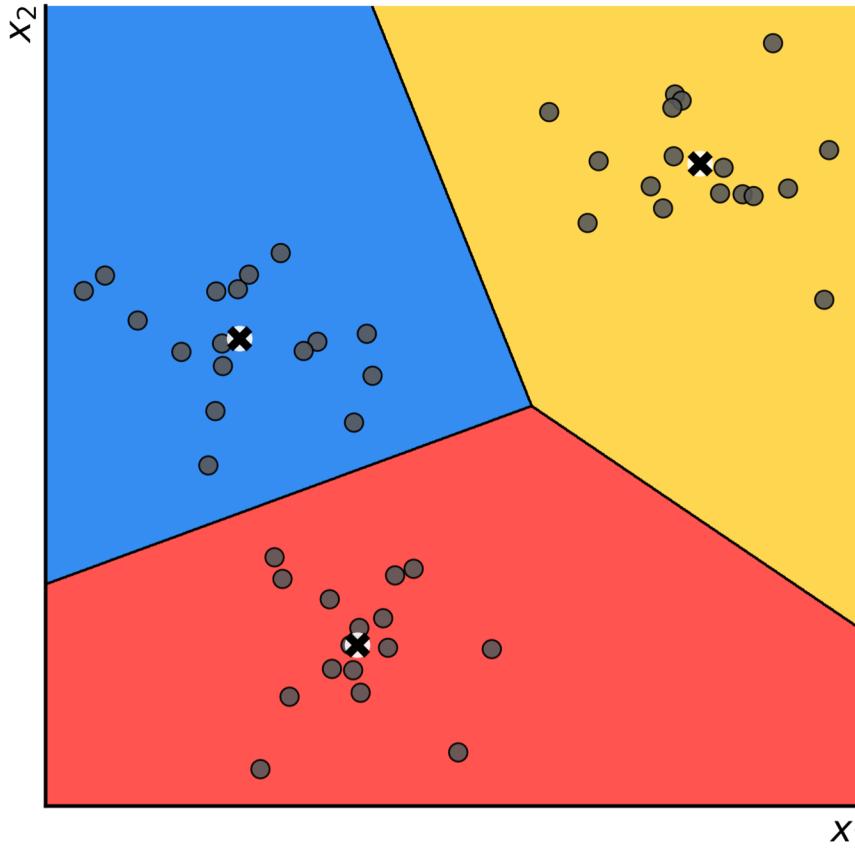
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

iteration 5

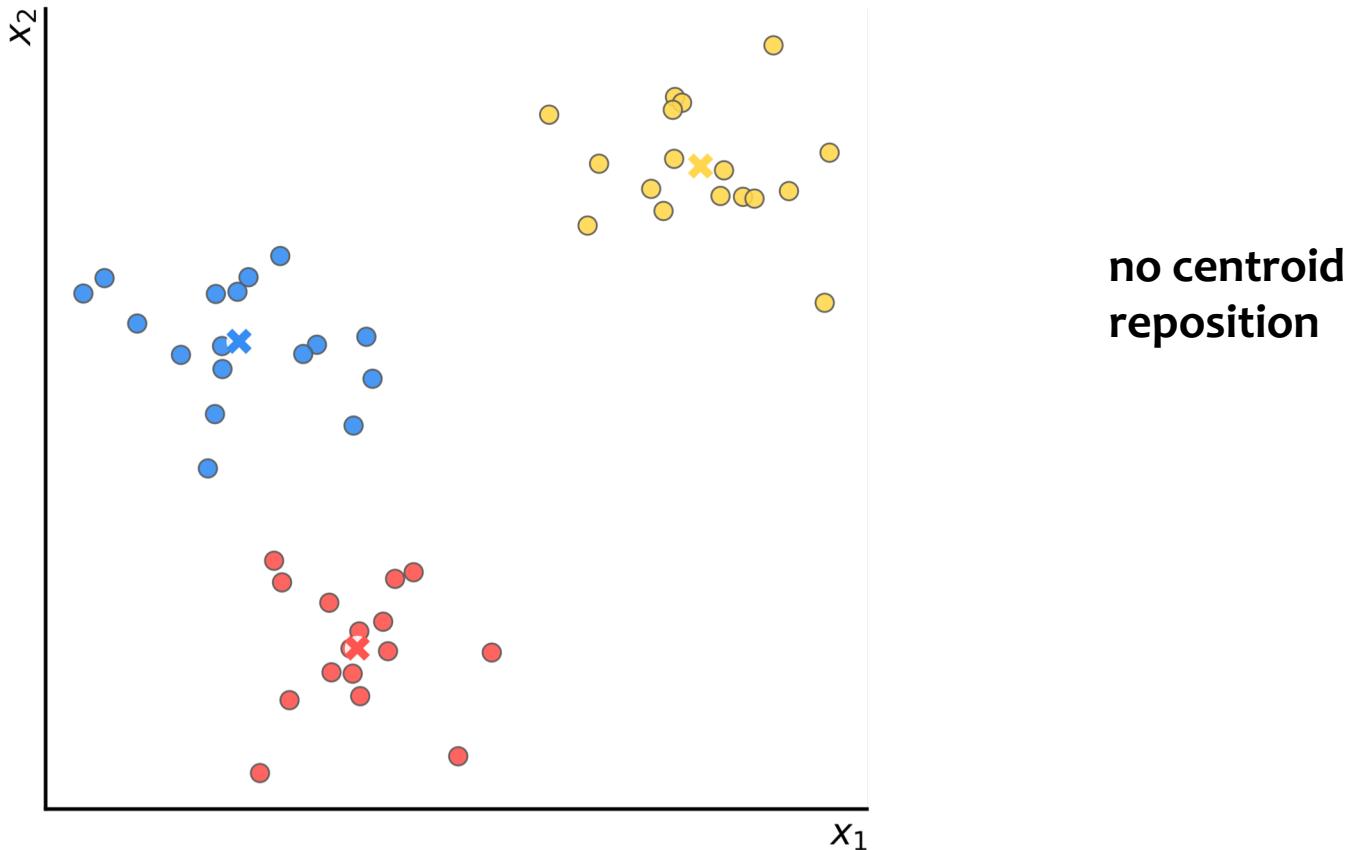
- 4) Assign each instance to **the nearest centroid**; 
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

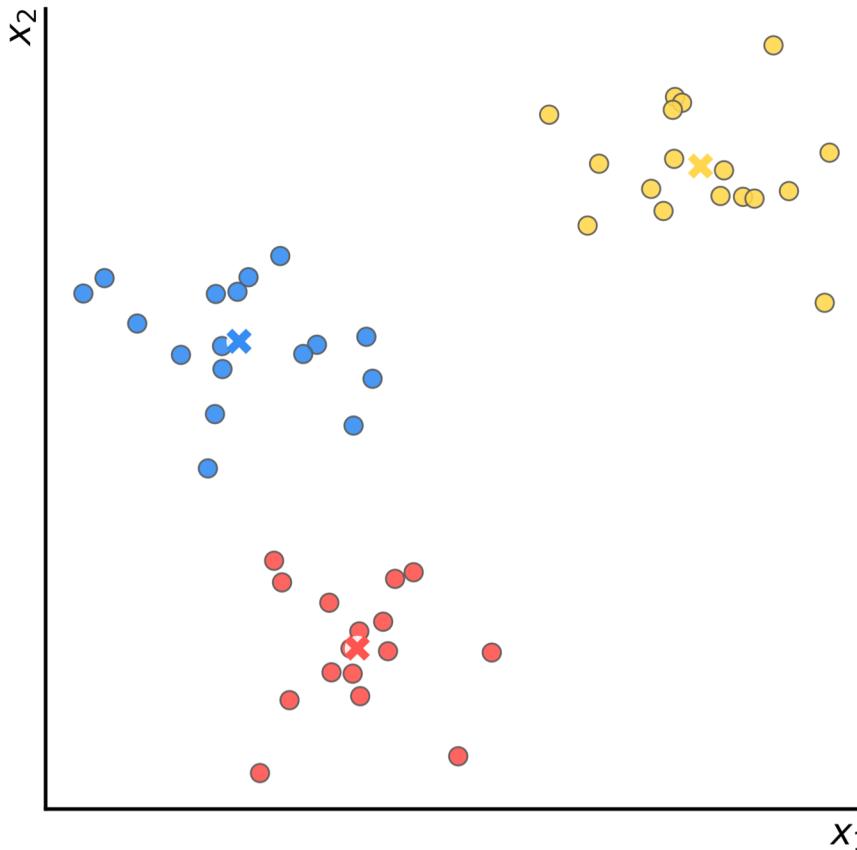
iteration 5

- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters** 



3) Repeat until **centroids do not move significantly**:

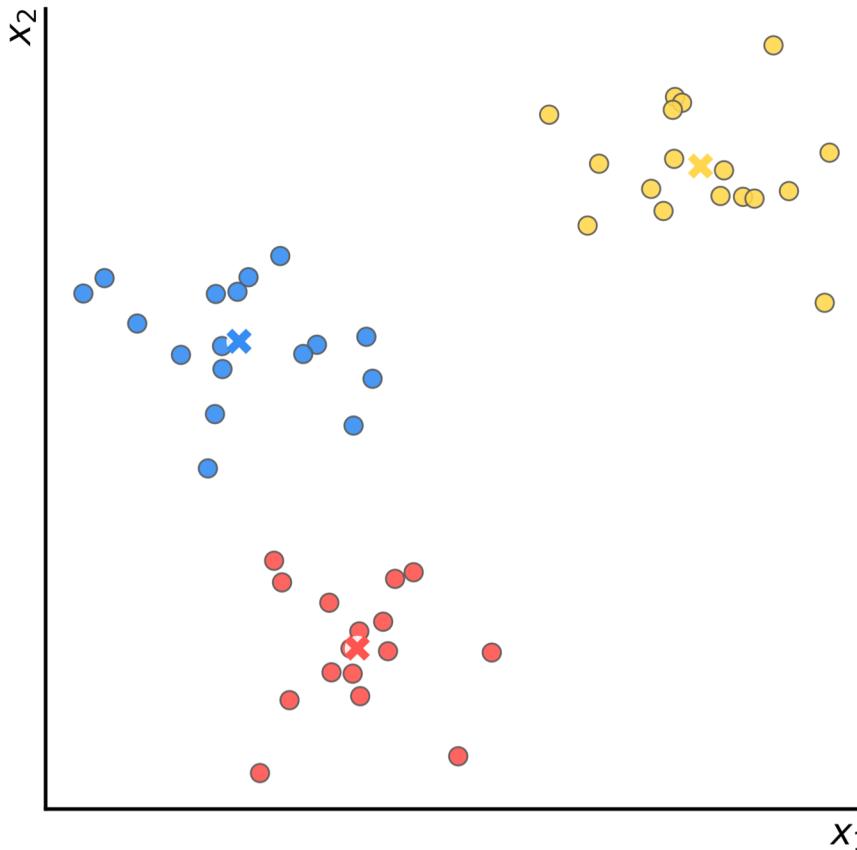
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**: ←

Stop!

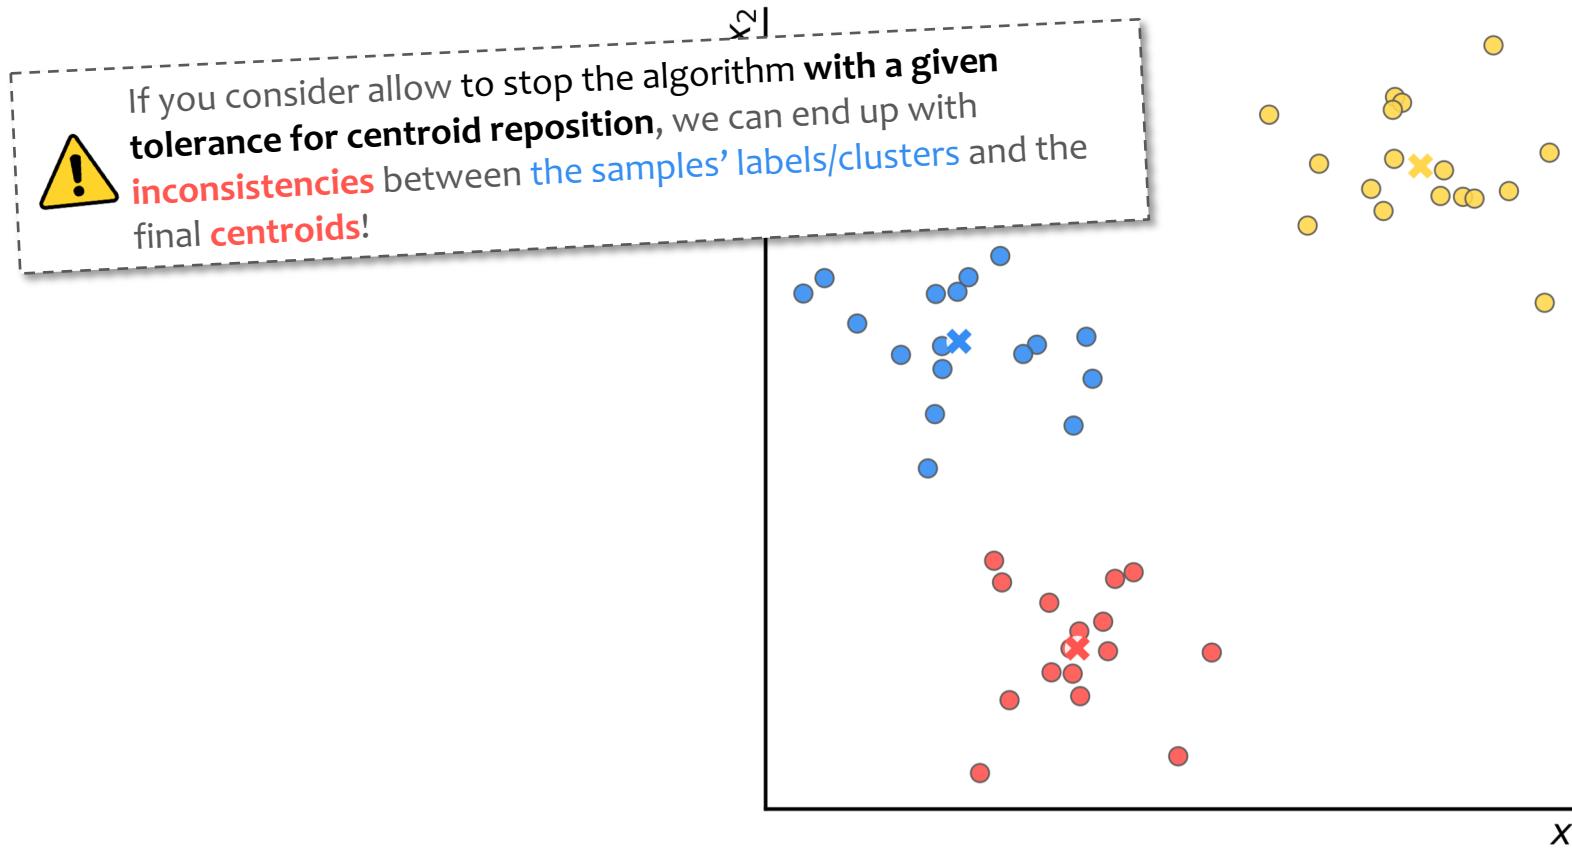
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

Stop!

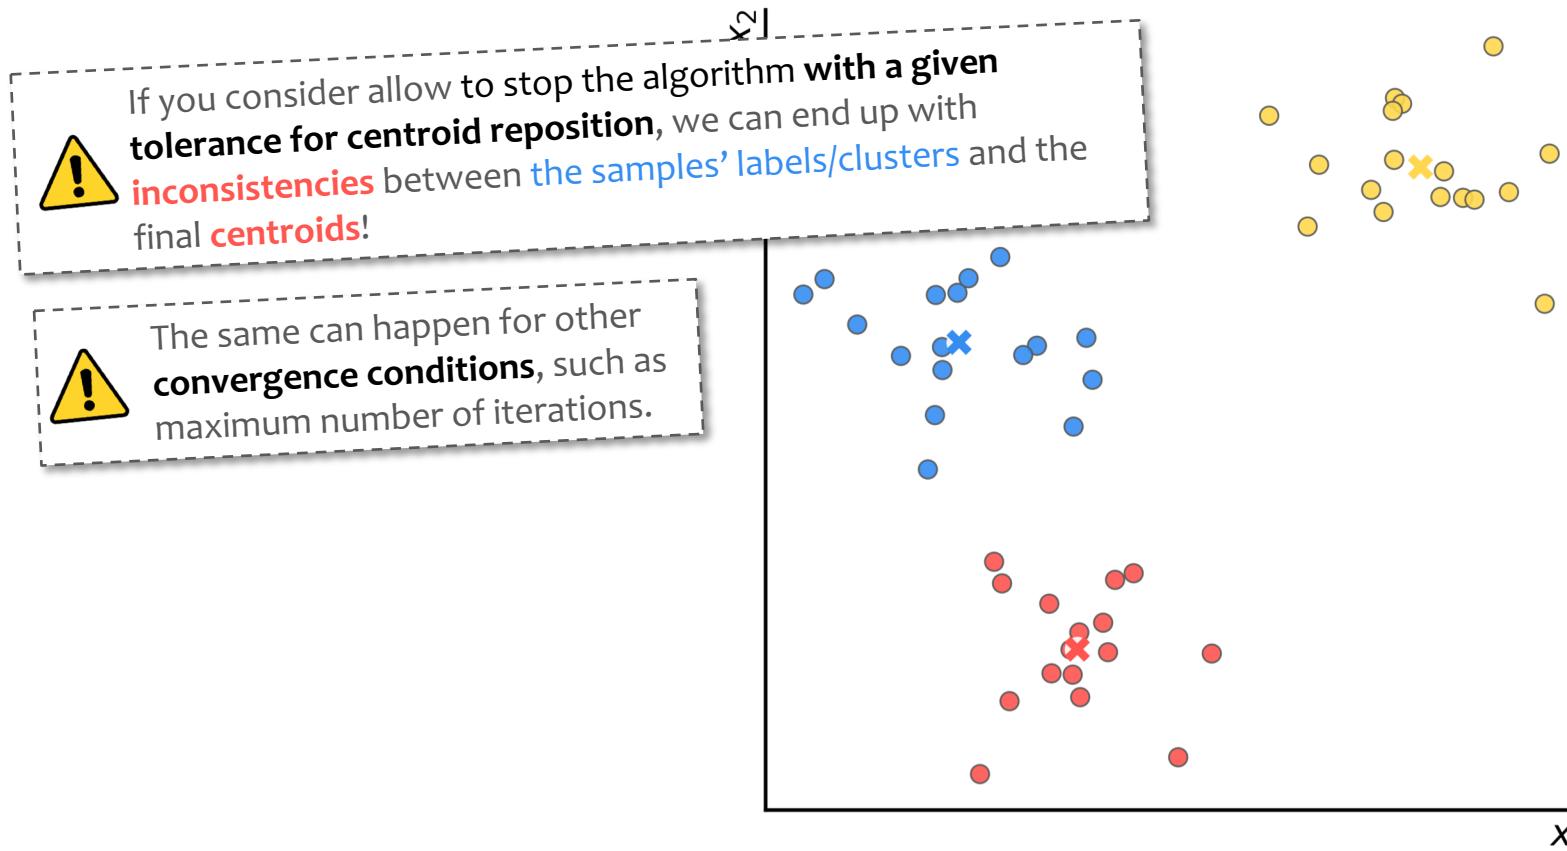
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

Stop!

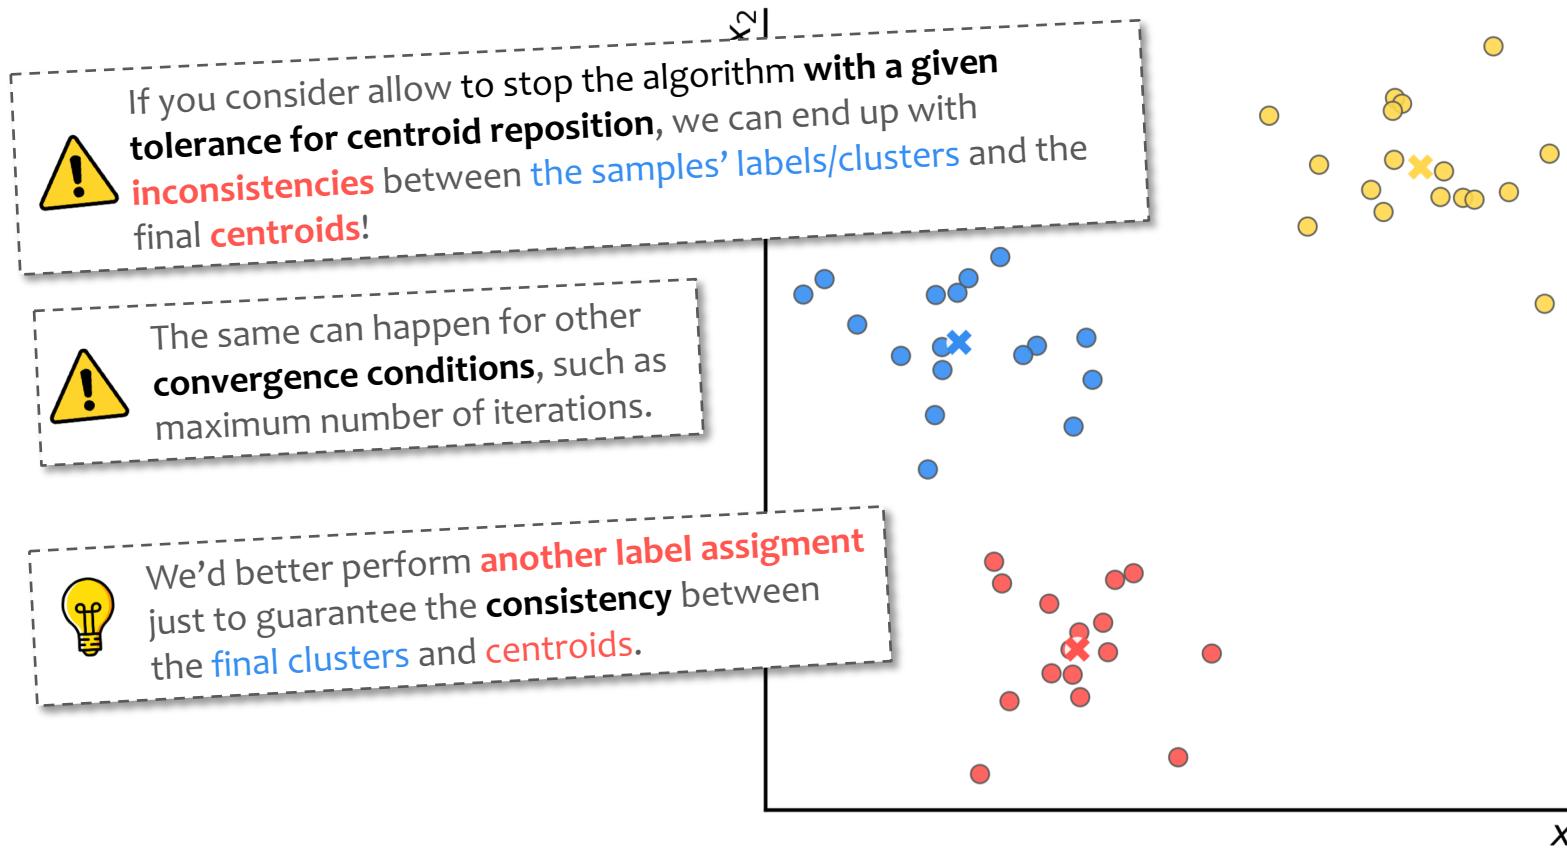
- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



3) Repeat until **centroids do not move significantly**:

Stop!

- 4) Assign each instance to **the nearest** centroid;
- 5) Move the centroids to the (new) centers of their **clusters**



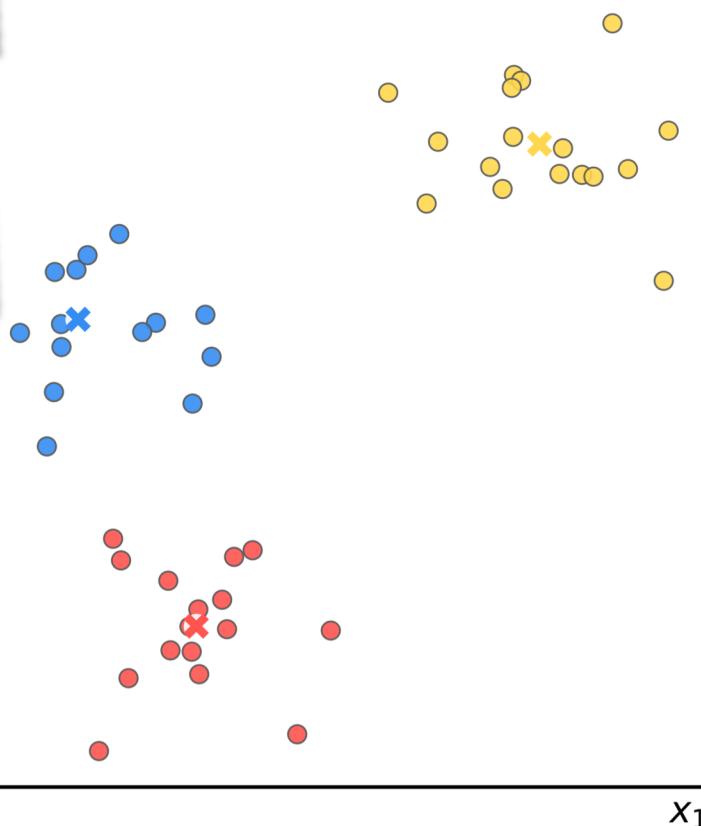
6) Assign each instance to the **final** nearest centroid;

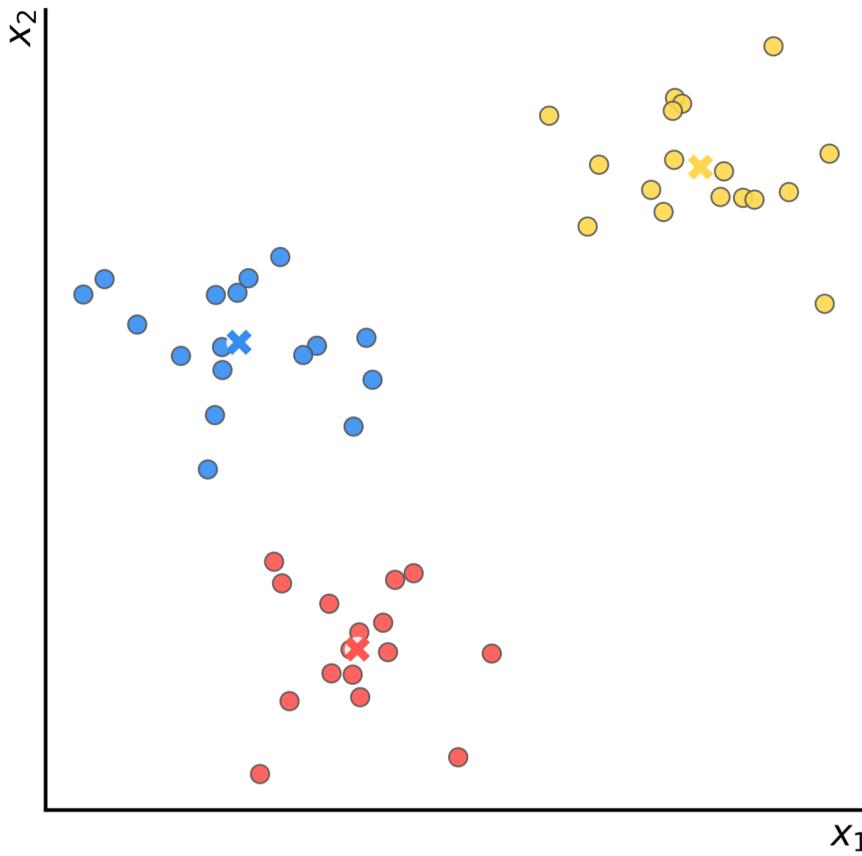


Since the **final centroids** are exactly the same from its previous iteration, for this example, this label assignment obtains **the same clusters**.



KMeans implementation in Scikit-Learn performs **this final label assignment**.

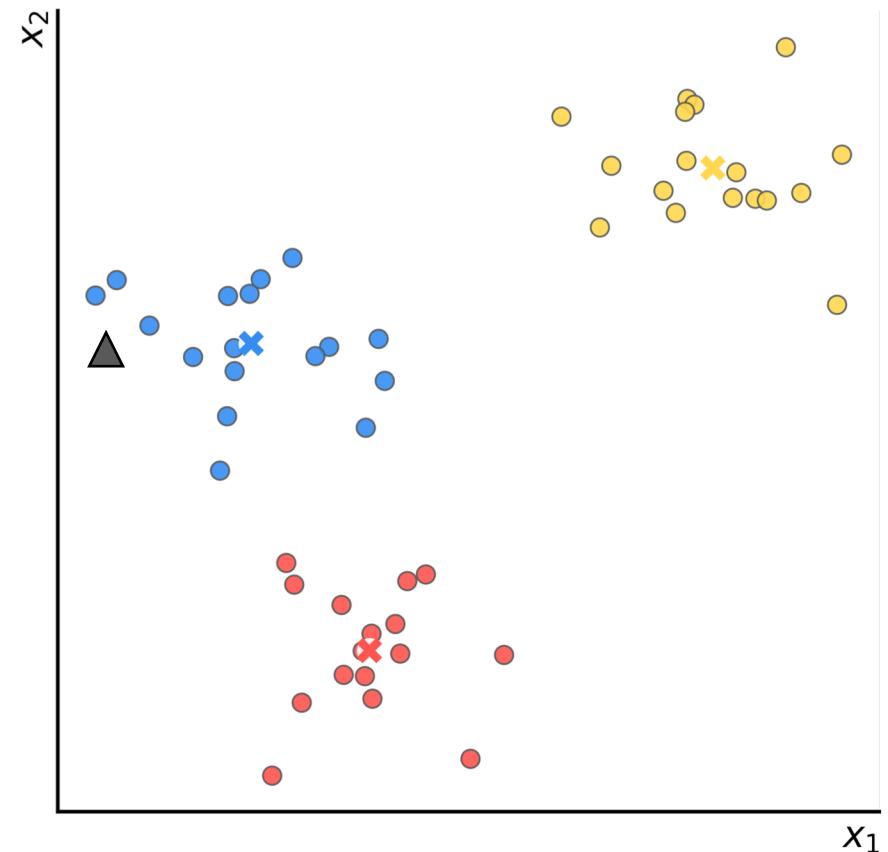




KMeans

Usage

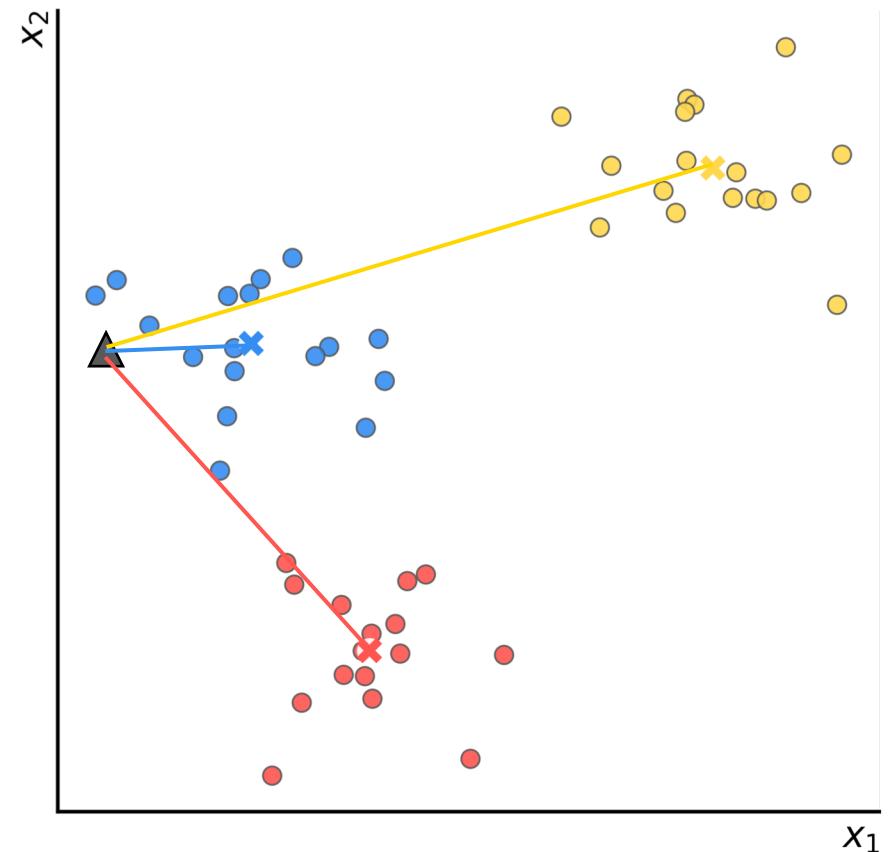
- For each new **testing sample**:
 - Assign each instance to **the nearest** centroid;



KMeans

Usage

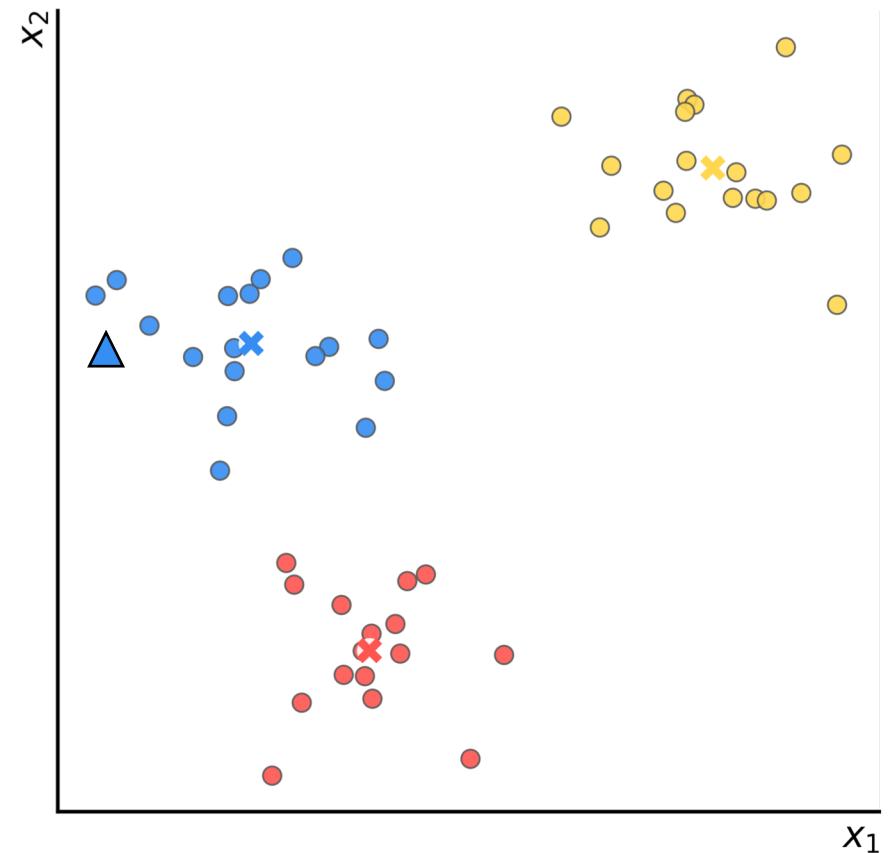
- For each new **testing sample**:
 - Assign each instance to **the nearest** centroid;



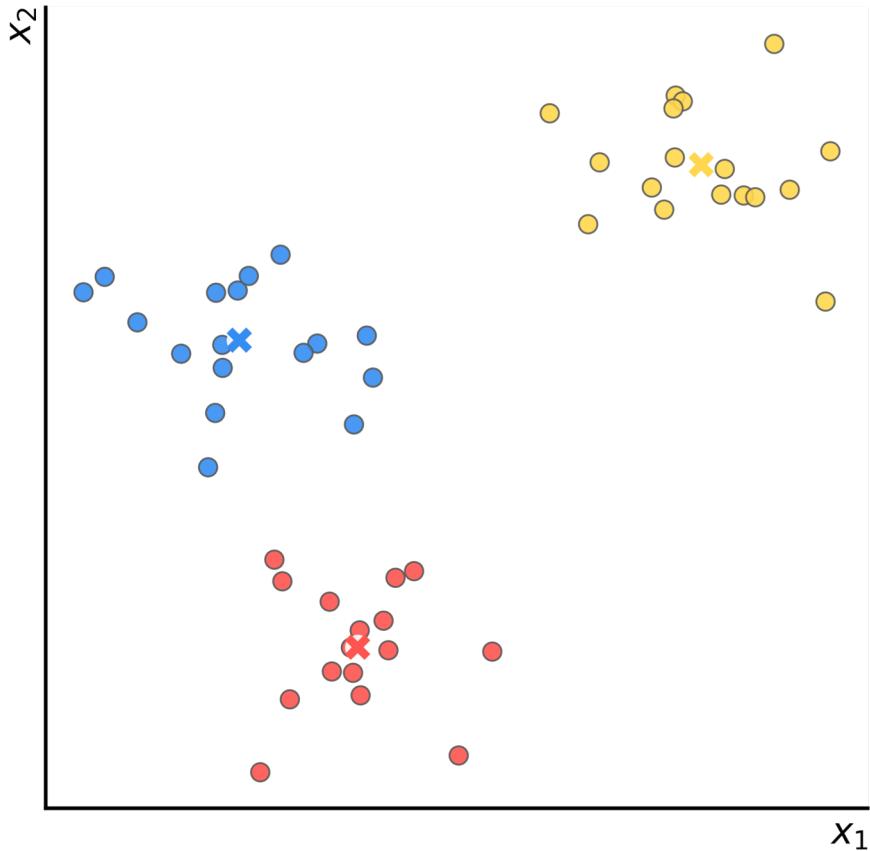
KMeans

Usage

- For each new **testing sample**:
 - Assign each instance to **the nearest** centroid;

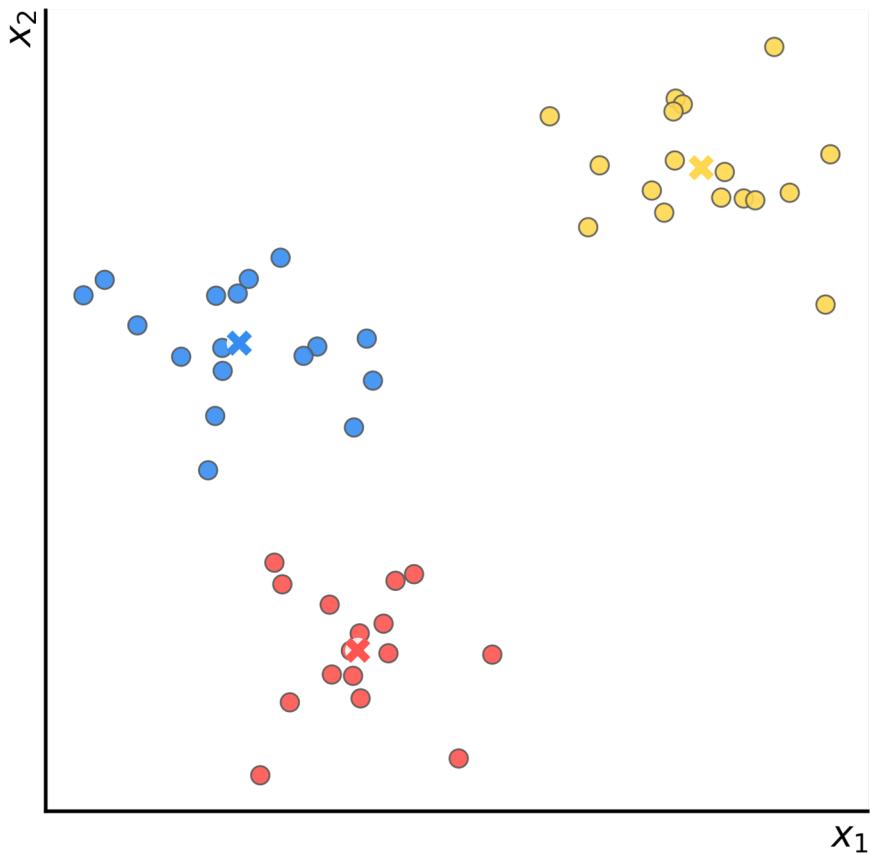


The algorithm is **guaranteed to converge**, but it may
not converge to the **right solution** (i.e., it may converge
to a local optimum).



The algorithm is **guaranteed to converge**, but it may **not converge** to the **right solution** (i.e., it may converge to a local optimum).

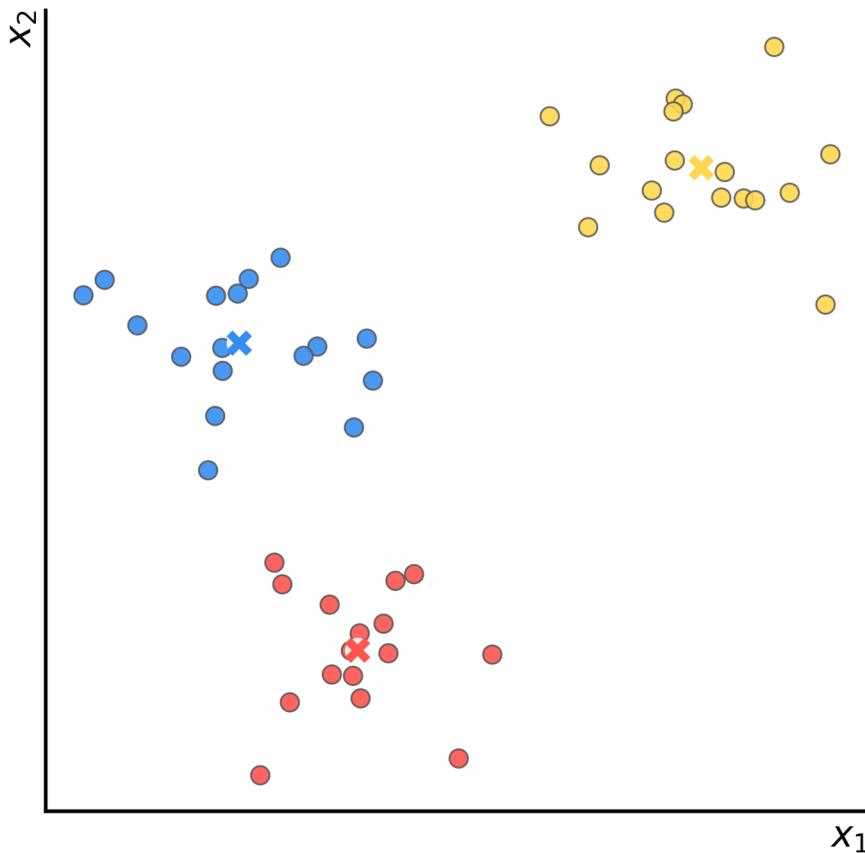
This depends on the **centroid initialization**.



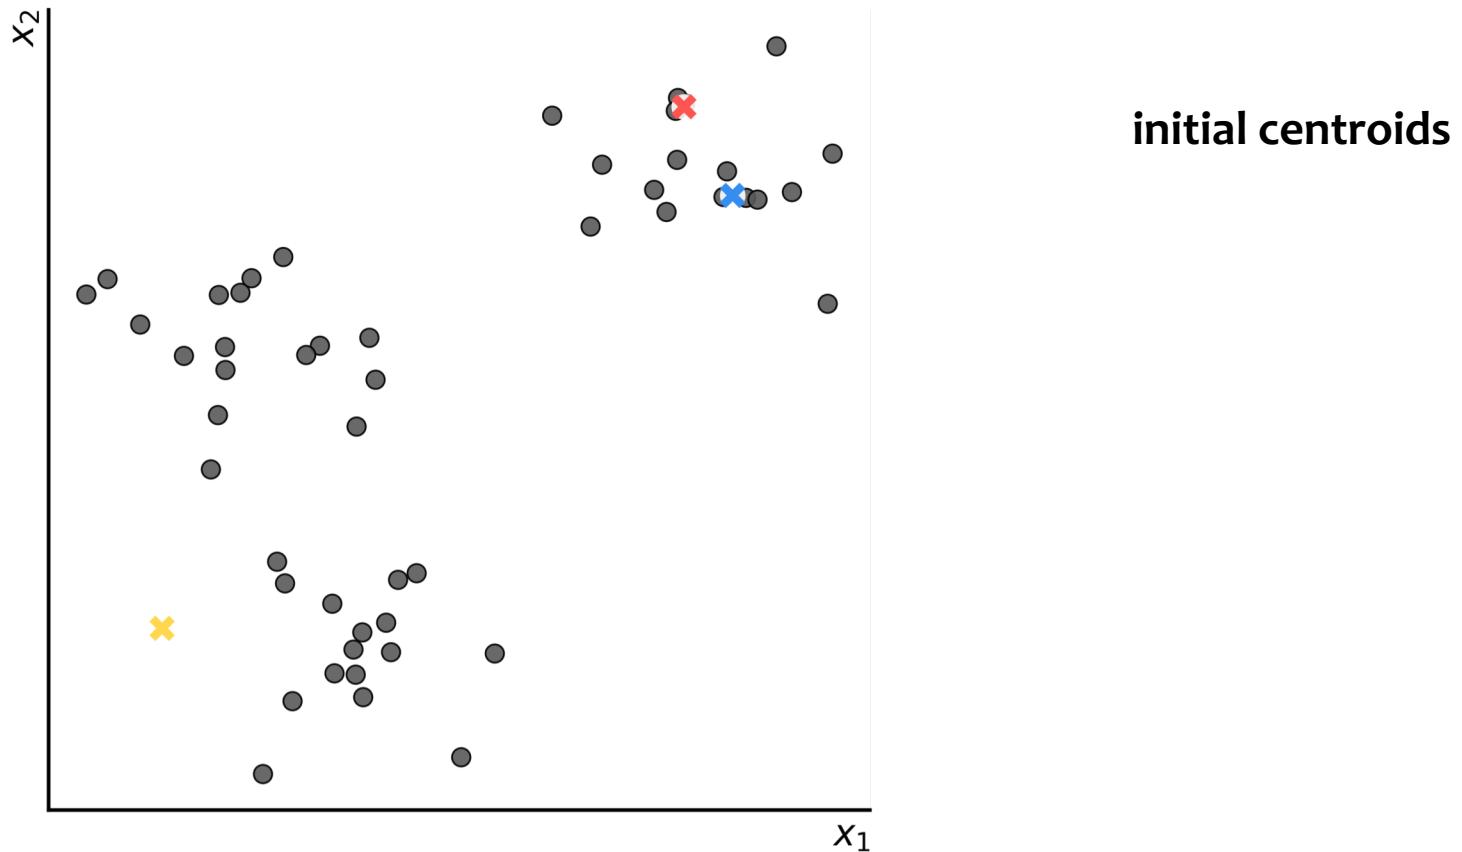
The algorithm is **guaranteed to converge**, but it may not converge to the **right solution** (i.e., it may converge to a local optimum).

This depends on the **centroid initialization**.

Poor initial centroids result to poor clusters.

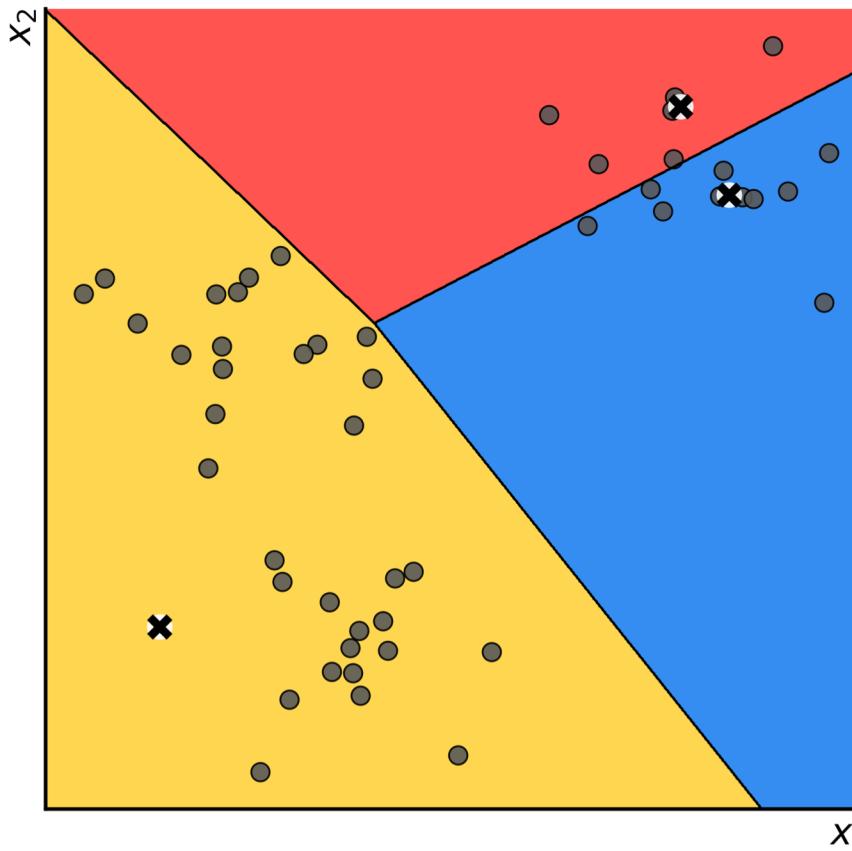


Example with poor initial centroids



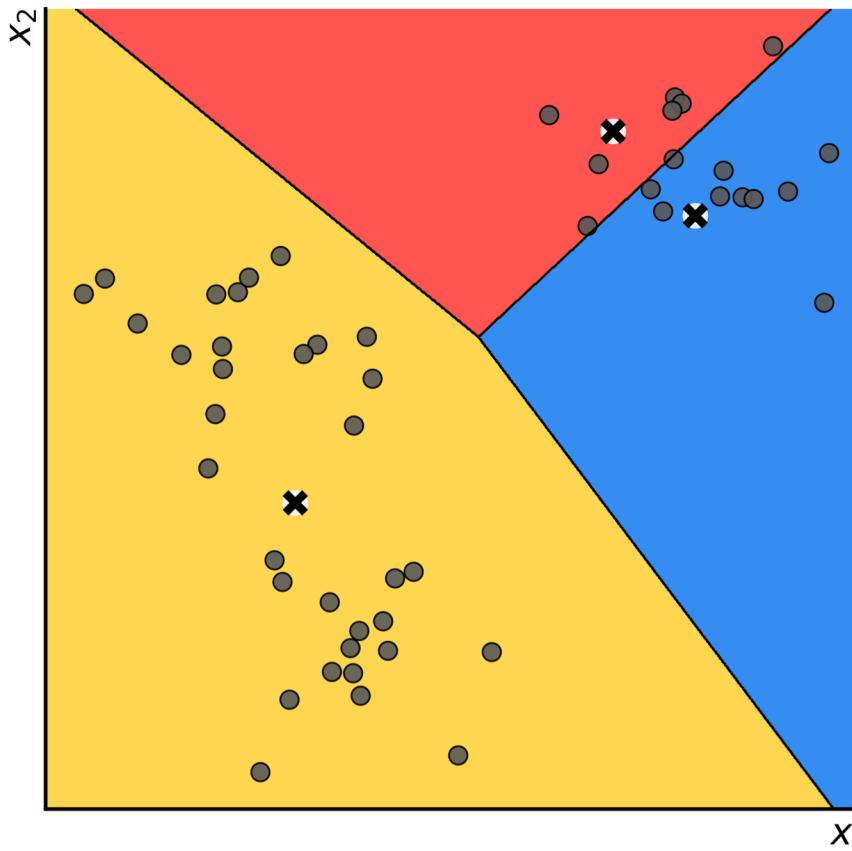
Example with poor initial centroids

Iteration 1



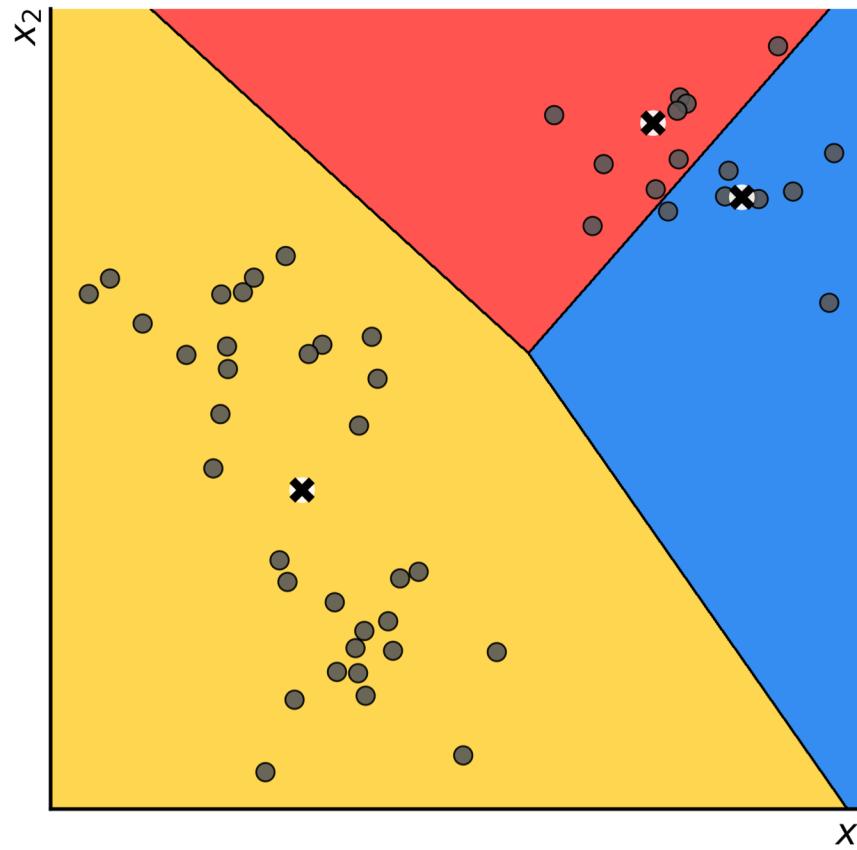
Example with poor initial centroids

Iteration 2



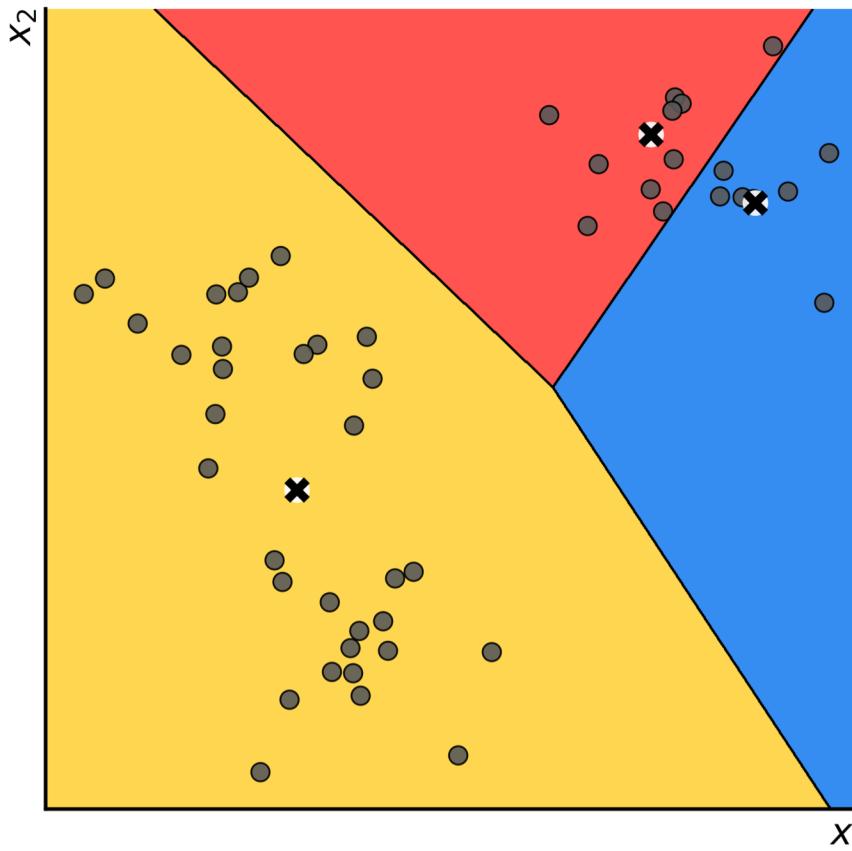
Example with poor initial centroids

Iteration 3

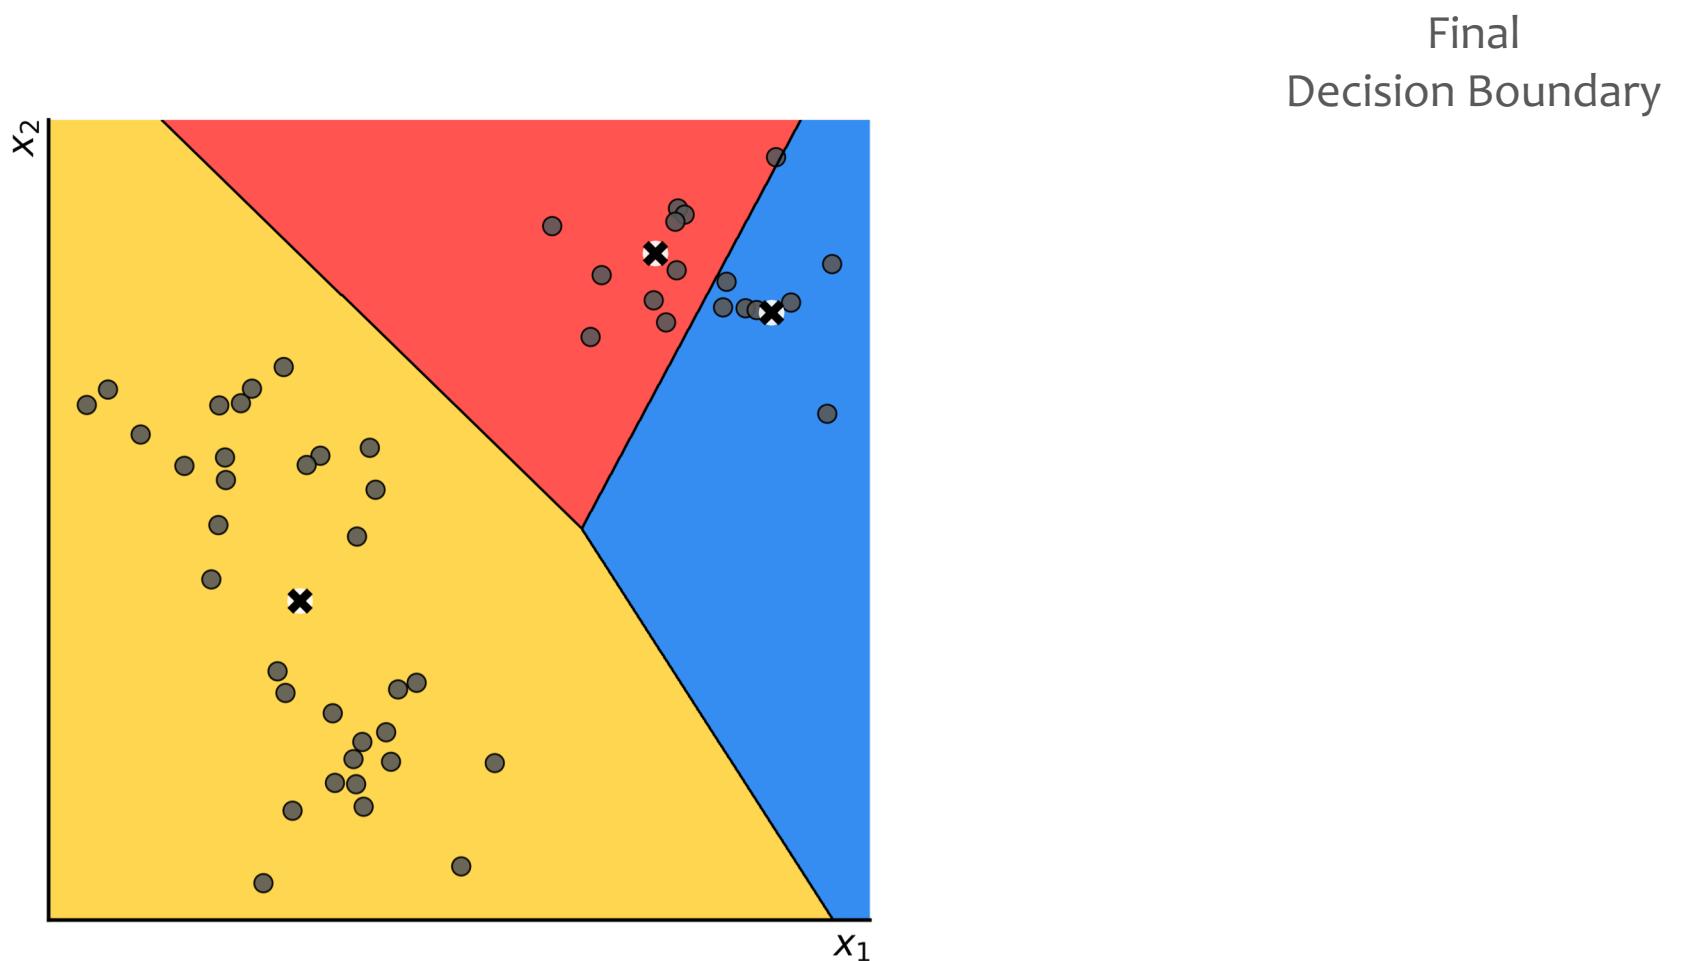


Example with poor initial centroids

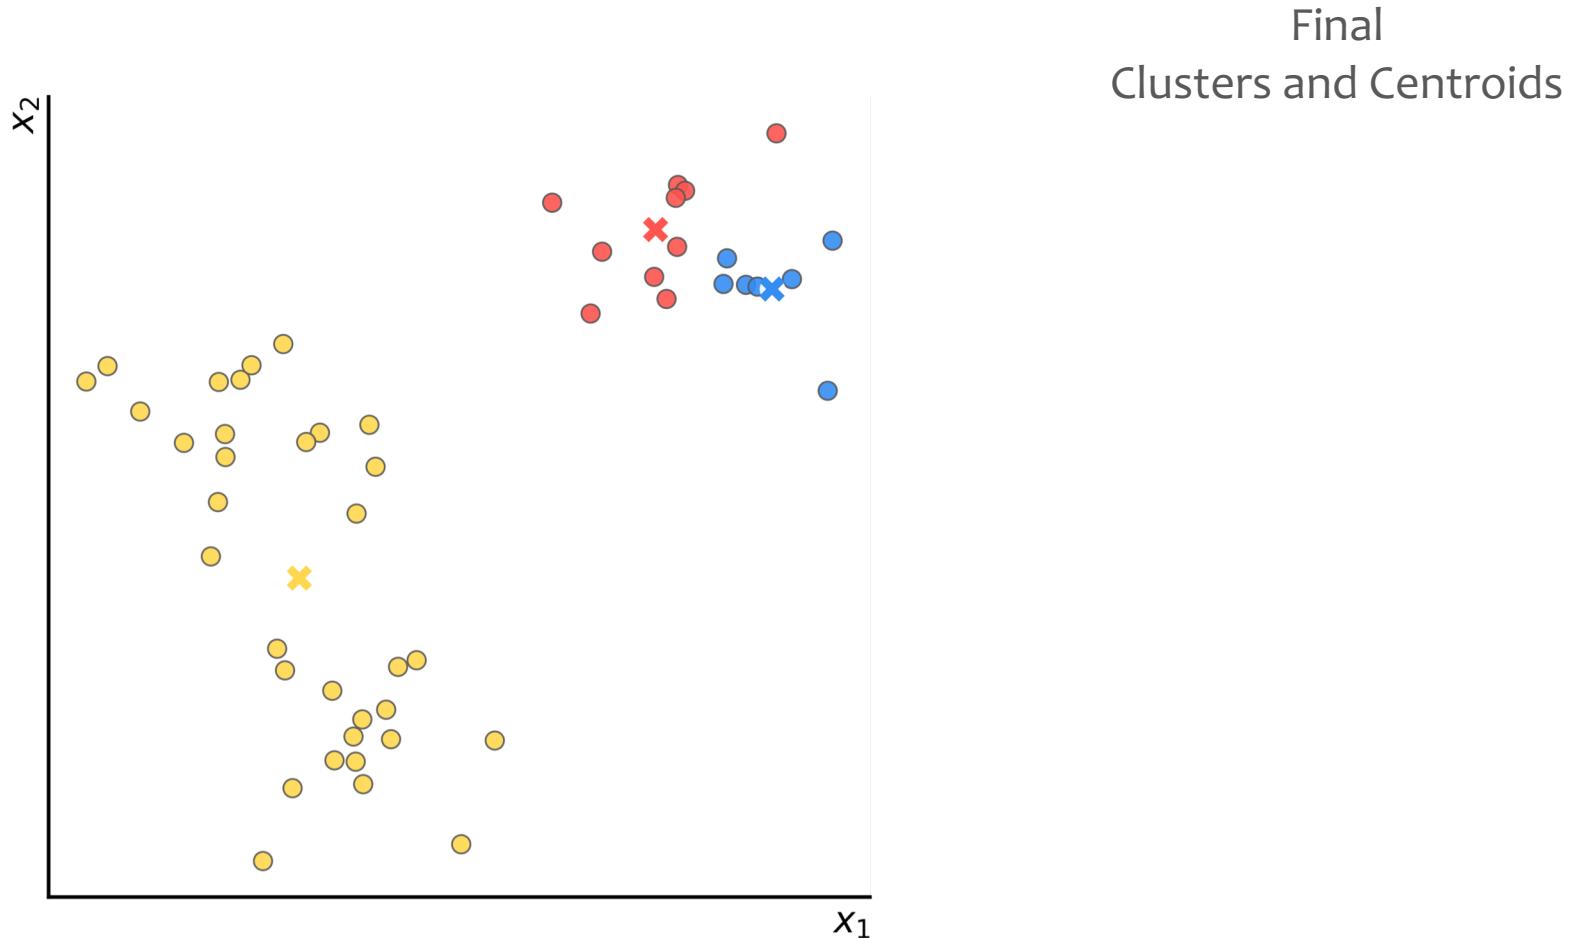
Iteration 4



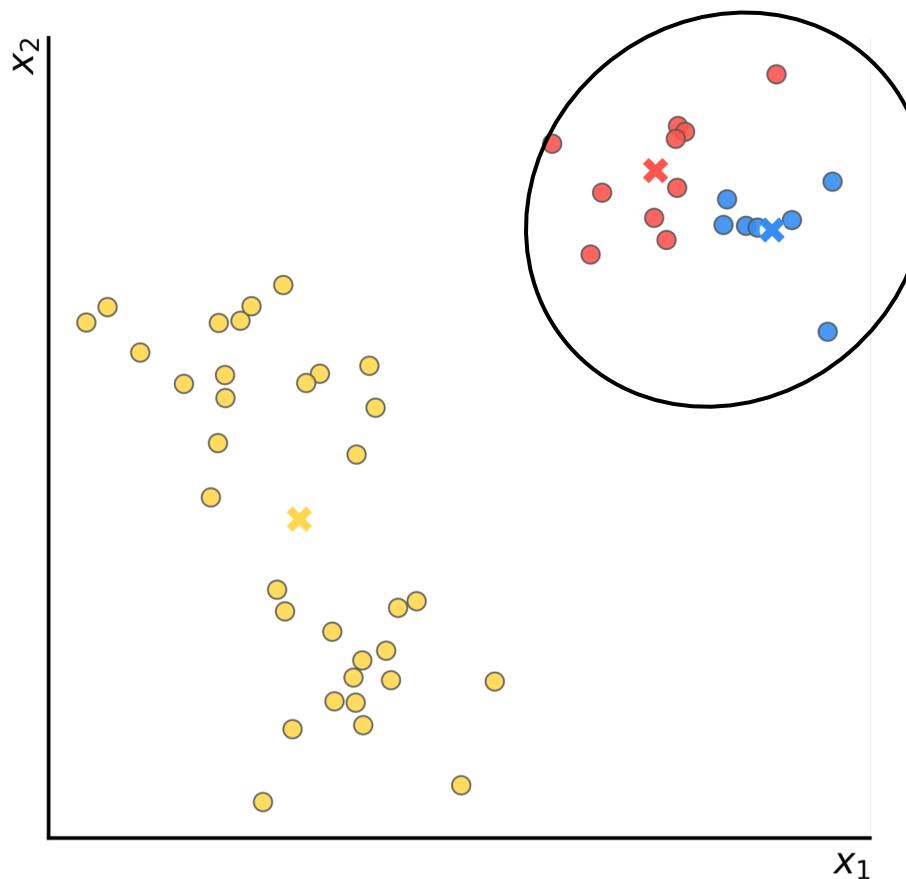
Example with poor initial centroids



Example with poor initial centroids



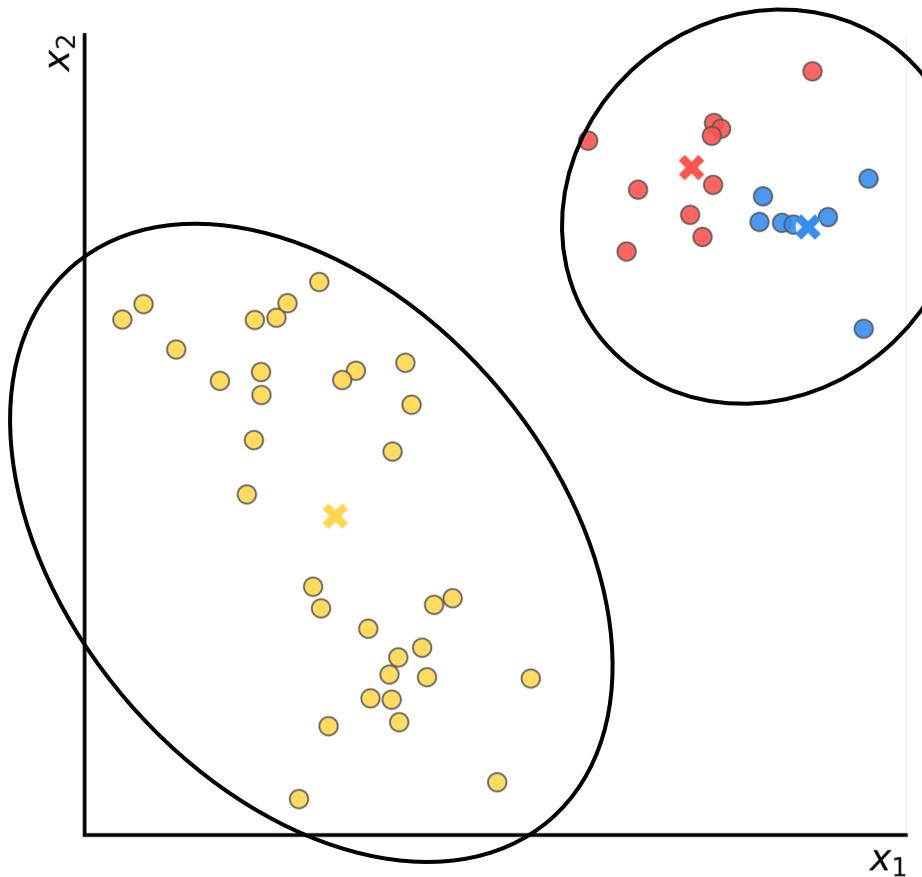
Example with poor initial centroids



A **cluster** is overpartitioned
into two clusters

Example with poor initial centroids

Two **clusters** is underpartitioned into a single cluster



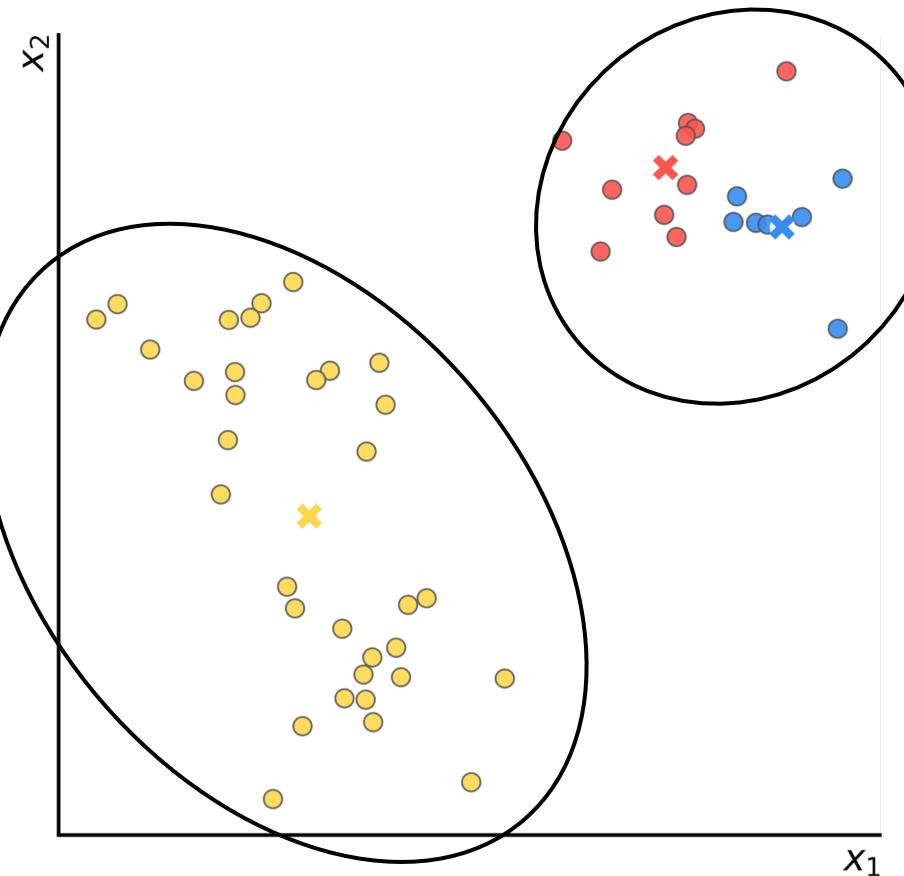
A **cluster** is overpartitioned into two clusters

Example with poor initial centroids

Two **clusters** is underpartitioned into a single cluster



If we are **unlucky**, our **random initialization** may select **poor initial centroids**.



A **cluster** is overpartitioned into two clusters

Example with poor initial centroids

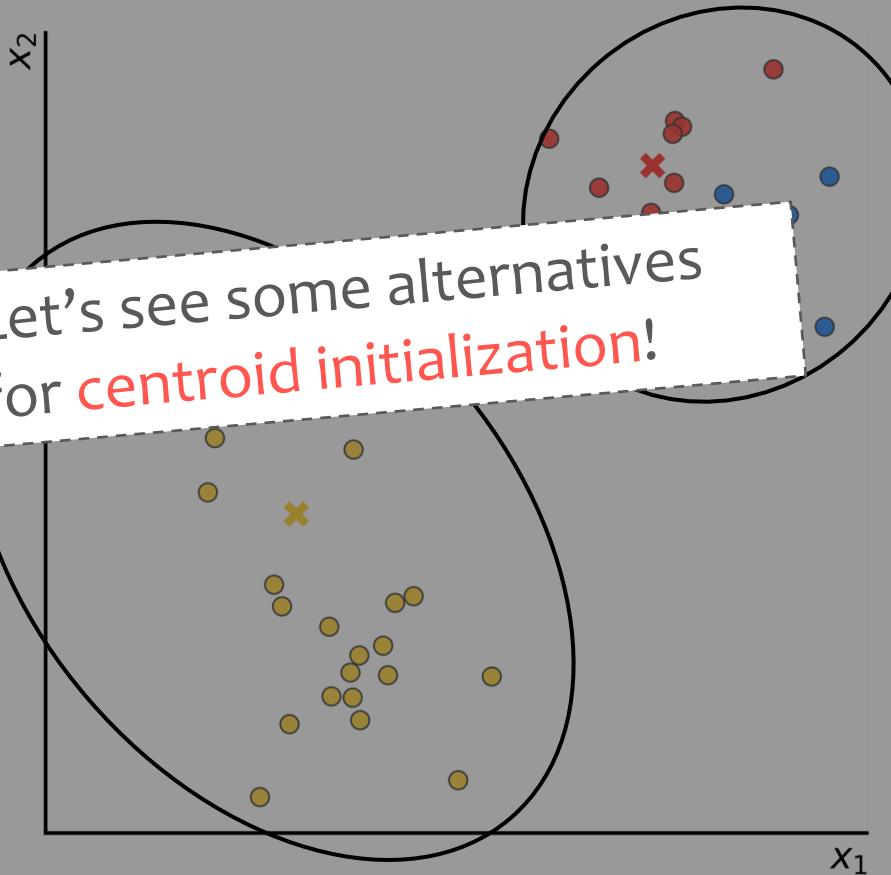
Two **clusters** is underpartitioned into a single cluster



If we are **unlucky**, our **random initialization** may select **poor initial centroids**.

Let's see some alternatives for **centroid initialization!**

A **cluster** is overpartitioned into two clusters



KMeans++

Initialization:

- 1) Take one centroid $c^{(1)}$, chosen uniformly at random from the dataset \mathbf{X} ;
- 2) Take a new centroid $c^{(l)}$, choosing an instance $x^{(i)} \in \mathbf{X}$ with probability $\frac{D(x^{(i)})^2}{\sum_{j=1}^m D(x^{(j)})^2}$;
 - where $D(x^{(i)})$ is the distance between the instance $x^{(i)}$ and **the nearest centroid that was already chosen**;
 - This probability distribution ensures that instances farther away from already chosen centroids are much more likely be selected as centroids.
- 3) Repeat the previous step until all **k centroids** have been chosen;

KMeans++

Initialization:

- 1) Take one centroid $c^{(1)}$, chosen uniformly at random from the dataset \mathbf{X} ;
- 2) Take a new centroid $c^{(l)}$, choosing an instance $x^{(i)} \in \mathbf{X}$ with probability $\frac{D(x^{(i)})^2}{\sum_{j=1}^m D(x^{(j)})^2}$;
 - where $D(x^{(i)})$ is the distance between the instance $x^{(i)}$ and **the nearest centroid that was already chosen**;
 - This probability distribution ensures that instances farther away from already chosen centroids are much more likely be selected as centroids.
- 3) Repeat the previous step until all **k centroids** have been chosen;



This is the **default initialization** for KMeans in Scikit-learn, but we can also force to use the **the original method** (i.e., picking k instances randomly as **initial centroids**).

Optimization Strategy

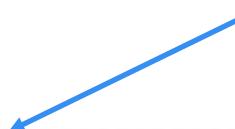
Another solution to get proper initial centroids and, consequently, good resulting clusters is to run KMeans multiple times with different random initializations and keep the best model, which is that with the lowest inertia.

Optimization Strategy

Another solution to get proper initial centroids and, consequently, good resulting clusters is to run KMeans multiple times with different random initializations and keep the best model, which is that with the lowest inertia.



Inertia is the mean squared distance between each instance and its closest centroid.



Optimization Strategy

Another solution to get proper initial centroids and, consequently, good resulting clusters is to run KMeans multiple times with different random initializations and keep the best model, which is that with the lowest inertia.



The KMeans implementation from Scikit-Learn runs the algorithm `n_init` times and keeps the model with the lowest inertia.



It uses KMeans++ for cluster initialization.



Inertia is the mean squared distance between each instance and its closest centroid.



Codes

Finding the optimal number of clusters

Finding the optimal number of clusters

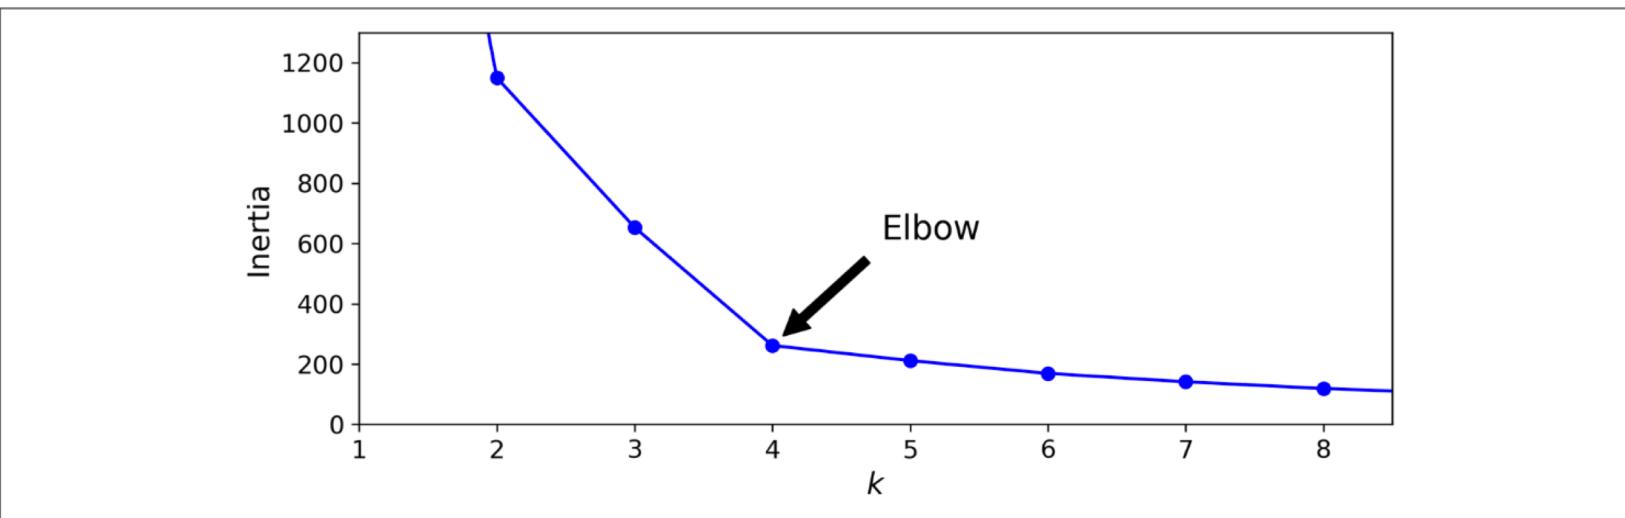


Figure 9-8. When plotting the inertia as a function of the number of clusters k , the curve often contains an inflection point called the “elbow”

https://github.com/ageron/handson-ml2/blob/master/09_unsupervised_learning.ipynb

Finding the optimal number of clusters

This technique for choosing the best value for the number of clusters is rather coarse.

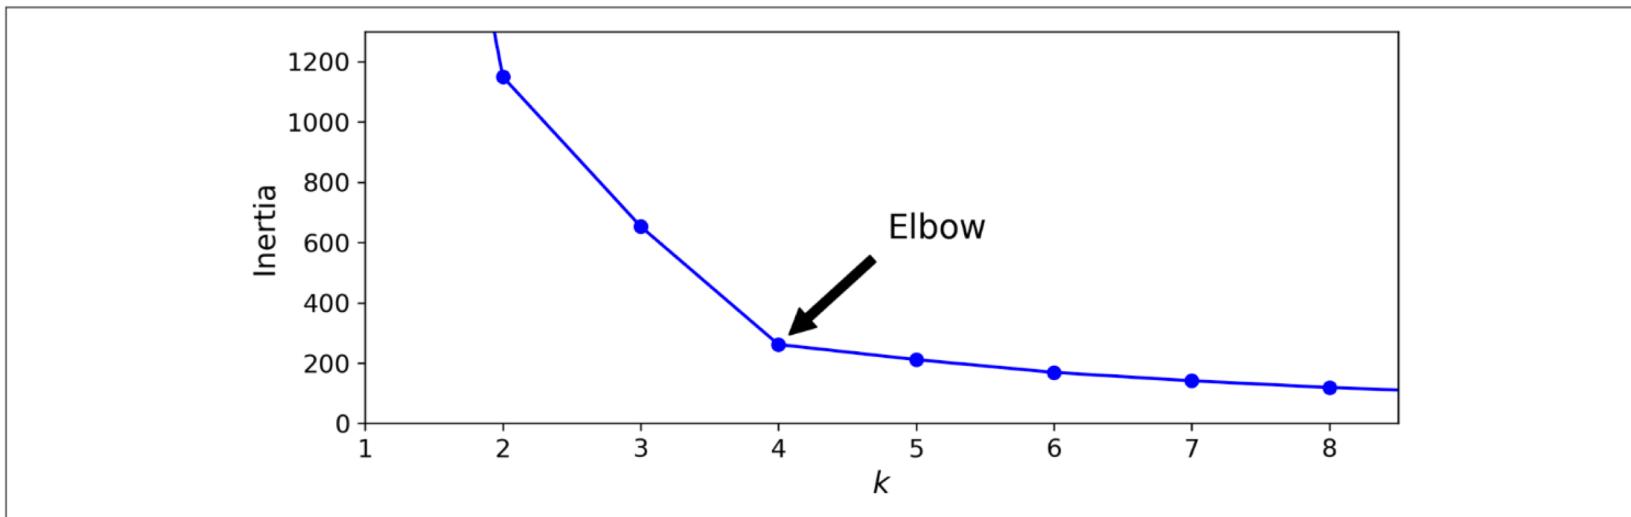


Figure 9-8. When plotting the inertia as a function of the number of clusters k , the curve often contains an inflection point called the “elbow”

https://github.com/ageron/handson-ml2/blob/master/09_unsupervised_learning.ipynb

Finding the optimal number of clusters

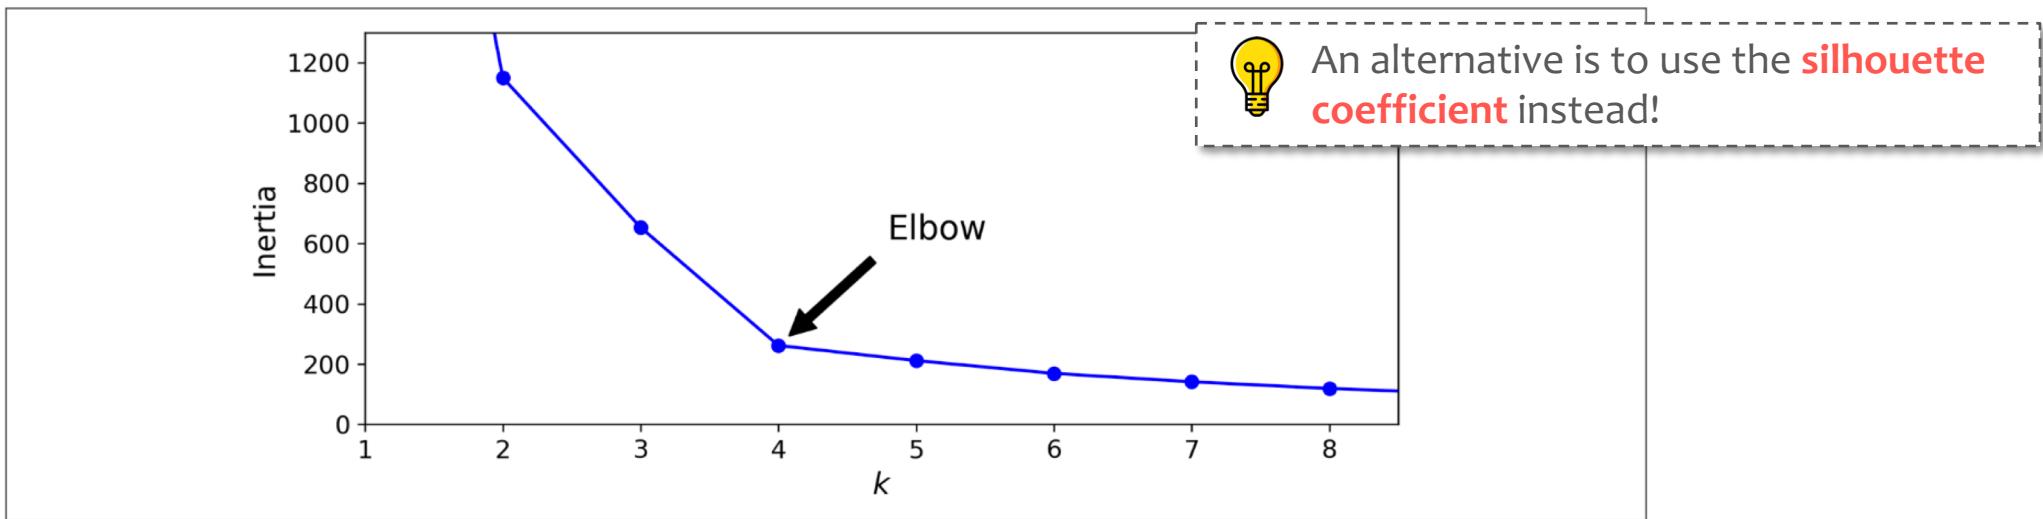
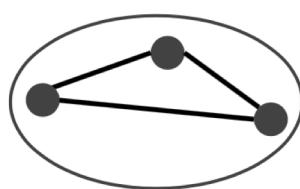


Figure 9-8. When plotting the inertia as a function of the number of clusters k , the curve often contains an inflection point called the “elbow”

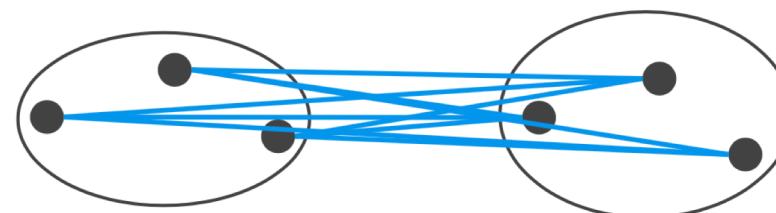
https://github.com/ageron/handson-ml2/blob/master/09_unsupervised_learning.ipynb

Silhouette Coefficient

The silhouette value is a measure of how similar a sample is to its own cluster (**cohesion**) compared to other clusters (**separation**)



Cohesion



Separation

Silhouette Coefficient

- The silhouette value is a measure of how similar a sample is to its own cluster (**cohesion**) compared to other clusters (**separation**)
- The silhouette ranges from -1 to $+1$
 - ▣ High value = the clustering configuration is appropriate
 - ▣ Low value = the clustering configuration may have too many or too few clusters

Silhouette Coefficient

- The Silhouette Coefficient is defined for each sample and is composed of two scores:
 - **a:** The mean distance between a sample and all other points in the same cluster
 - **b:** The mean distance between a sample and all other points in the next nearest cluster
- The Silhouette Coefficient s for a single sample is given as:
$$s = \frac{b - a}{\max(a, b)}$$
- The score is bounded between -1 for incorrect clustering and +1 for highly dense clustering ($a \ll b$).
- Scores around zero indicate overlapping clusters

Silhouette Coefficient

```
>>> from sklearn.metrics import silhouette_score  
>>> silhouette_score(X, kmeans.labels_)  
0.655517642572828
```

Let's compare the silhouette scores for different numbers of clusters (see Figure 9-9).

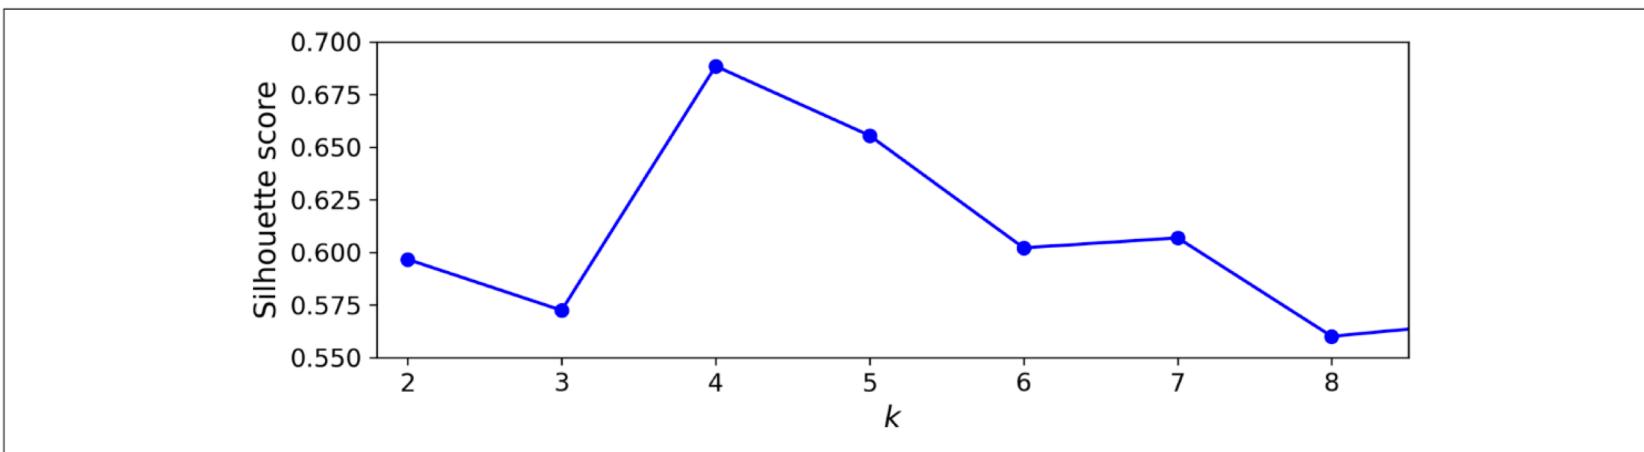


Figure 9-9. Selecting the number of clusters k using the silhouette score

Silhouette Coefficient

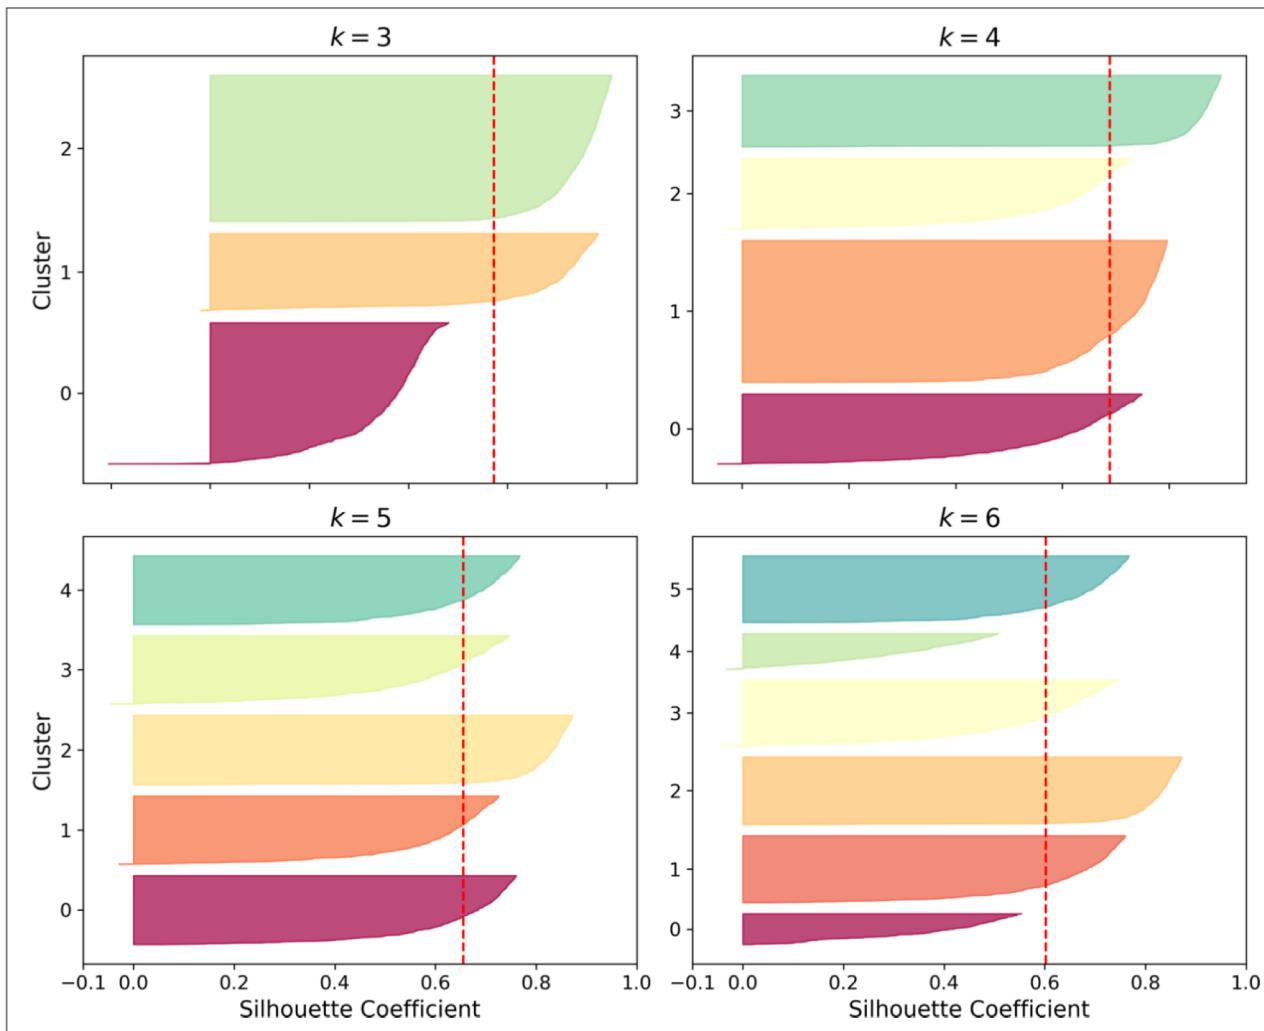


Figure 9-10. Analyzing the silhouette diagrams for various values of k

Aprendizado de Máquina e Reconhecimento de Padrões



Kmeans

Prof. Dr. Samuel Martins (Samuka)
samuel.martins@ifsp.edu.br

